

# Computer Architecture 1

Computer Organization and Design  
THE HARDWARE/SOFTWARE INTERFACE

[Adapted from *Computer Organization and Design, RISC-V Edition*, Patterson & Hennessy, © 2018, MK]

[Adapted from *Great ideas in Computer Architecture (CS 61C)* lecture slides, Garcia and Nikolic, © 2020, UC Berkeley]

# RISC-V Processor Design

# Great Idea #1: Abstraction (Levels of Representation/Interpretation)

High Level Language  
Program (e.g., C)

```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

Compiler

Assembly Language  
Program (e.g., RISC-V)

```
lw    x3, 0(x10)  
lw    x4, 4(x10)  
sw    x4, 0(x10)  
sw    x3, 4(x10)
```

Assembler

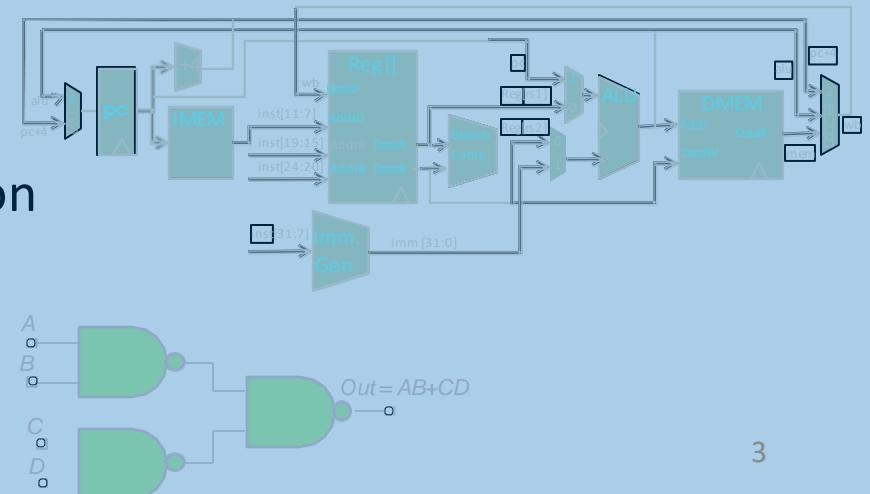
Machine Language  
Program (RISC-V)

|      |      |      |      |      |      |      |      |      |      |
|------|------|------|------|------|------|------|------|------|------|
| 1000 | 1101 | 1110 | 0010 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 |
| 1000 | 1110 | 0001 | 0000 | 0000 | 0000 | 0000 | 0000 | 0100 |      |
| 1010 | 1110 | 0001 | 0010 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 |
| 1010 | 1101 | 1110 | 0010 | 0000 | 0000 | 0000 | 0000 | 0100 |      |

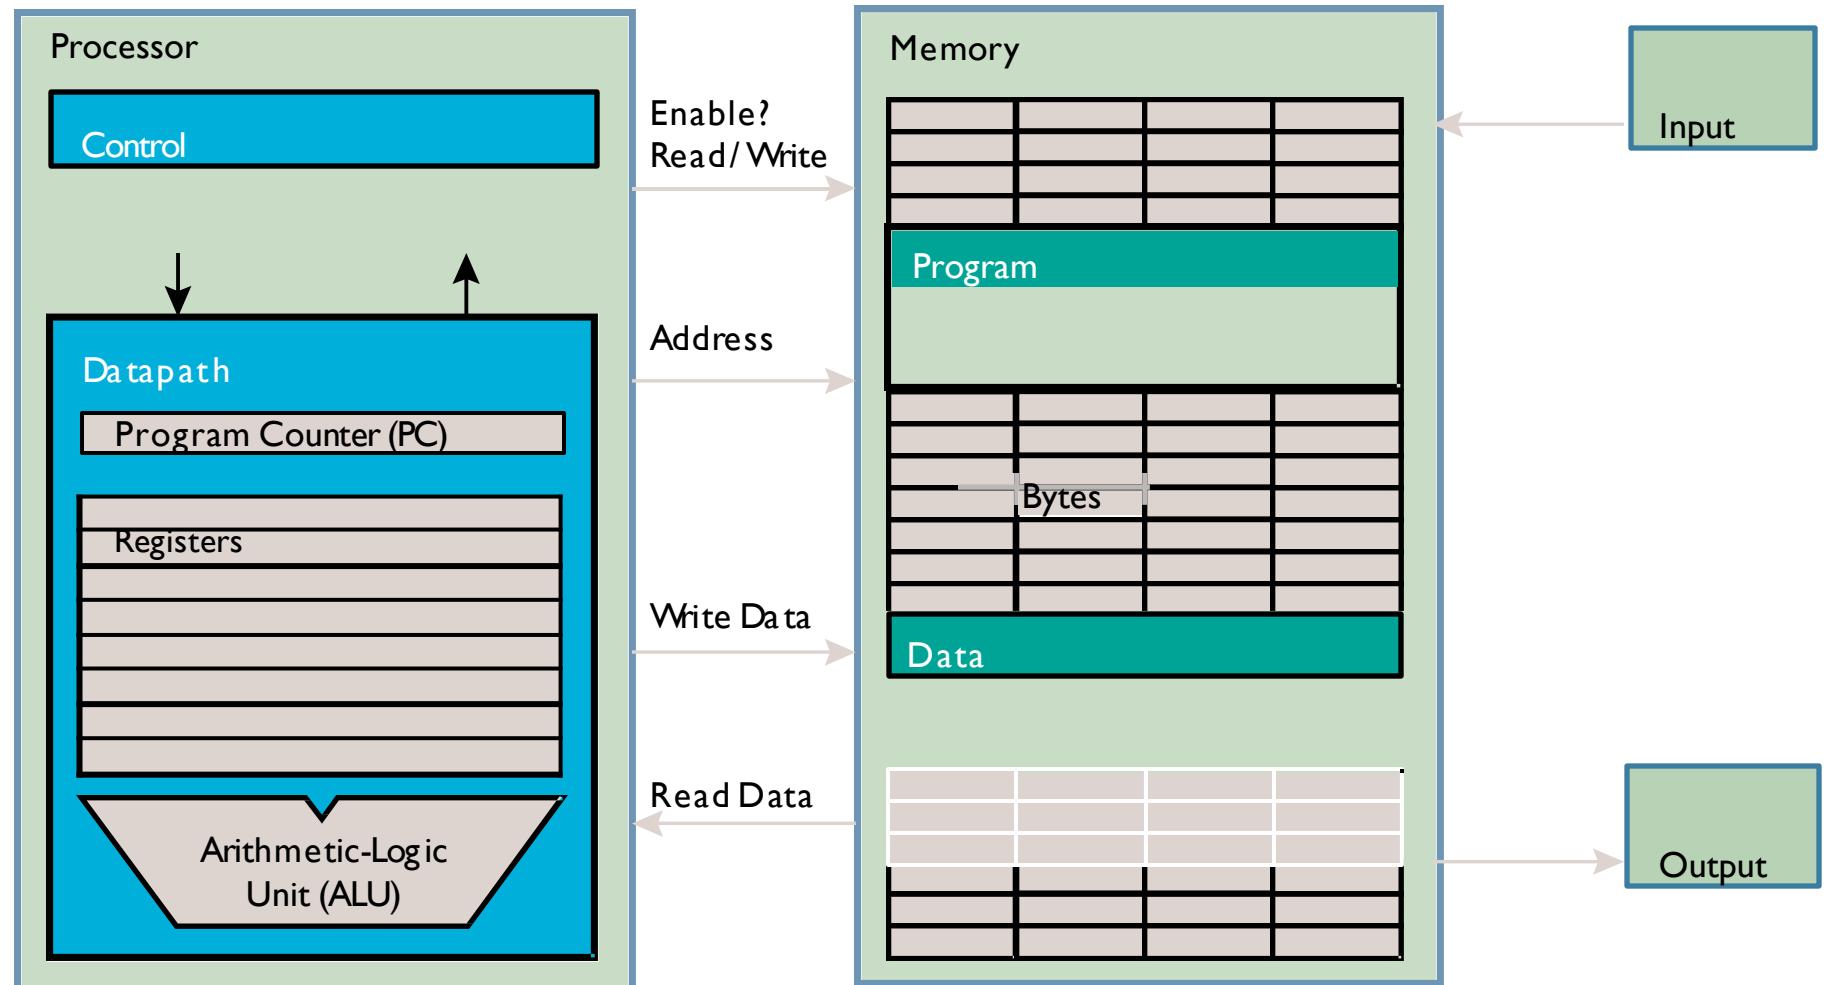
Hardware Architecture Description  
(e.g., block diagrams)

Architecture Implementation

Logic Circuit Description  
(Circuit Schematic Diagrams)



# Our Single-Core Processor So Far...



# The CPU

- **Processor (CPU):** the active part of the computer that does all the work (data manipulation and decision-making)
- **Datapath:** portion of the processor that contains hardware necessary to perform operations required by the processor (the brawn)
- **Control:** portion of the processor (also in hardware) that tells the datapath what needs to be done (the brain)

# Need to Implement All RV32I Instructions

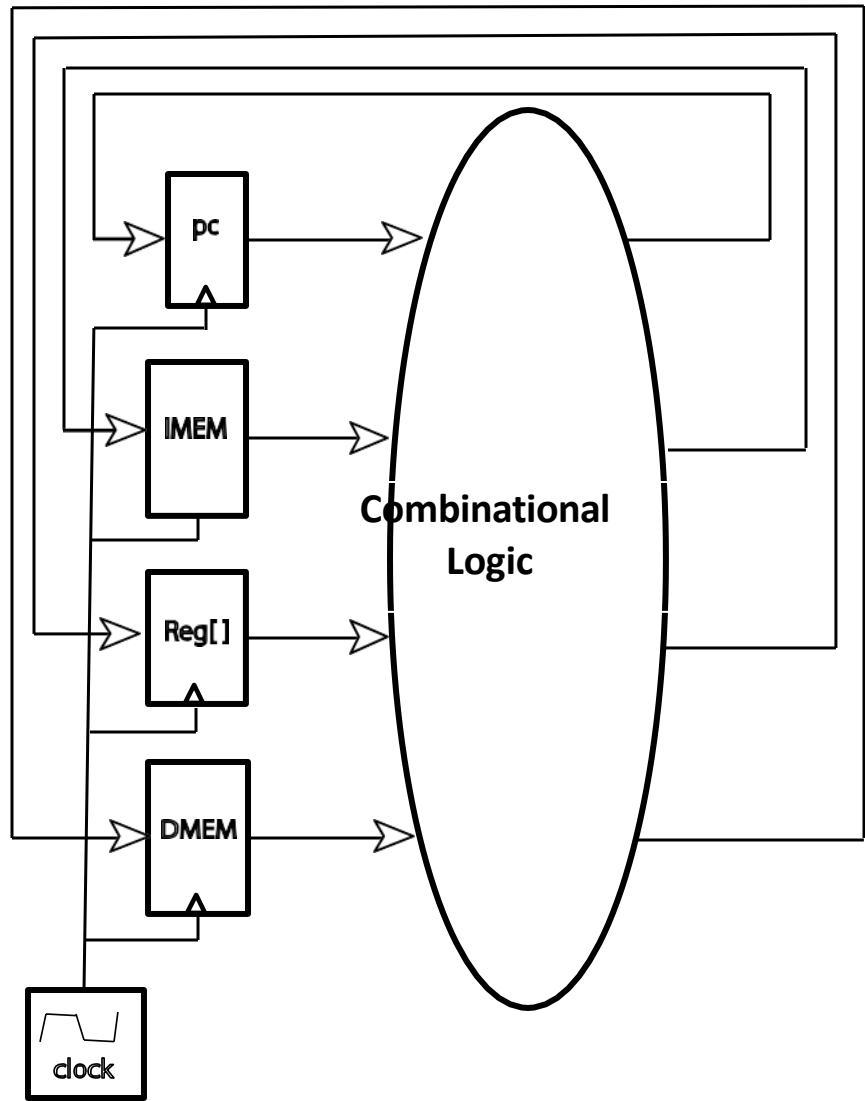


| Open Reference Card              |                         |     |                   |             |                      |     |                  |
|----------------------------------|-------------------------|-----|-------------------|-------------|----------------------|-----|------------------|
| Base Integer Instructions: RV32I |                         |     |                   |             |                      |     |                  |
| Category                         | Name                    | Fmt | RV32I Base        | Category    | Name                 | Fmt | RV32I Base       |
| Shifts                           | Shift Left Logical      | R   | SLL rd,rs1,rs2    | Loads       | Load Byte            | I   | LB rd,rs1,imm    |
|                                  | Shift Left Log. Imm.    | I   | SLLI rd,rs1,shamt |             | Load Halfword        | I   | LH rd,rs1,imm    |
|                                  | Shift Right Logical     | R   | SRL rd,rs1,rs2    |             | Load Byte Unsigned   | I   | LBU rd,rs1,imm   |
|                                  | Shift Right Log. Imm.   | I   | SRLI rd,rs1,shamt |             | Load Half Unsigned   | I   | LHU rd,rs1,imm   |
|                                  | Shift Right Arithmetic  | R   | SRA rd,rs1,rs2    |             | Load Word            | I   | LW rd,rs1,imm    |
|                                  | Shift Right Arith. Imm. | I   | SRAI rd,rs1,shamt | Stores      | Store Byte           | S   | SB rs1,rs2,imm   |
| Arithmetic                       | ADD                     | R   | ADD rd,rs1,rs2    |             | Store Halfword       | S   | SH rs1,rs2,imm   |
|                                  | ADD Immediate           | I   | ADDI rd,rs1,imm   |             | Store Word           | S   | SW rs1,rs2,imm   |
|                                  | SUBtract                | R   | SUB rd,rs1,rs2    | Branches    | Branch =             | B   | BEQ rs1,rs2,imm  |
|                                  | Load Upper Imm          | U   | LUI rd,imm        |             | Branch ≠             | B   | BNE rs1,rs2,imm  |
|                                  | Add Upper Imm to PC     | U   | AUIPC rd,imm      |             | Branch <             | B   | BLT rs1,rs2,imm  |
| Logical                          | XOR                     | R   | XOR rd,rs1,rs2    |             | Branch ≥             | B   | BGE rs1,rs2,imm  |
|                                  | XOR Immediate           | I   | XORI rd,rs1,imm   |             | Branch < Unsigned    | B   | BLTU rs1,rs2,imm |
|                                  | OR                      | R   | OR rd,rs1,rs2     |             | Branch ≥ Unsigned    | B   | BGEU rs1,rs2,imm |
|                                  | OR Immediate            | I   | ORI rd,rs1,imm    | Jump & Link | J&L                  | J   | JAL rd,imm       |
|                                  | AND                     | R   | AND rd,rs1,rs2    |             | Jump & Link Register | I   | JALR rd,rs1,imm  |
|                                  | AND Immediate           | I   | ANDI rd,rs1,imm   |             |                      |     |                  |
| Compare                          | Set <                   | R   | SLT rd,rs1,rs2    | Synch       | Synch thread         | I   | FENCE            |
|                                  | Set < Immediate         | I   | SLTI rd,rs1,imm   |             |                      |     | Not in           |
|                                  | Set < Unsigned          | R   | SLTU rd,rs1,rs2   | Environment | CALL                 | I   | ECALL            |
|                                  | Set < Imm Unsigned      | I   | SLTIU rd,rs1,imm  |             | BREAK                | I   | EBREAK           |



# **Building a RISC-V Processor**

# One-Instruction-Per-Cycle RISC-V Machine



- On every tick of the clock, the computer executes one instruction
- Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
- At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle

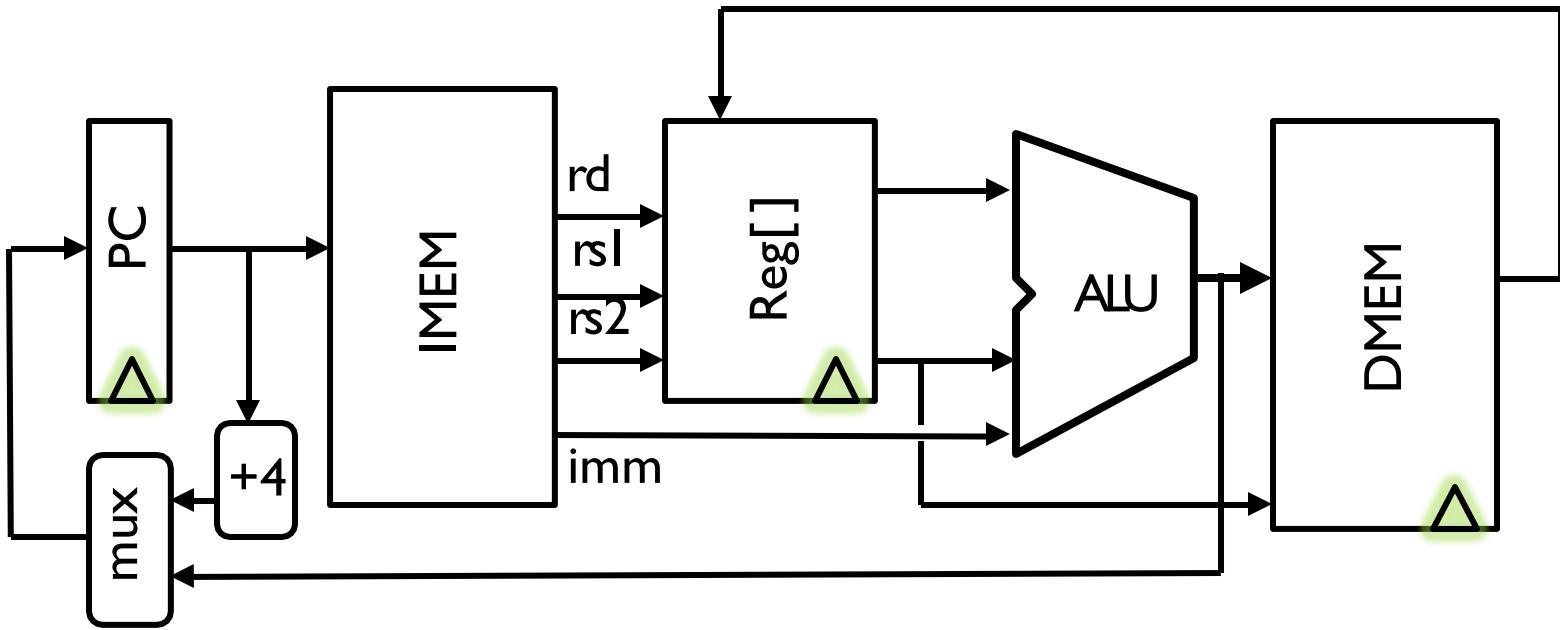
# Stages of the Datapath : Overview

- Problem: a single, “monolithic” block that “executes an instruction” (performs all necessary operations beginning with fetching the instruction) would be too bulky and inefficient
- Solution: break up the process of “executing an instruction” into stages, and then connect the stages to create the whole datapath
  - smaller stages are easier to design
  - easy to optimize (change) one stage without touching the others (modularity)

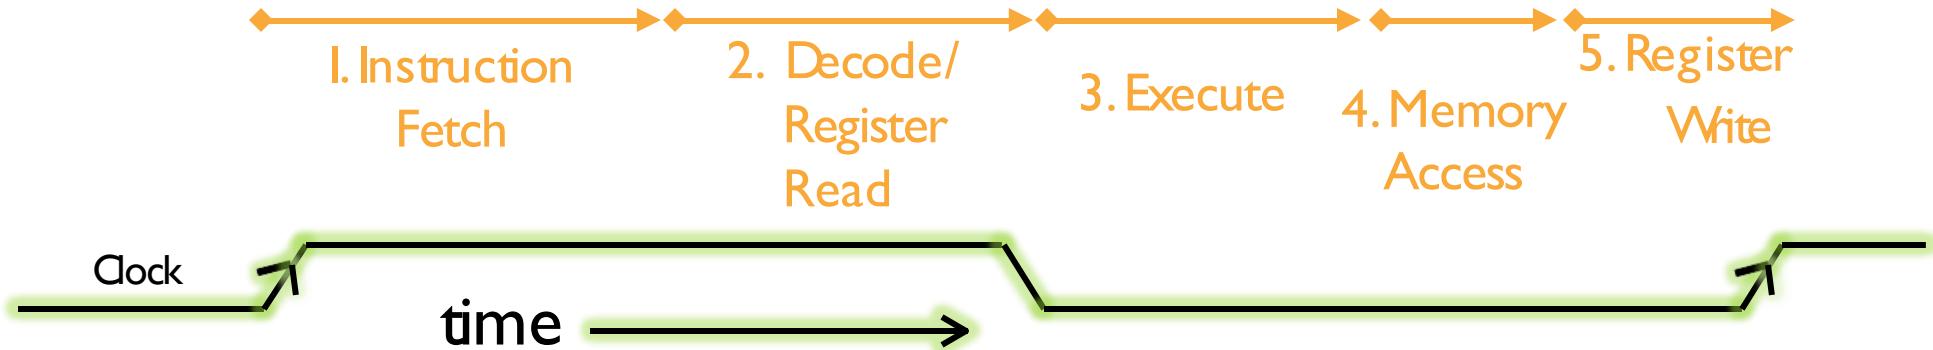
# Five Stages of the Datapath

- Stage 1: *Instruction Fetch (IF)*
- Stage 2: *Instruction Decode (ID)*
- Stage 3: *Execute (EX) - ALU (Arithmetic-Logic Unit)*
- Stage 4: *Memory Access (MEM)*
- Stage 5: *Write Back to Register (WB)*

# Basic Phases of Instruction Execution

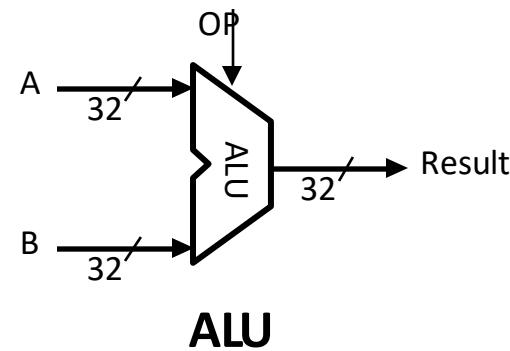
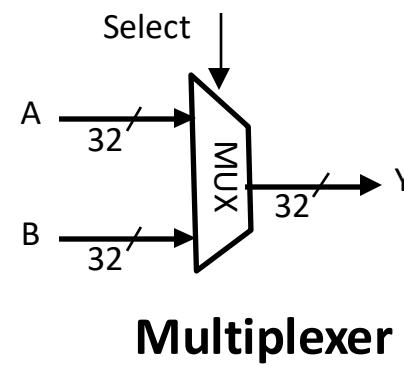
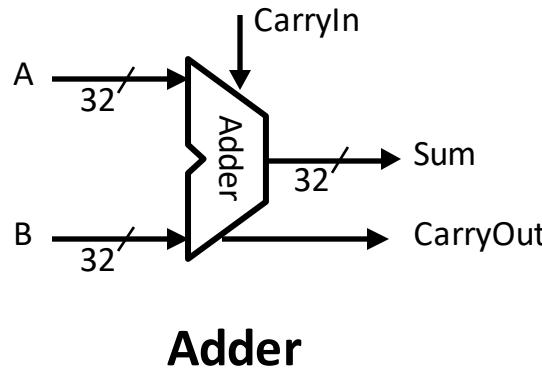


◆ I. Instruction Fetch  
◆ 2. Decode/ Register Read  
◆ 3. Execute  
◆ 4. Memory Access  
◆ 5. Register Write



# Datapath Components: Combinational

- Combinational elements



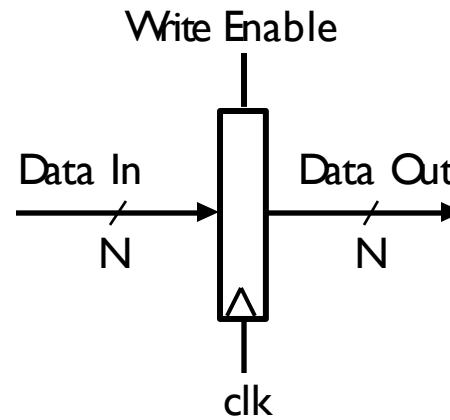
- Storage elements + clocking methodology
- Building blocks

# Datapath Elements: State and Sequencing (1/3)

- Register

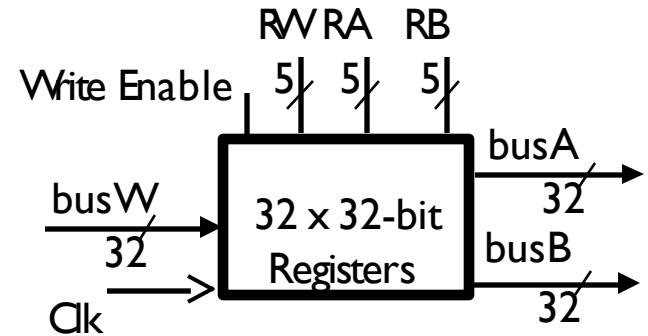
- Write Enable:

- Low (or deasserted) (0): Data Out will not change
- Asserted (1): Data Out will become Data In on positive edge of clock



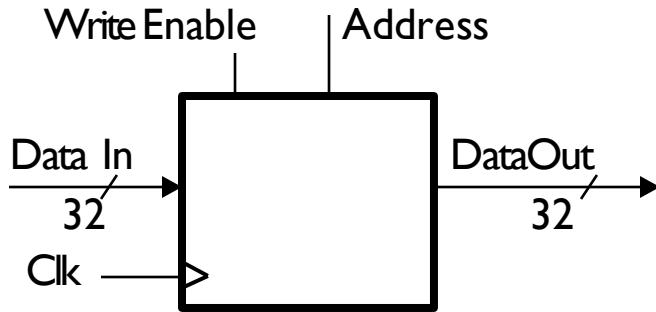
# Datapath Elements: State and Sequencing (2/3)

- Register file (regfile, RF) consists of 32 registers:
  - Two 32-bit output busses: busA and busB
  - One 32-bit input bus: busW
- Register is selected by:
  - RA (number) selects the register to put on busA (data)
  - RB (number) selects the register to put on busB (data)
  - RW (number) selects the register to be written via busW (data) when Write Enable is 1
- Clock input (Clk)
  - Clk input is a factor ONLY during write operation
  - During read operation, behaves as a combinational logic block:
    - RA or RB valid  $\Rightarrow$  busA or busB valid after “access time.”



# Datapath Elements: State and Sequencing (3/3)

- “Magic” Memory
  - One input bus: Data In
  - One output bus: Data Out
- Memory word is found by:
  - For Read: Address selects the word to put on Data Out
  - For Write: Set Write Enable = 1: address selects the memory word to be written via the Data In bus
- Clock input (CLK)
  - CLK input is a factor ONLY during write operation
  - During read operation, behaves as a combinational logic block: Address valid  $\Rightarrow$  Data Out valid after “access time”



# State Required by RV32I ISA (1/2)

Each instruction during execution reads and updates the state of :

(1) Registers, (2) Program counter, (3) Memory

- Registers (**x0 . . x31**)
  - Register file (*regfile*) **Reg** holds 32 registers x 32 bits/register:  
**Reg [0] . . Reg [31]**
  - First register read specified by **rs1** field in instruction
  - Second register read specified by **rs2** field in instruction
  - Write register (destination) specified by *rd* field in instruction
  - **x0** is always 0 (writes to **Reg [0]** are ignored)
- Program Counter (**PC**)
  - Holds address of current instruction

# State Required by RV32I ISA (2/2)

- Memory (**MEM**)
  - Holds both instructions & data, in one 32-bit byte-addressed memory space
  - We'll use separate memories for instructions (**IMEM**) and data (**DMEM**)
    - *These are placeholders for instruction and data caches*
  - Instructions are read (*fetched*) from instruction memory (assume **IMEM** read-only)
  - Load/store instructions access data memory

# R-Type Add Datapath

# Review: R-Type Instructions

| 31             | 30  | 29 | 28 | 27 | 26      | 25  | 24 | 23 | 22 | 21      | 20         | 19 | 18 | 17 | 16      | 15 | 14 | 13 | 12 | 10     | 9            | 8 | 7 | 6 | 5      | 4 | 3 | 2 | 1 | 0 |
|----------------|-----|----|----|----|---------|-----|----|----|----|---------|------------|----|----|----|---------|----|----|----|----|--------|--------------|---|---|---|--------|---|---|---|---|---|
| R-format : ALU |     |    |    |    |         |     |    |    |    |         |            |    |    |    |         |    |    |    |    |        |              |   |   |   |        |   |   |   |   |   |
| [31:25]        |     |    |    |    | [24:20] |     |    |    |    | [19:15] |            |    |    |    | [14:12] |    |    |    |    | [11:7] |              |   |   |   | [6:0]  |   |   |   |   |   |
| 7              |     |    |    |    | 5       |     |    |    |    | 5       |            |    |    |    | 3       |    |    |    |    | 5      |              |   |   |   | 7      |   |   |   |   |   |
| func7          |     |    |    |    | rs2     |     |    |    |    | rs1     |            |    |    |    | func3   |    |    |    |    | rd     |              |   |   |   | opcode |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 000 : ADD  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0100000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 000 : SUB  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 001 : SLL  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 010 : SLT  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 011 : SLTU |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 100 : XOR  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 101 : SRL  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0100000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 101 : SRA  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 110 : OR   |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |
| 0000000        | rs2 |    |    |    |         | rs1 |    |    |    |         | 111 : AND  |    |    |    |         | rd |    |    |    |        | 0110011:OP-R |   |   |   |        |   |   |   |   |   |

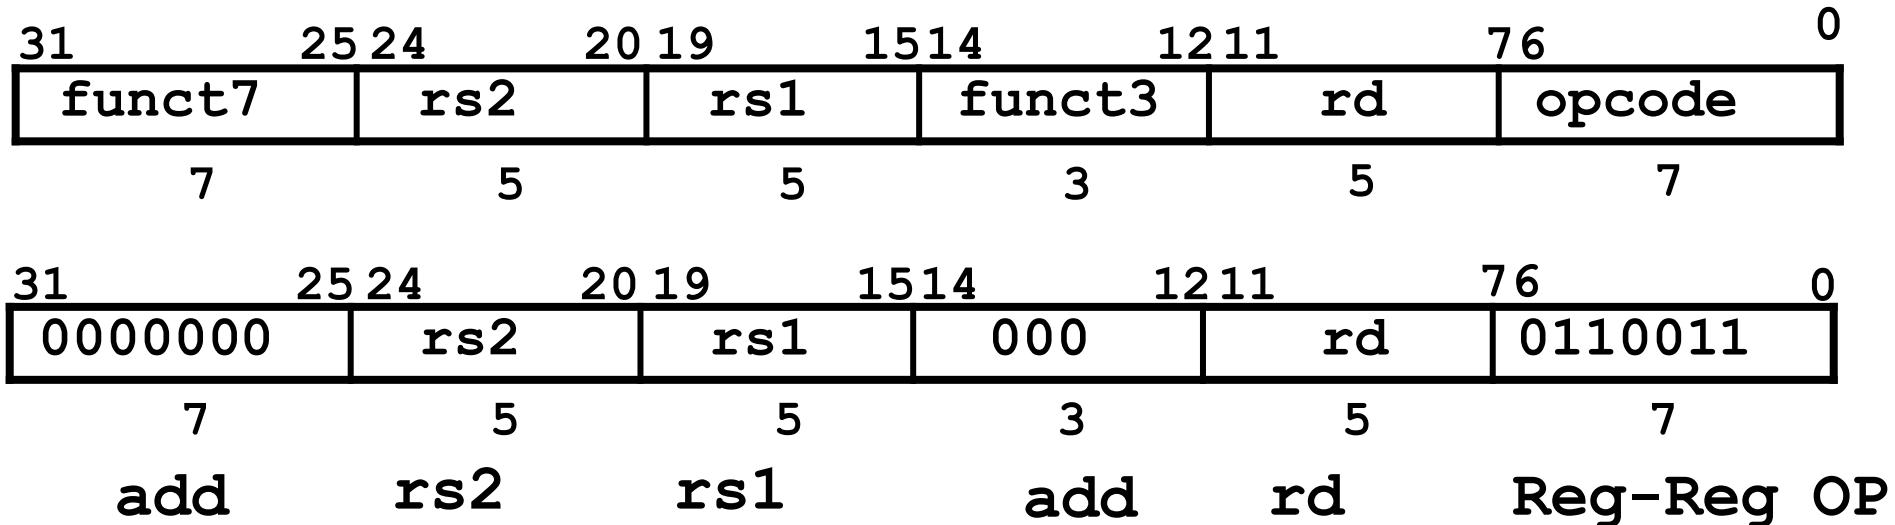
- E.g. Addition/subtraction `add rd, rs1, rs2`

$$R[rd] = R[rs1] + R[rs2]$$

`sub rd, rs1, rs2`

$$R[rd] = R[rs1] - R[rs2]$$

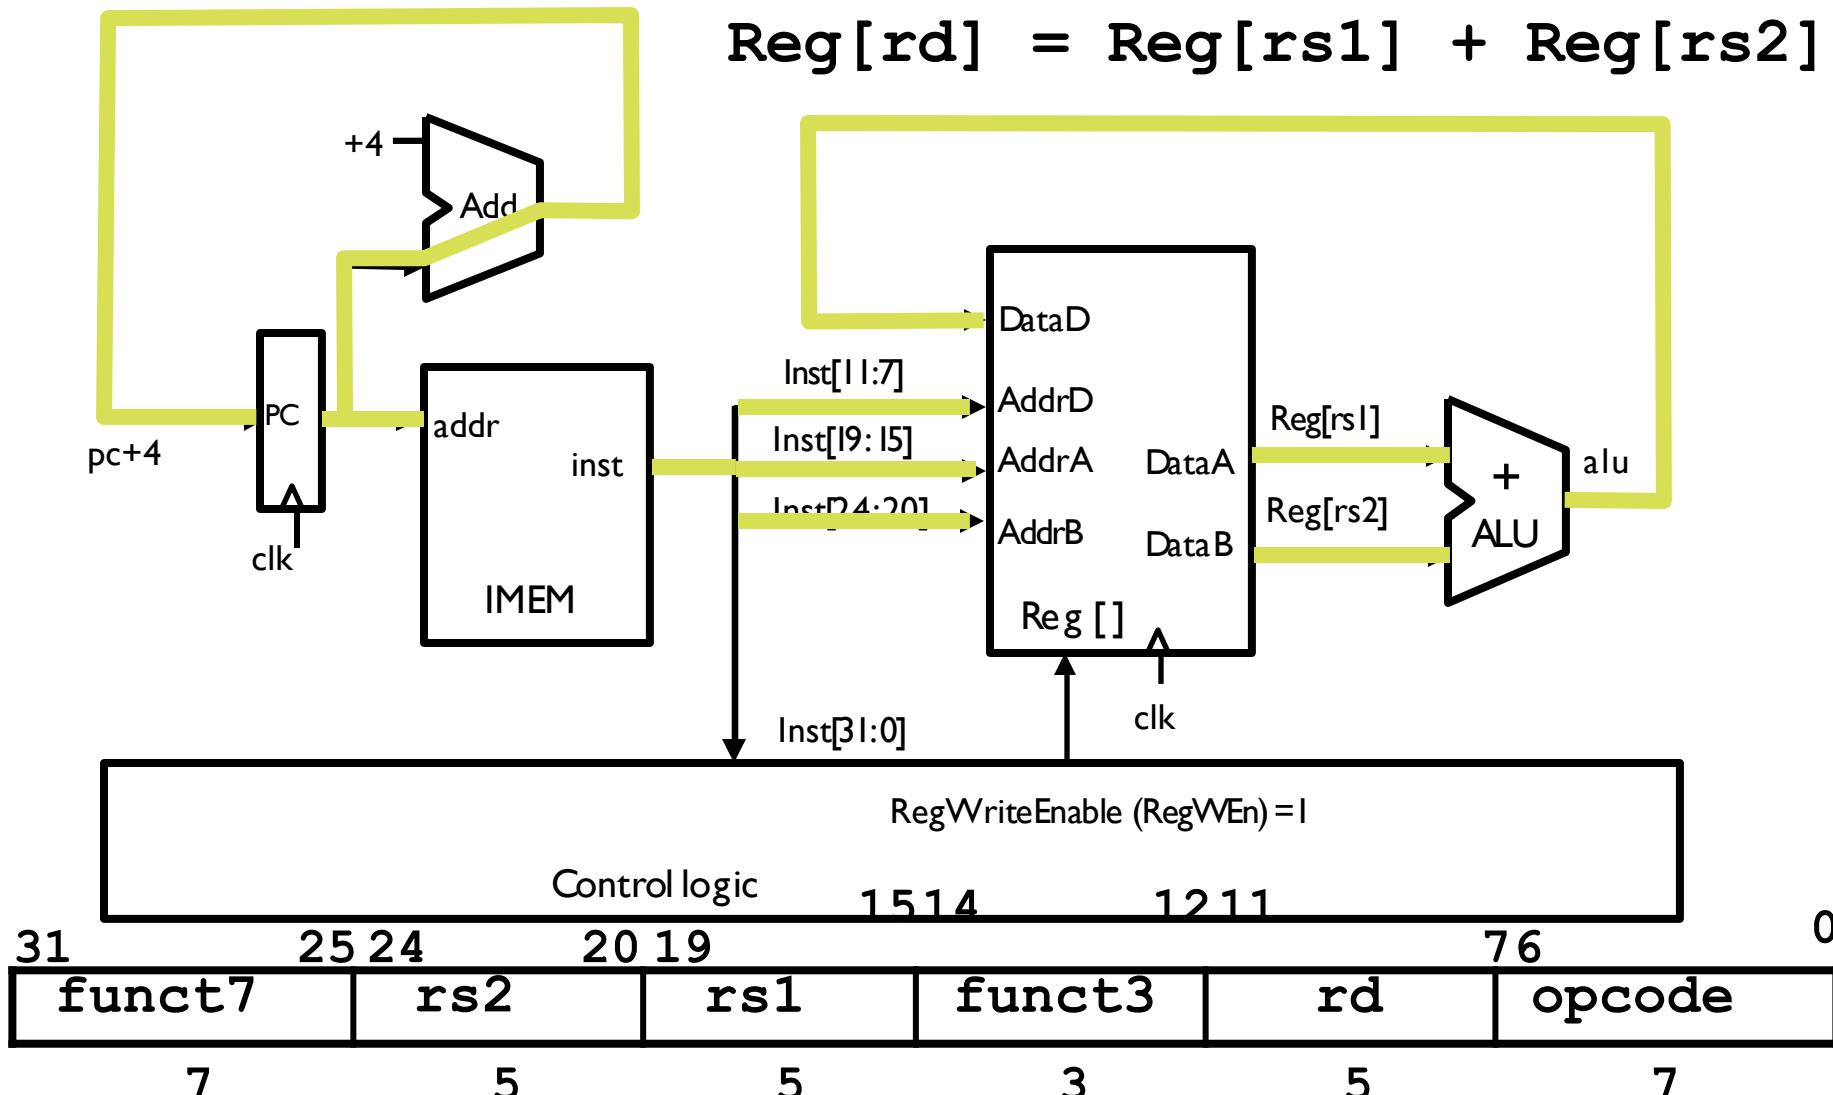
# Implementing the add instruction



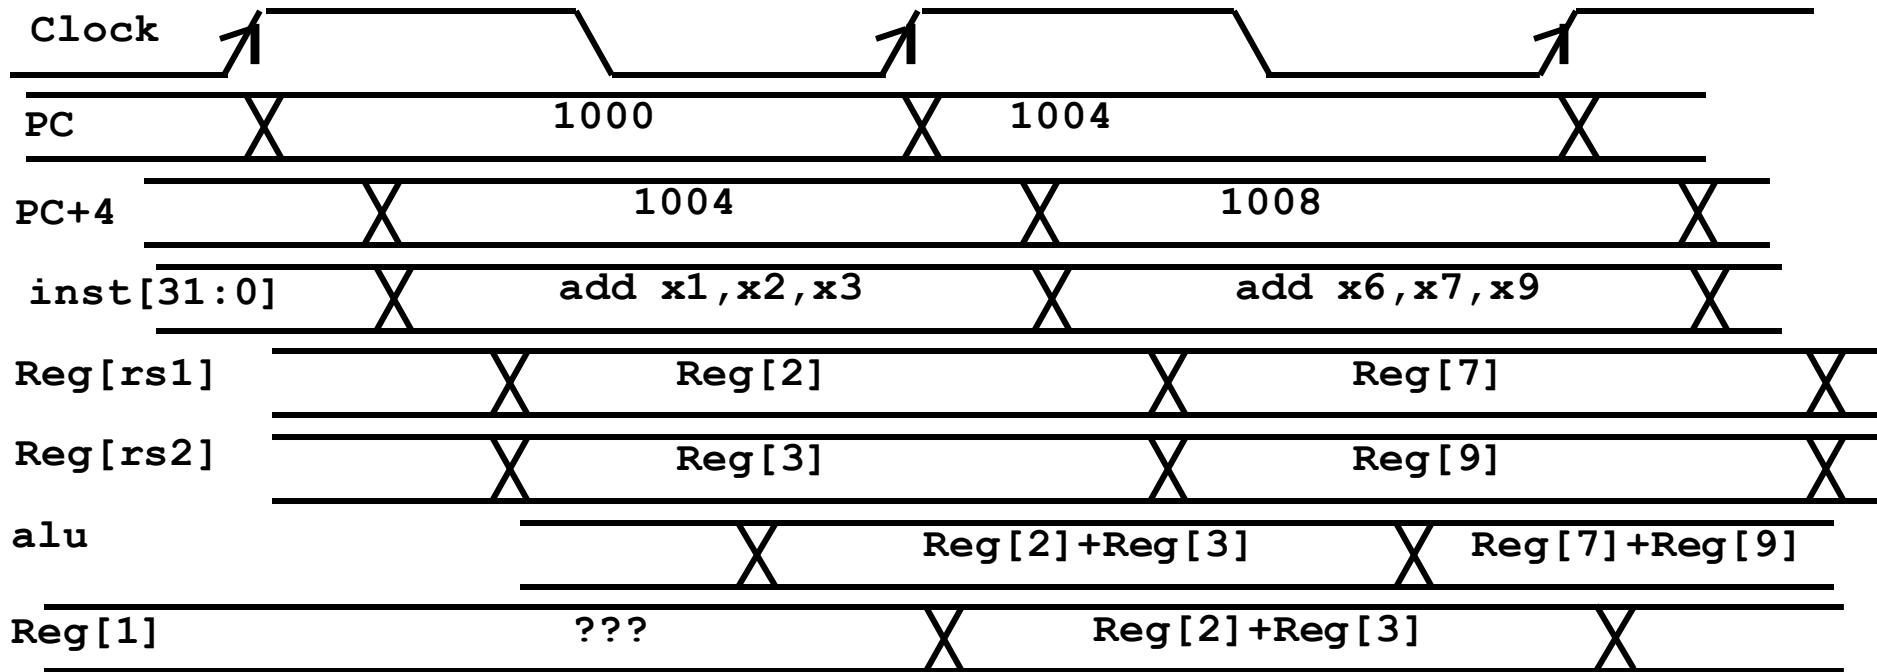
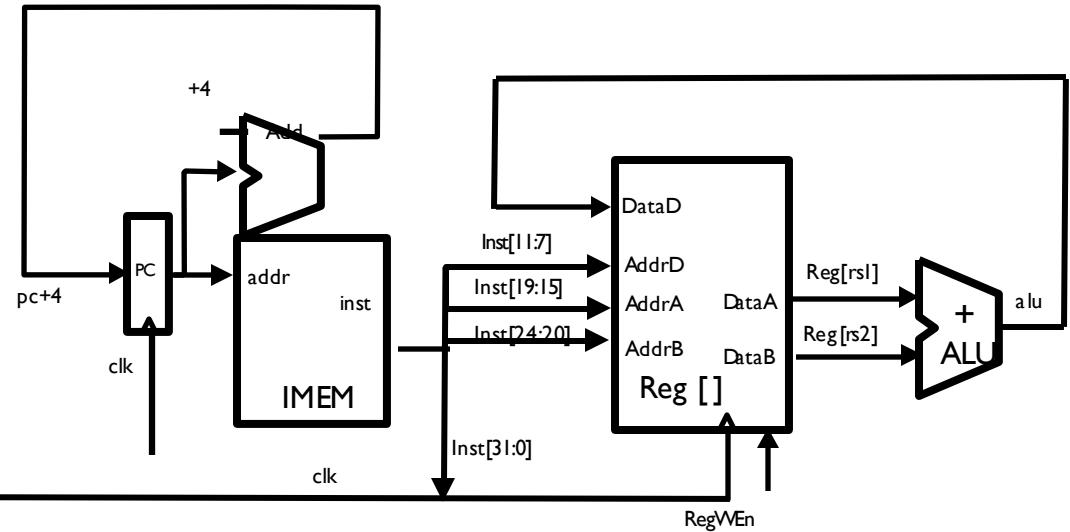
add rd, rs1, rs2

- Instruction makes two changes to machine's state:
  - $\text{Reg}[rd] = \text{Reg}[rs1] + \text{Reg}[rs2]$
  - $\text{PC} = \text{PC} + 4$

# Datapath for add



# Timing Diagram for add



time →

# **Sub Datapath**

# Implementing the sub instruction

|         |     |     |     |    |         |     |
|---------|-----|-----|-----|----|---------|-----|
| 0000000 | rs2 | rs1 | 000 | rd | 0110011 | add |
| 0100000 | rs2 | rs1 | 000 | rd | 0110011 | sub |

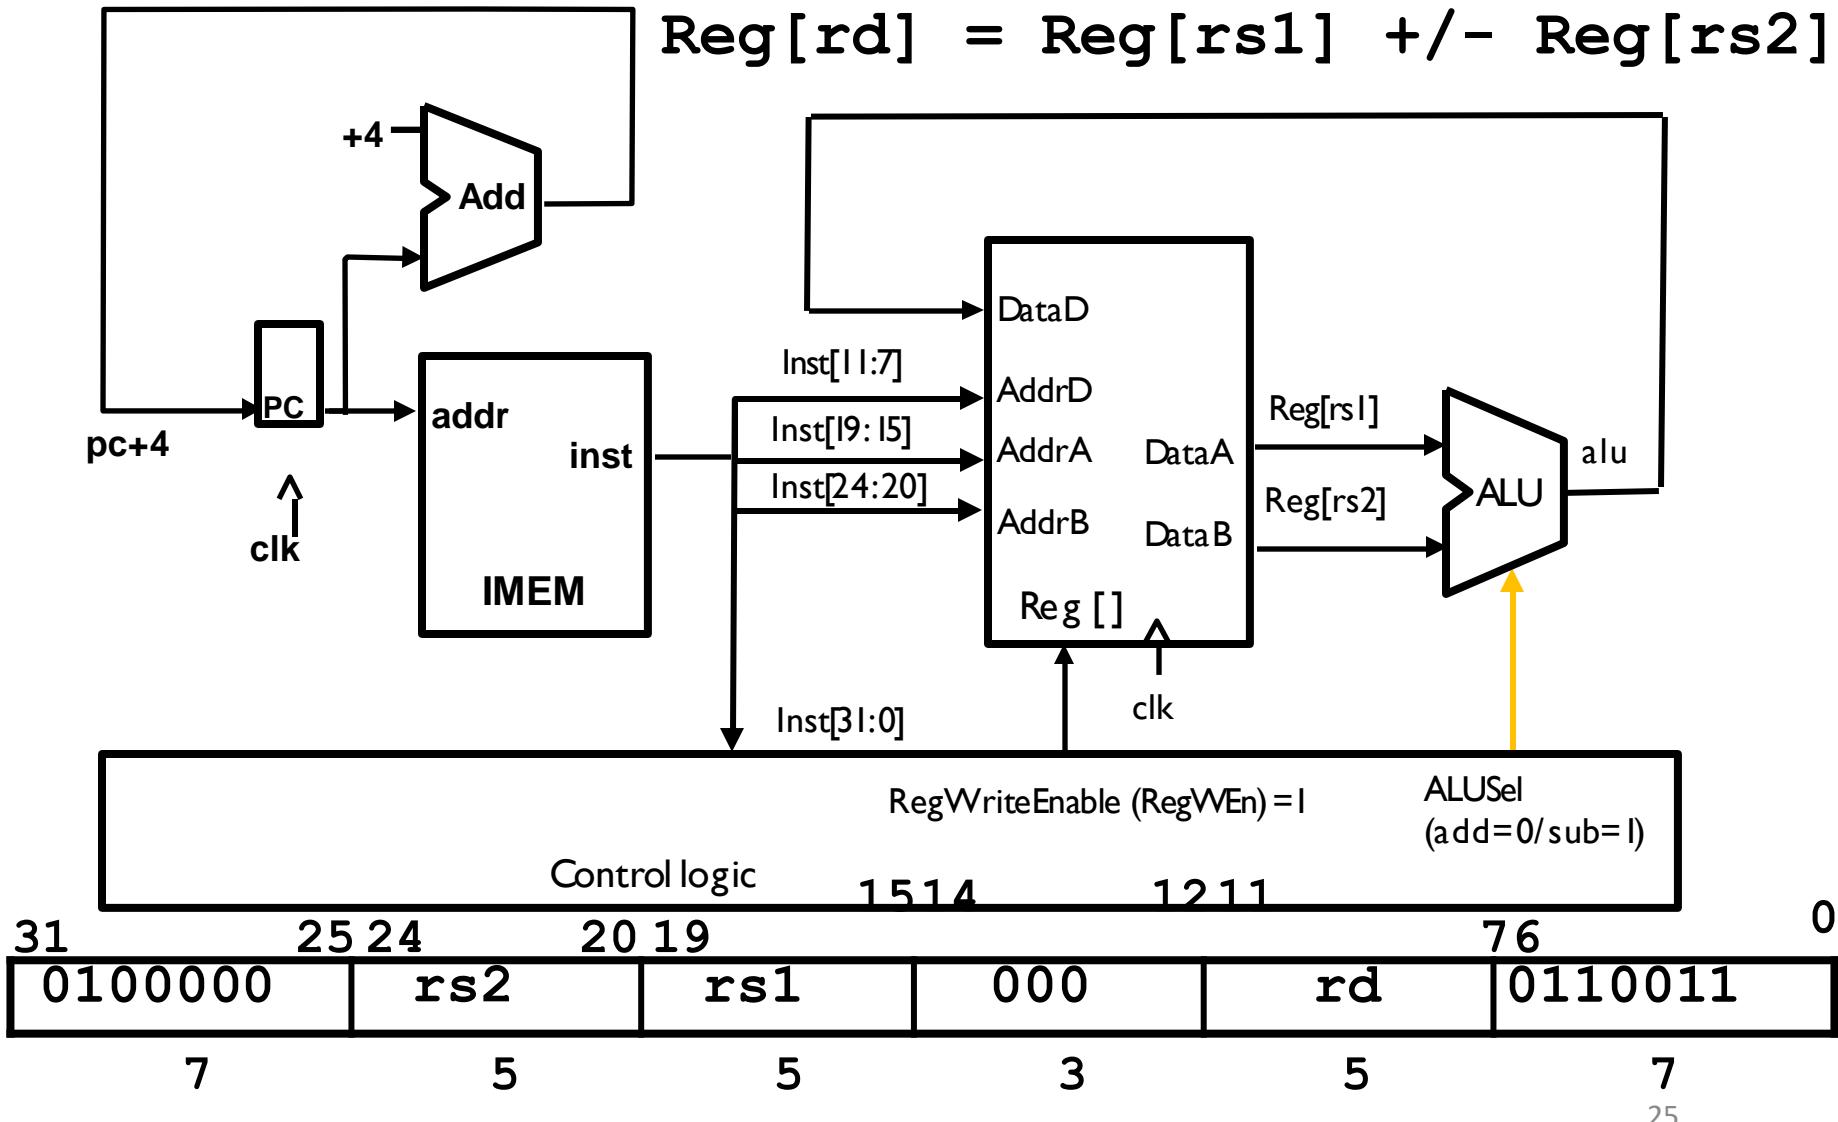
**sub rd, rs1, rs2**

- Almost the same as add, except now have to subtract operands instead of adding them
- **inst[30]** selects between add and subtract

# Datapath for add/sub

$$PC = PC + 4$$

$$Reg[rd] = Reg[rs1] +/- Reg[rs2]$$



# Implementing Other R-Format Instructions

|         |     |     |     |    |         |      |
|---------|-----|-----|-----|----|---------|------|
| 0000000 | rs2 | rs1 | 000 | rd | 0110011 | add  |
| 0100000 | rs2 | rs1 | 000 | rd | 0110011 | sub  |
| 0000000 | rs2 | rs1 | 001 | rd | 0110011 | sll  |
| 0000000 | rs2 | rs1 | 010 | rd | 0110011 | slt  |
| 0000000 | rs2 | rs1 | 011 | rd | 0110011 | sltu |
| 0000000 | rs2 | rs1 | 100 | rd | 0110011 | xor  |
| 0000000 | rs2 | rs1 | 101 | rd | 0110011 | srl  |
| 0100000 | rs2 | rs1 | 101 | rd | 0110011 | sra  |
| 0000000 | rs2 | rs1 | 110 | rd | 0110011 | or   |
| 0000000 | rs2 | rs1 | 111 | rd | 0110011 | and  |

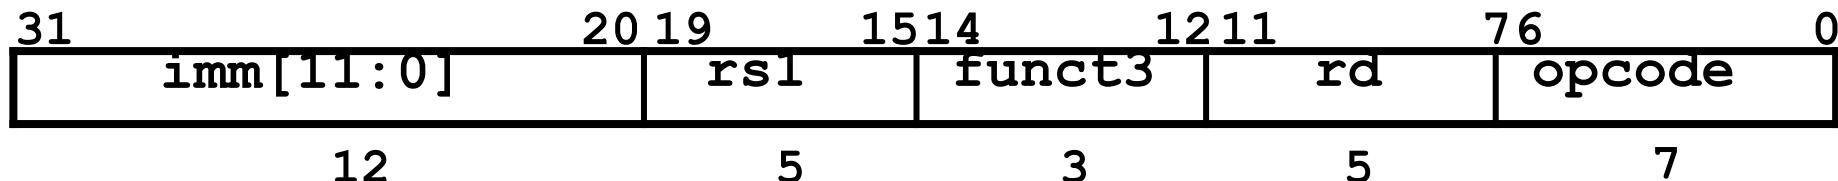
All implemented by decoding funct3 and funct7 fields  
and selecting appropriate ALU function

# Datapath With Immediates

# Implementing I-Format - addi instruction

- RISC-V Assembly Instruction:

**addi x15,x1,-50**



|              |       |     |       |         |
|--------------|-------|-----|-------|---------|
| 111111001110 | 00001 | 000 | 01111 | 0010011 |
|--------------|-------|-----|-------|---------|

imm=-50

rs1=1

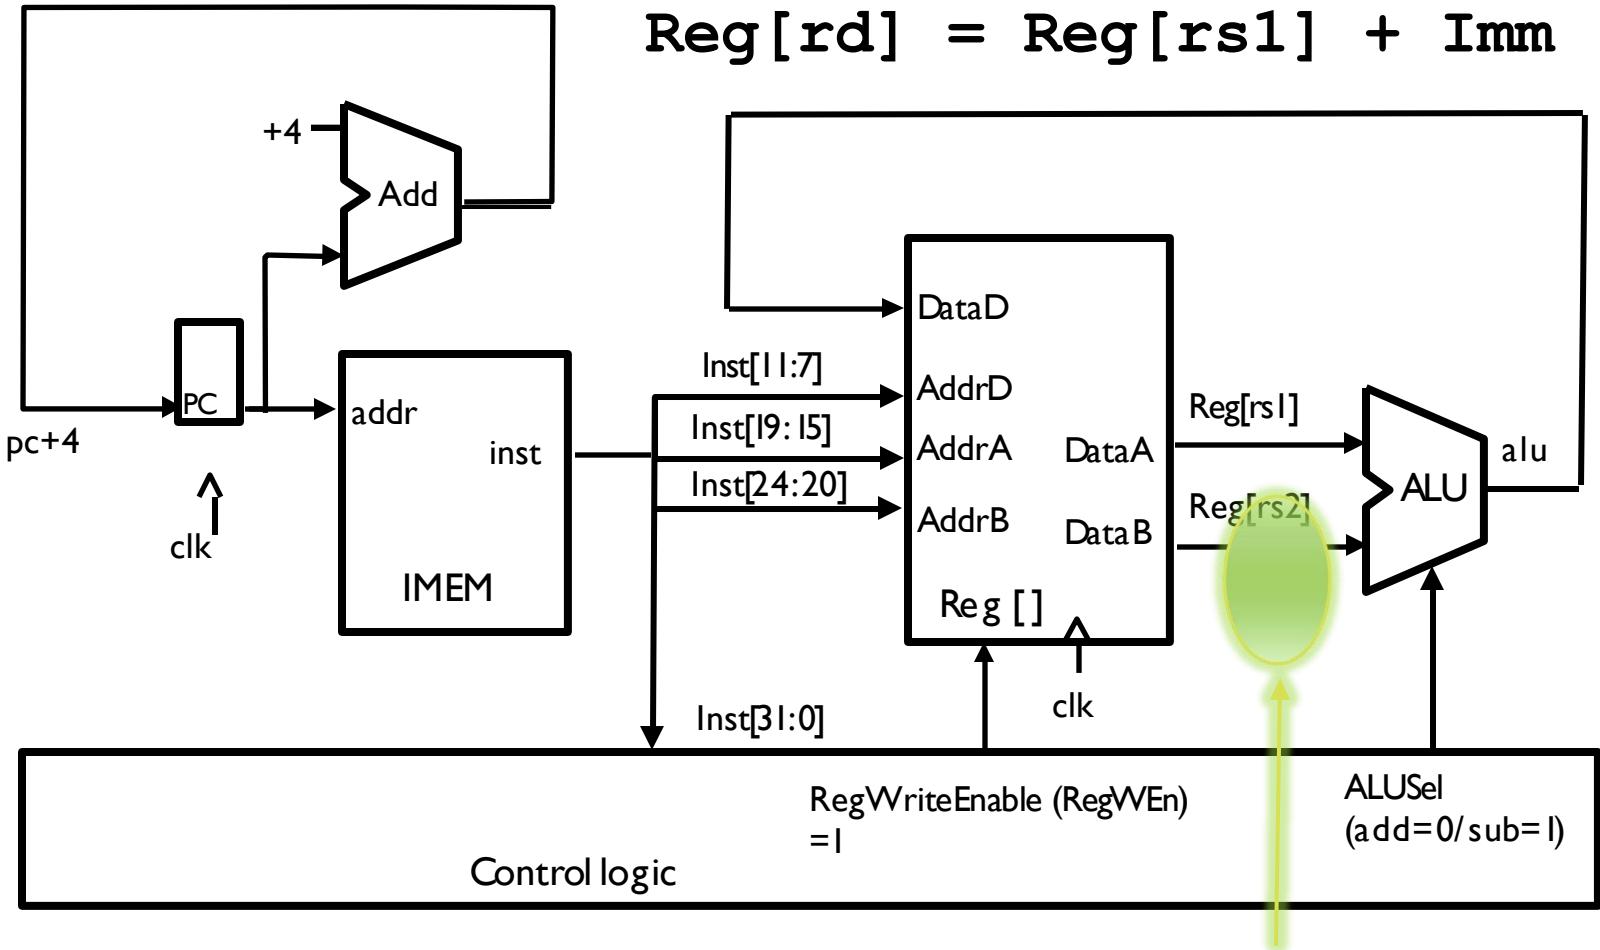
add

rd=15

OP-Imm

# Datapath for add/sub

$$\text{PC} = \text{PC} + 4$$

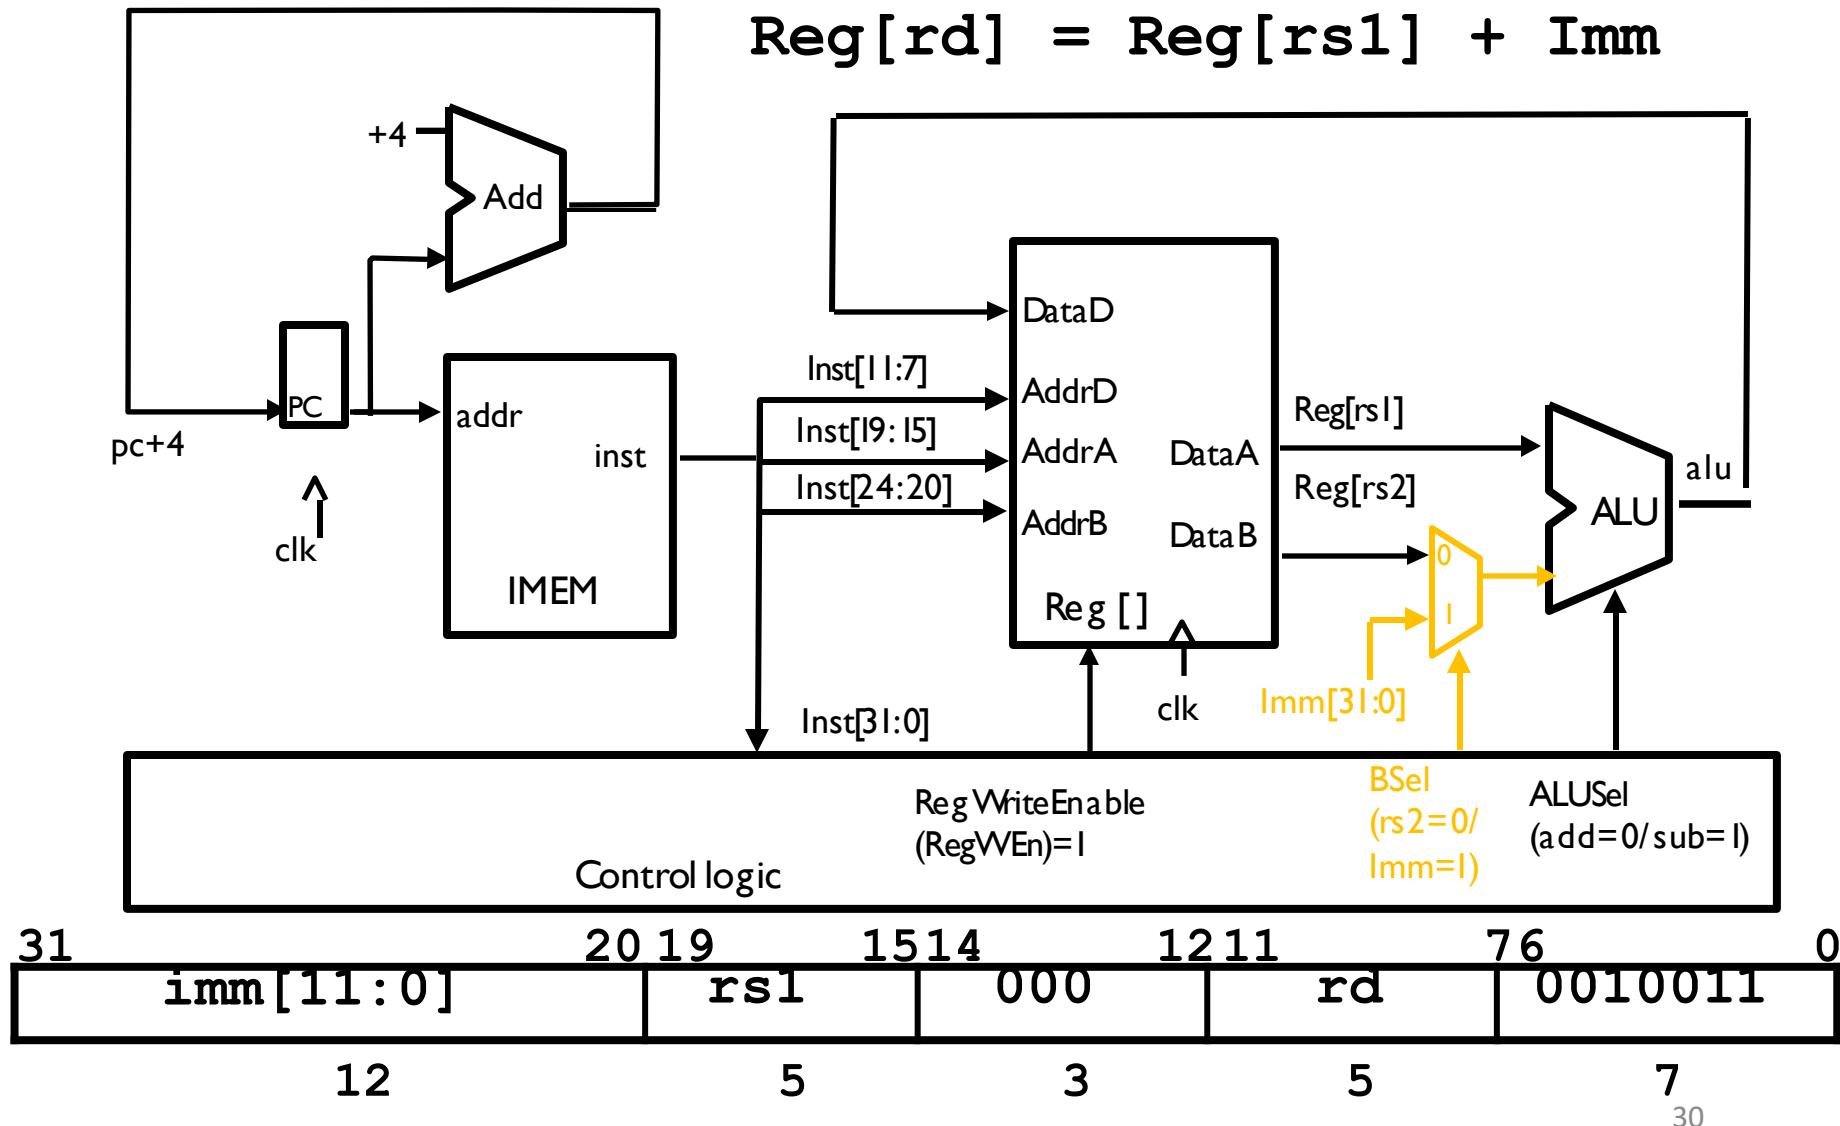


Immediate should  
be here

# Adding **addi** to Datapath

$$\text{PC} = \text{PC} + 4$$

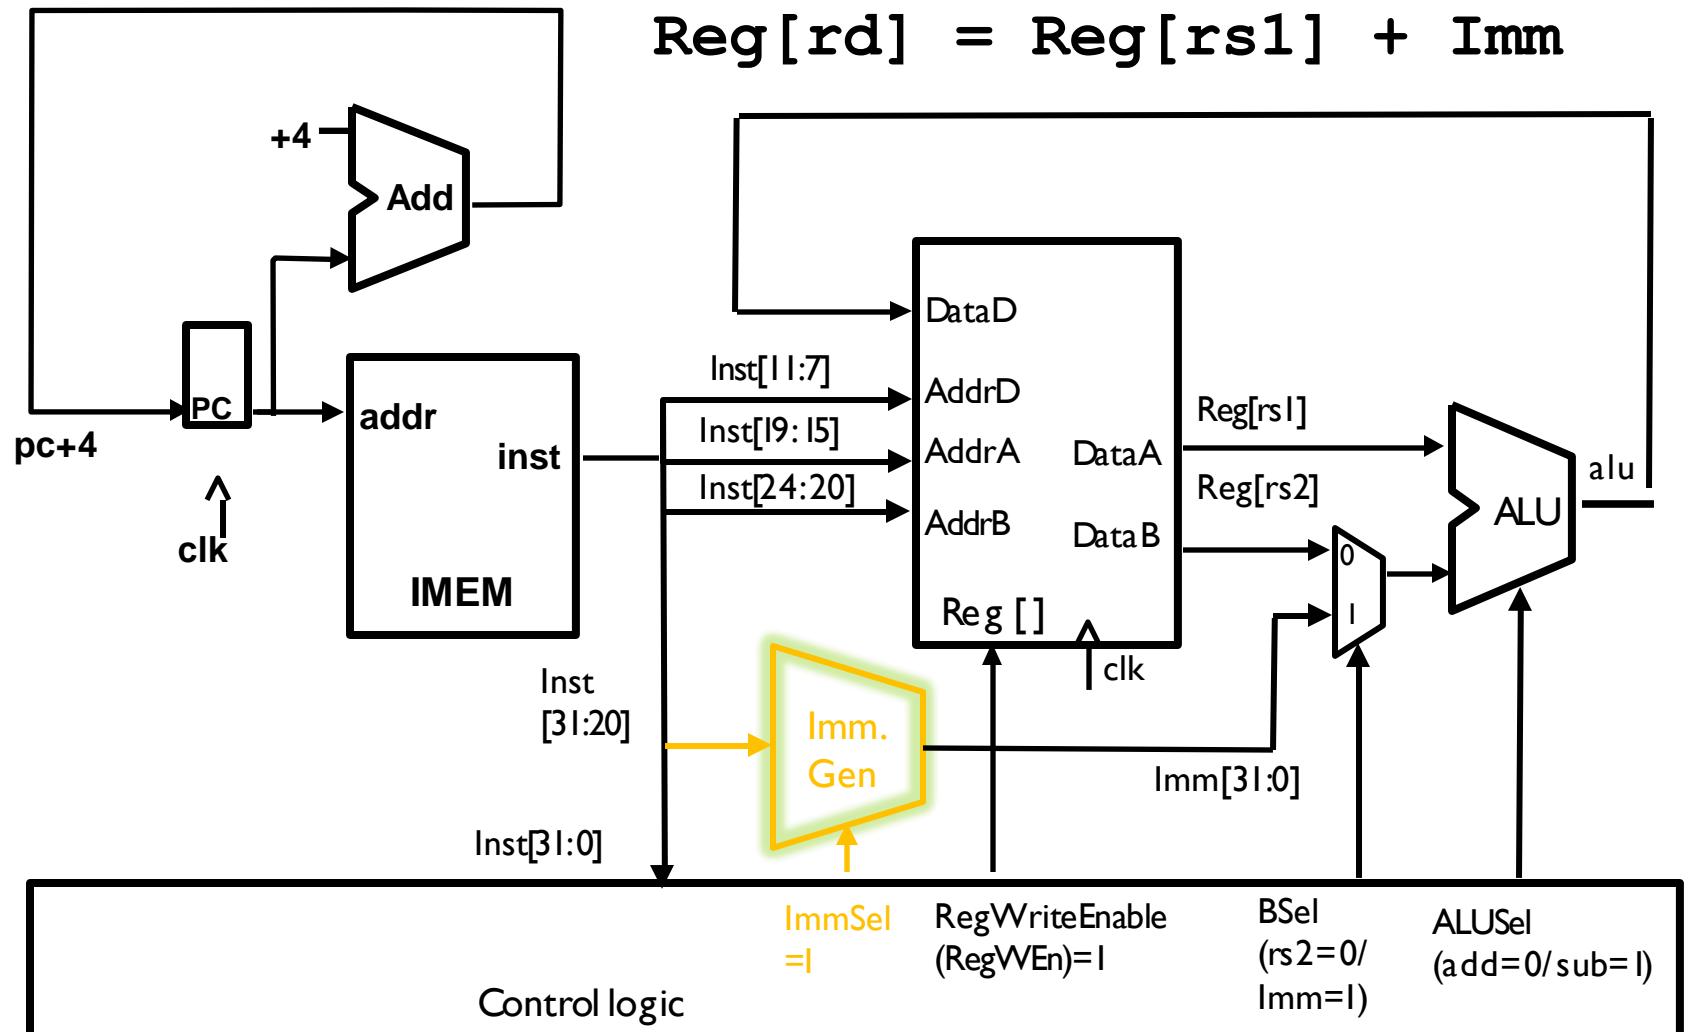
$$\text{Reg}[rd] = \text{Reg}[rs1] + \text{Imm}$$



# Adding **addi** to Datapath

$$PC = PC + 4$$

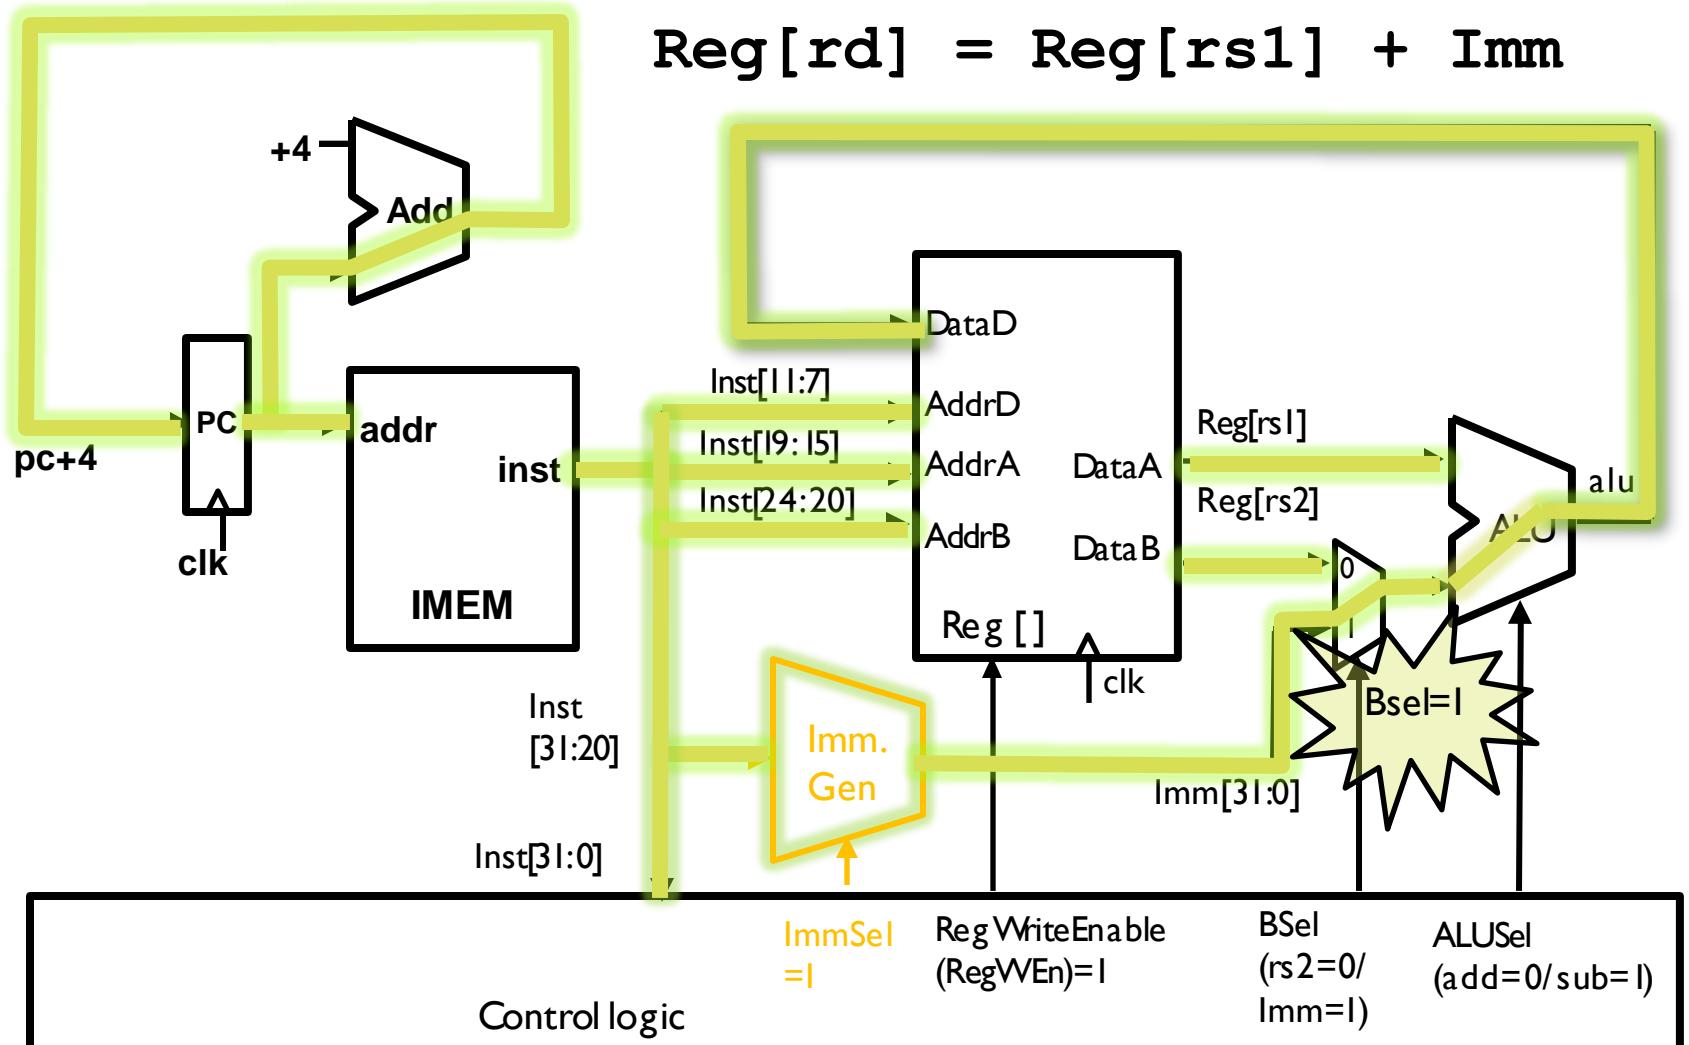
$$Reg[rd] = Reg[rs1] + Imm$$



# Adding **addi** to Datapath

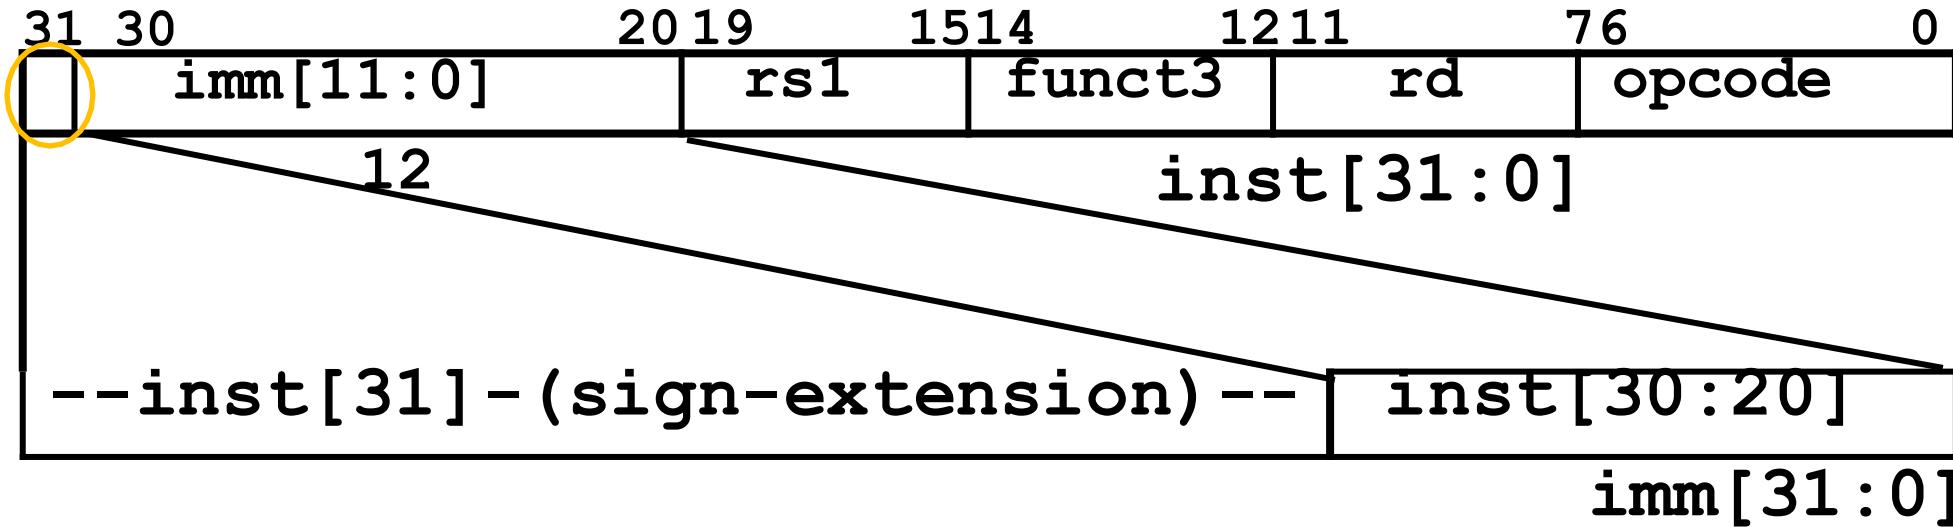
$$PC = PC + 4$$

$$Reg[rd] = Reg[rs1] + Imm$$



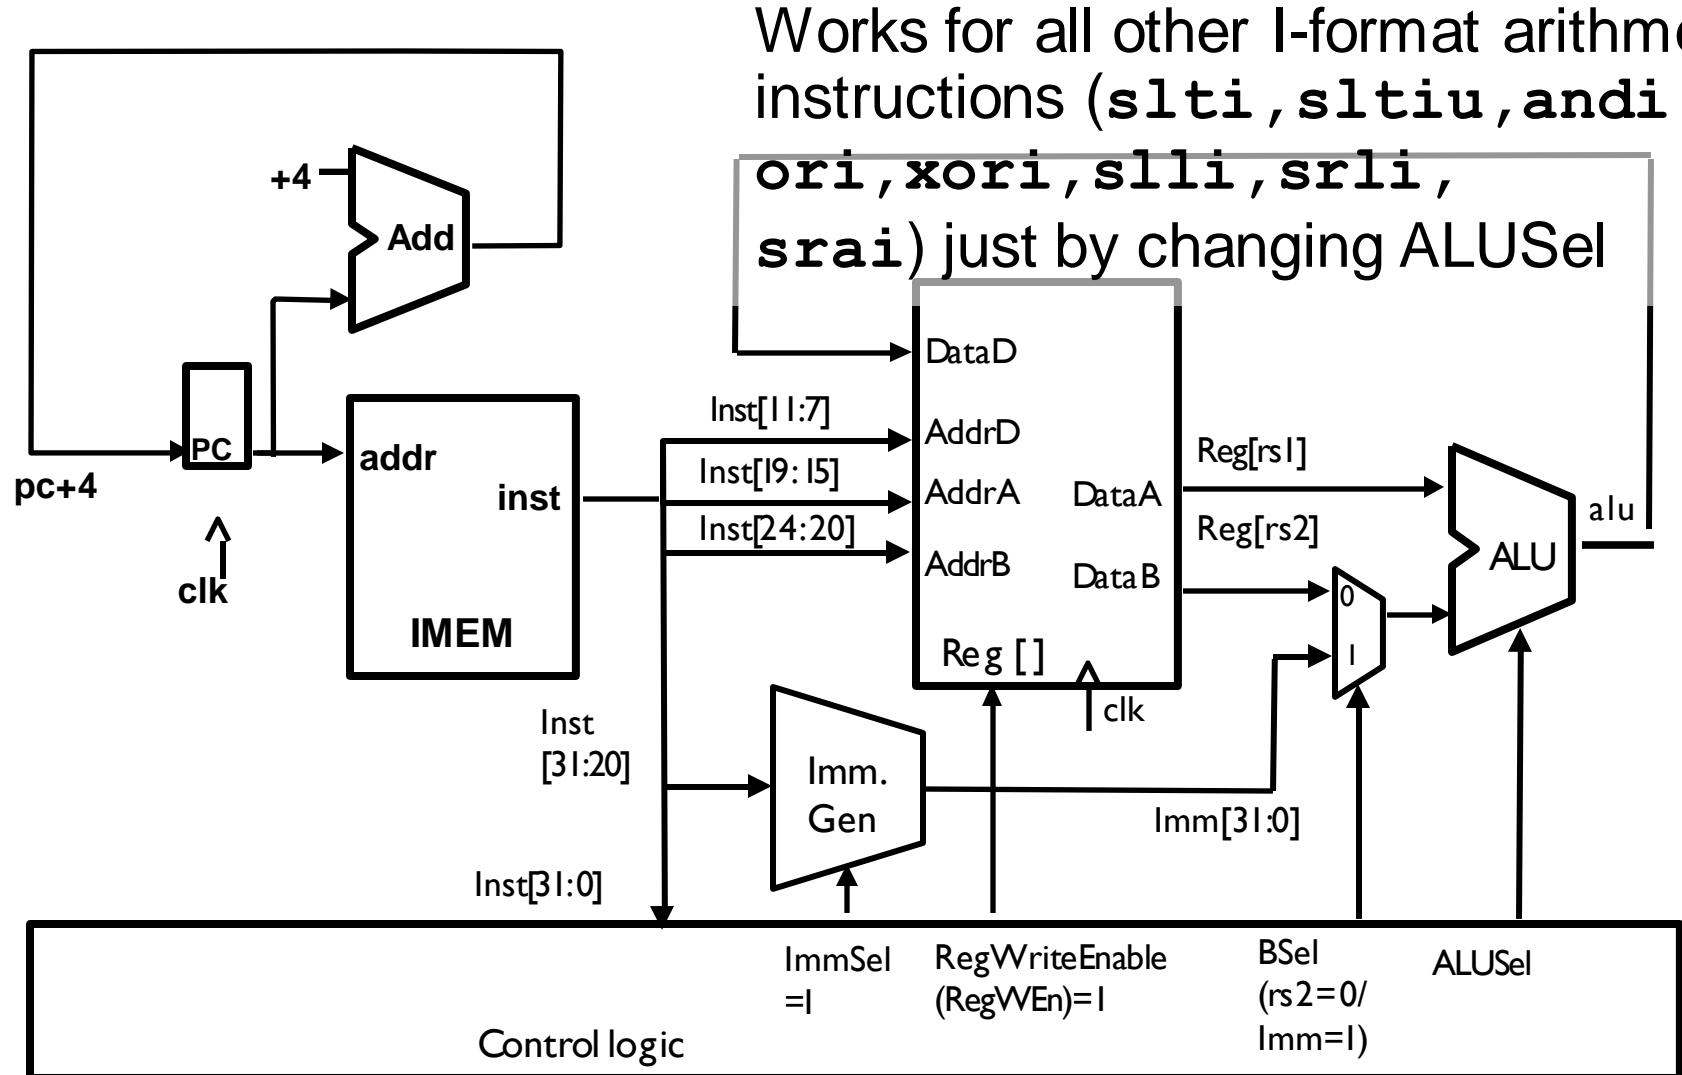
# I-Format Immediates

-inst[31]-



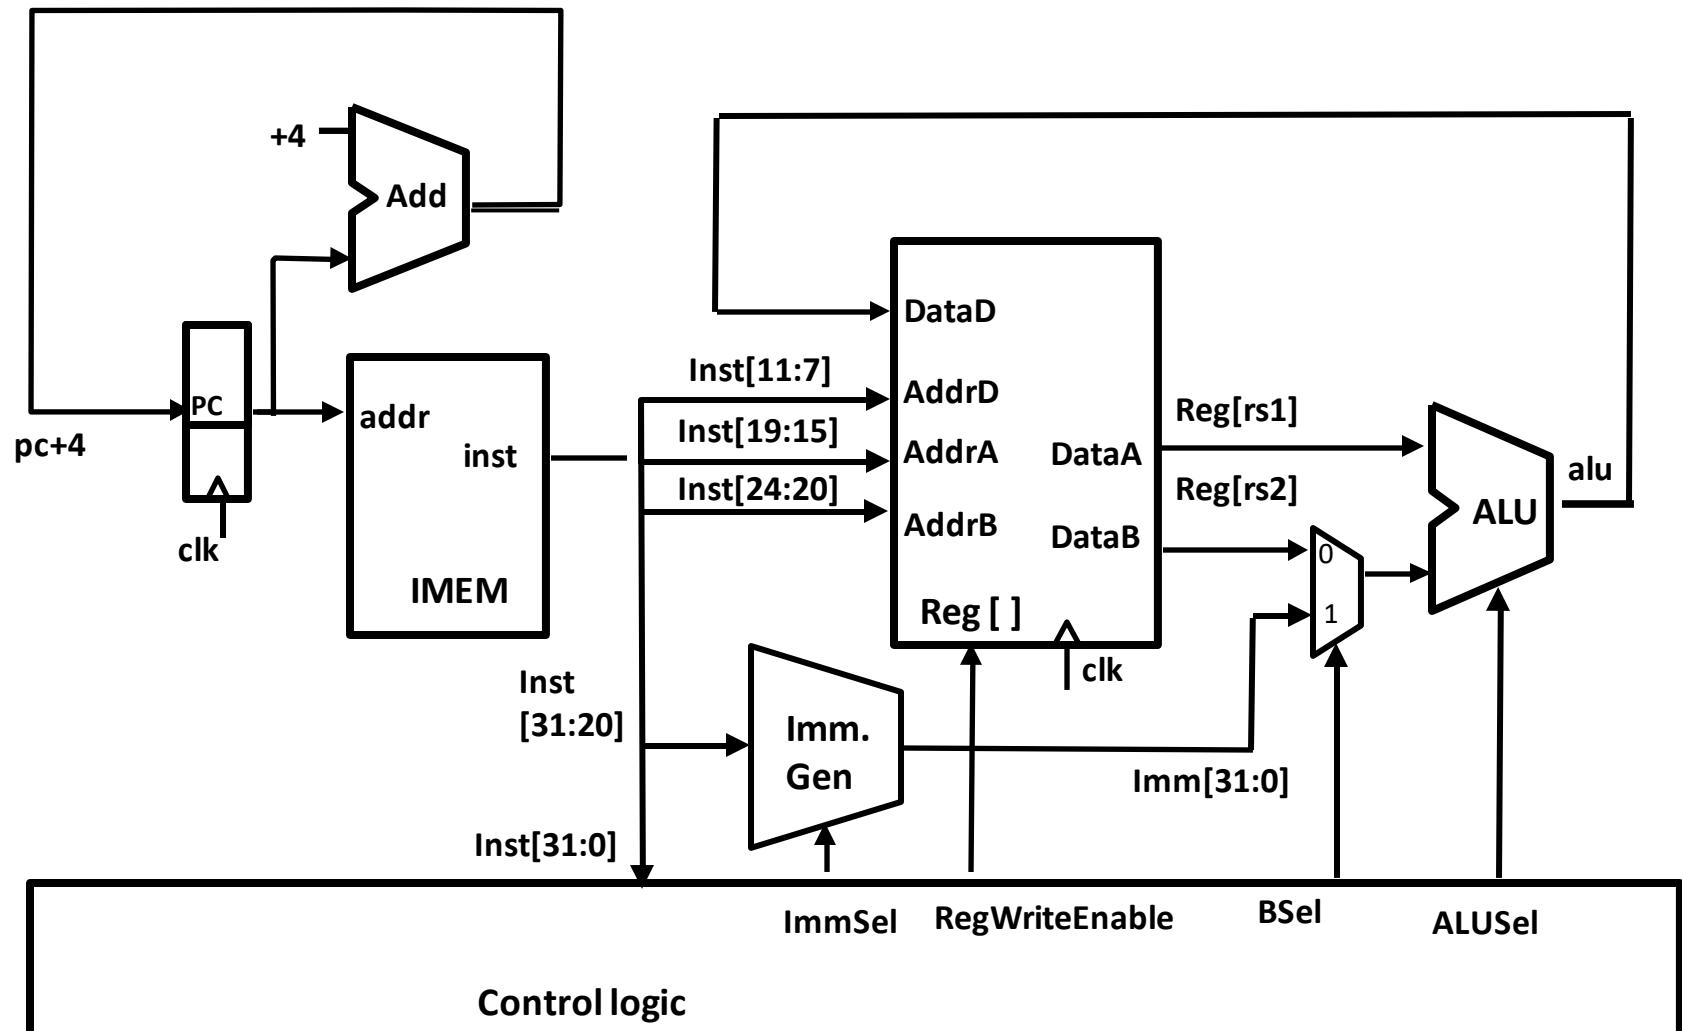
- High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])
- Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])

# Adding **addi** to Datapath



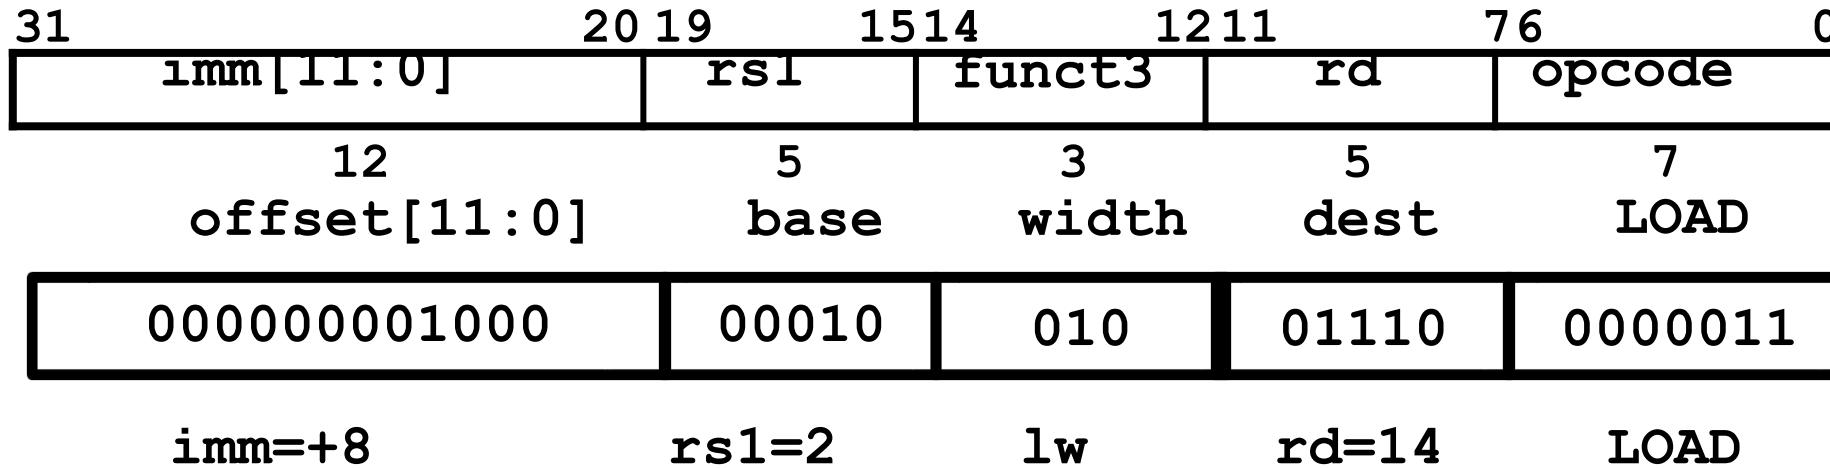
# **Supporting Loads**

# R+I Arithmetic/Logic Datapath



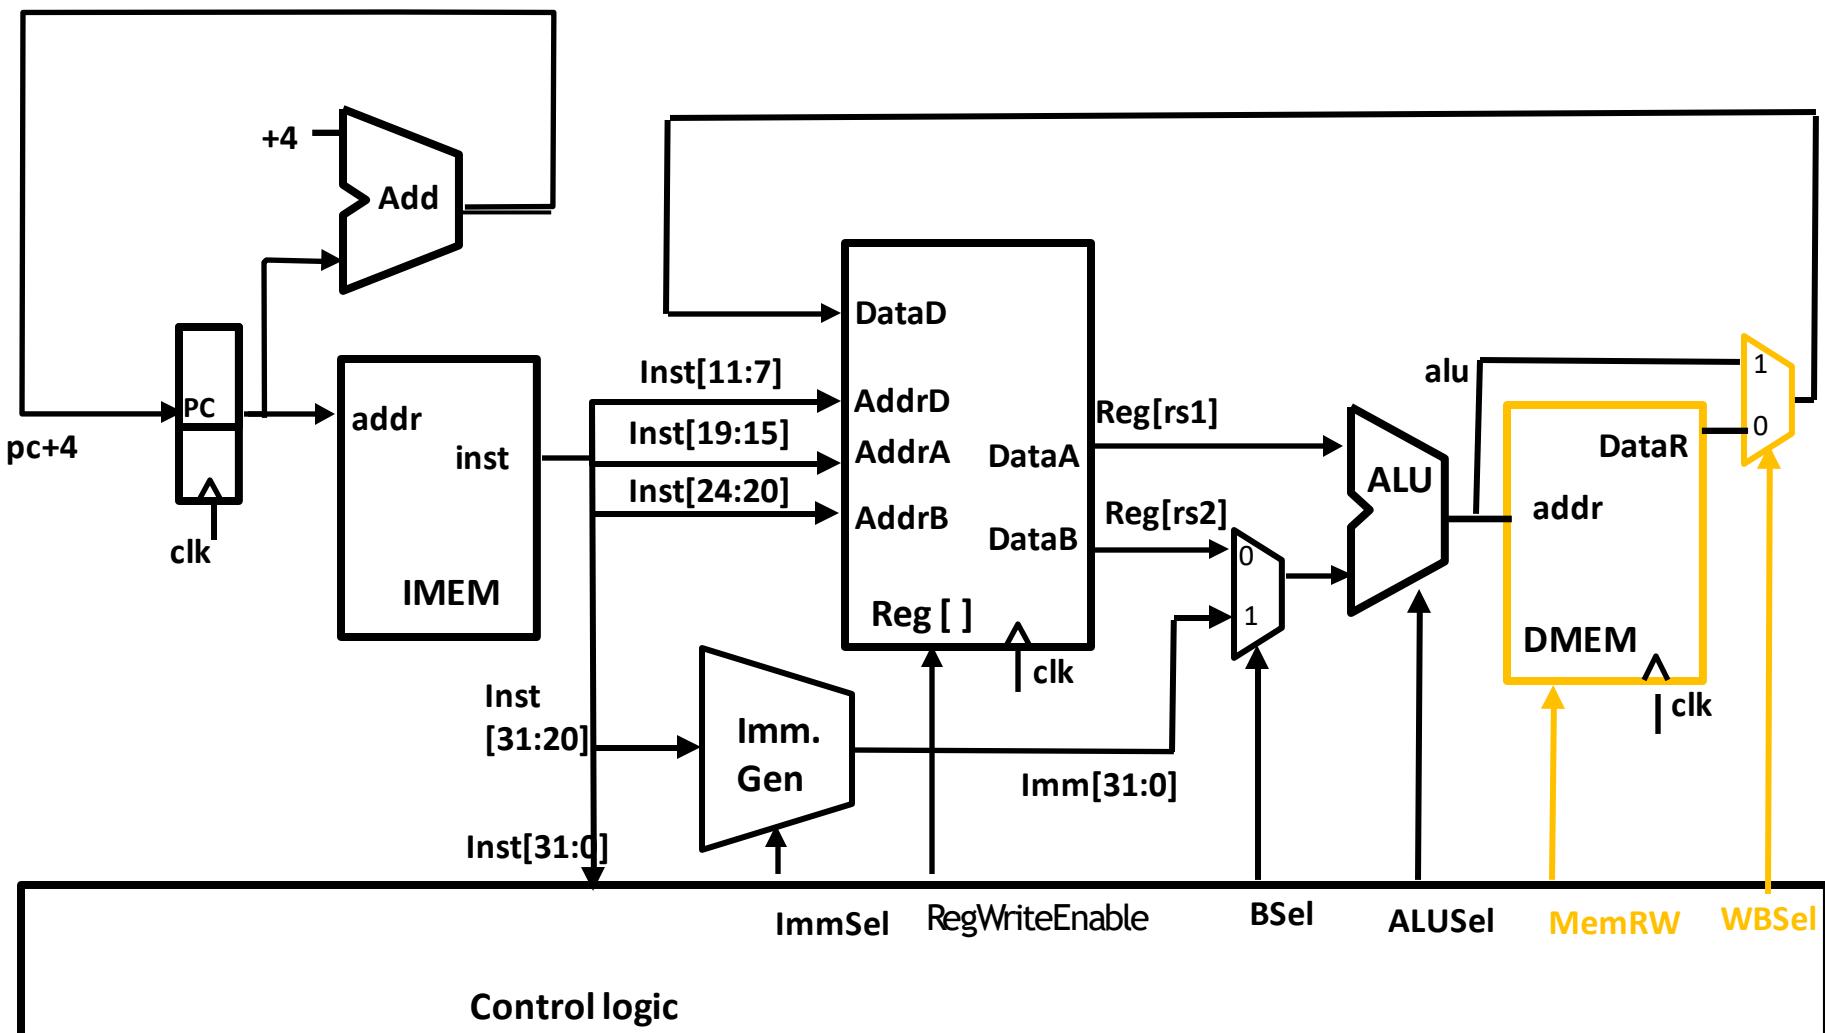
# Add **lw**

- RISC-V Assembly Instruction (I-type): **lw x14, 8(x2)**

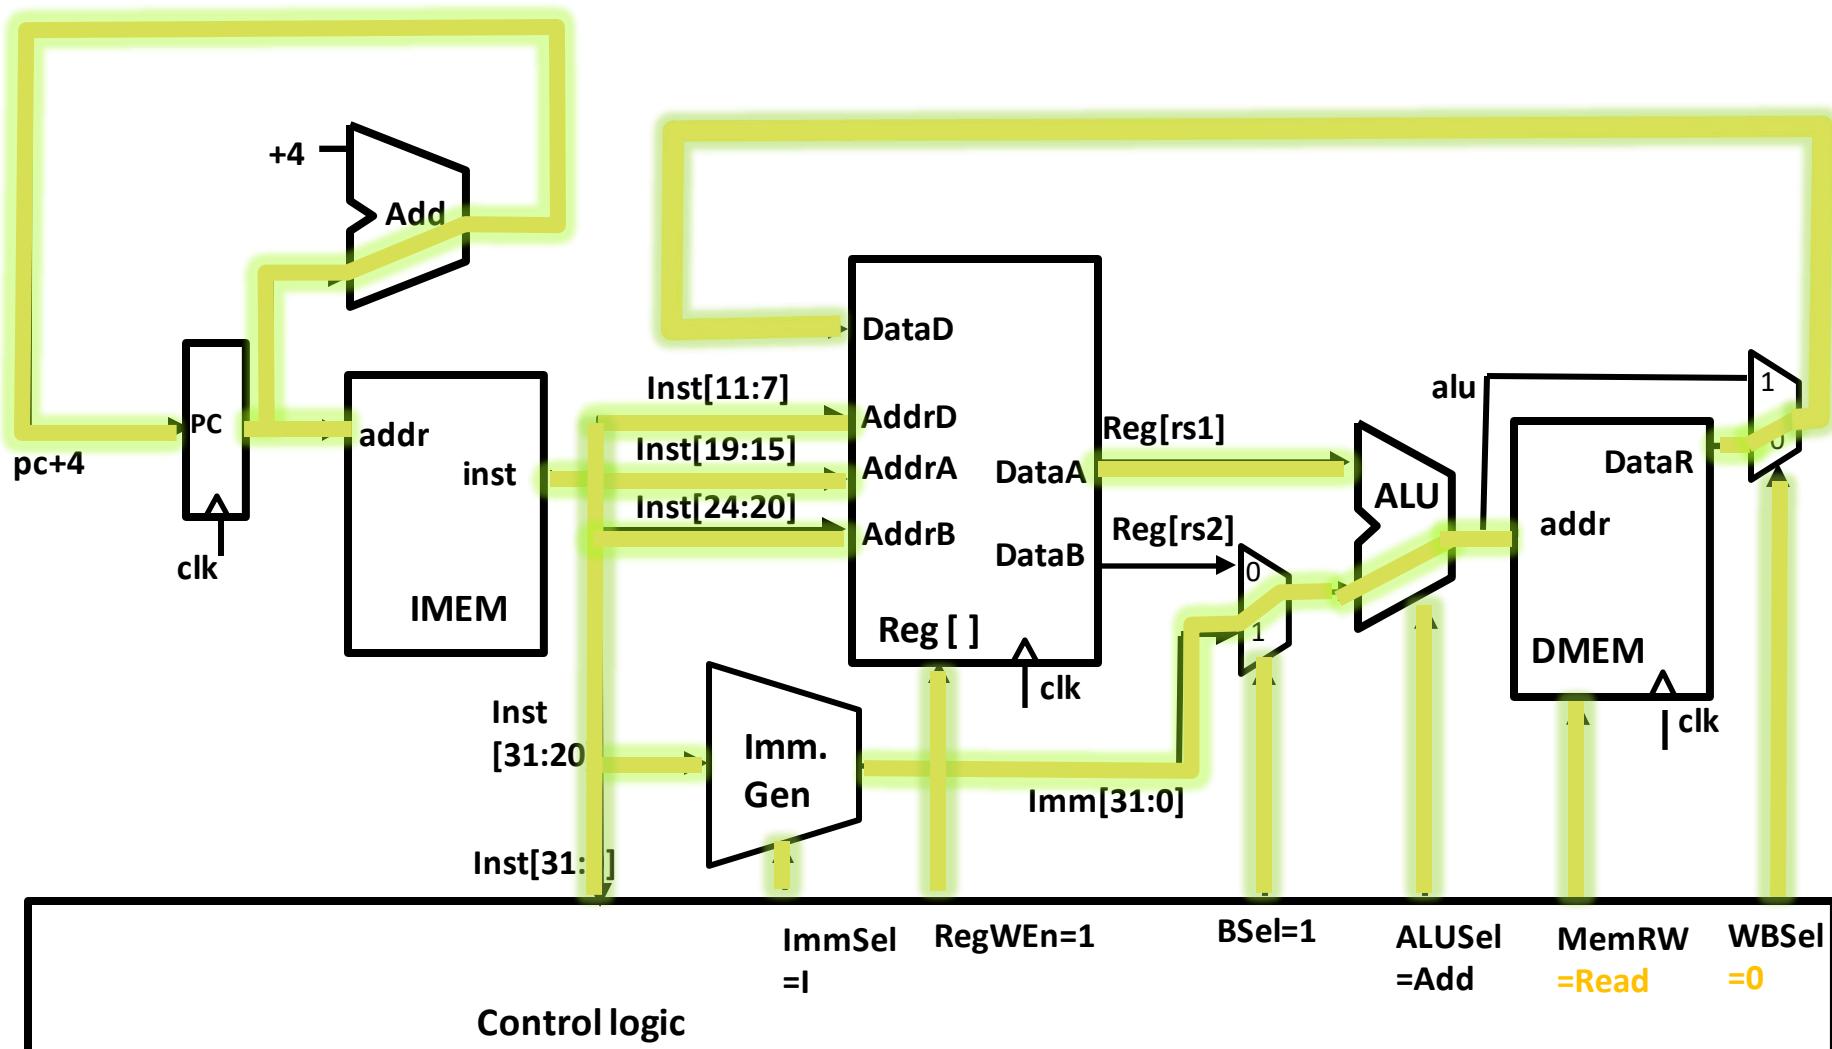


- The 12-bit signed immediate is added to the base address in register **rs1** to form the memory address
  - This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from memory is stored in register **rd**

# R+I Arithmetic/Logic Datapath



# R+I Arithmetic/Logic Datapath



# All RV32 Load Instructions

|           |     |     |    |         |     |
|-----------|-----|-----|----|---------|-----|
| imm[11:0] | rs1 | 000 | rd | 0000011 | lb  |
| imm[11:0] | rs1 | 001 | rd | 0000011 | lh  |
| imm[11:0] | rs1 | 010 | rd | 0000011 | lw  |
| imm[11:0] | rs1 | 100 | rd | 0000011 | lbu |
| imm[11:0] | rs1 | 101 | rd | 0000011 | lhu |

- funct3 field encodes size and ‘signedness’ of load data

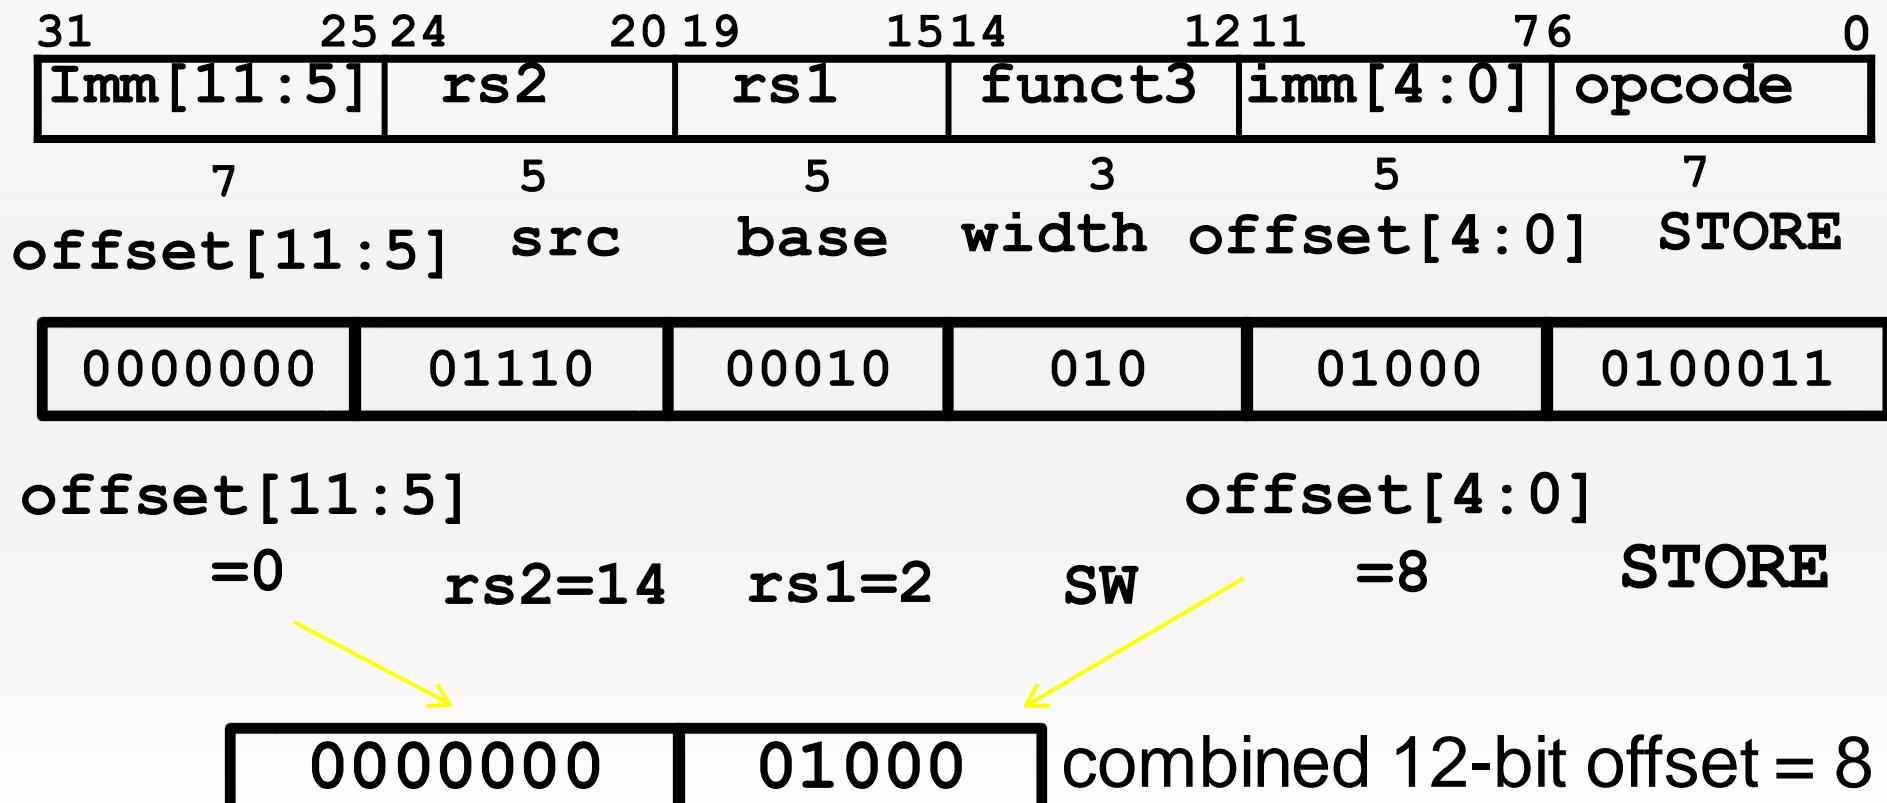
- Supporting the narrower loads requires additional logic to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to registerfile.
  - It is just a mux + a few gates

# **Datapath for Stores**

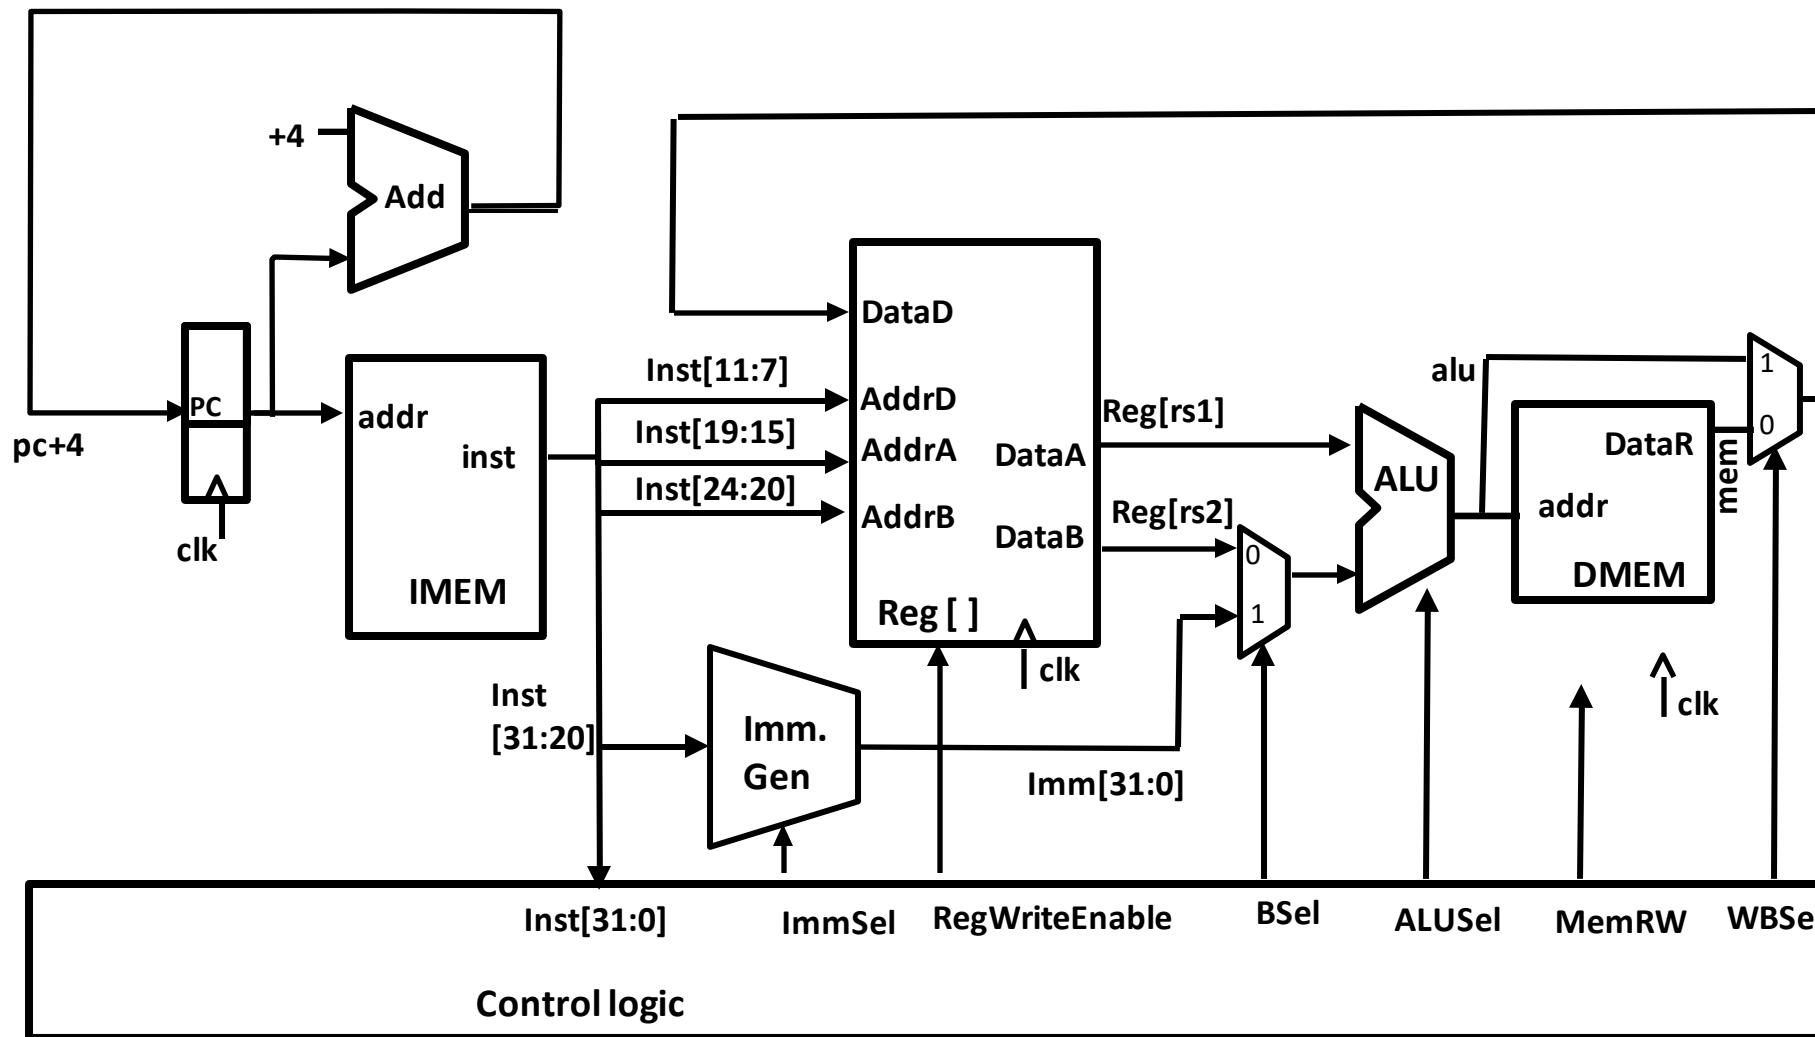
# Adding **sw** instruction

**sw**: Reads two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!

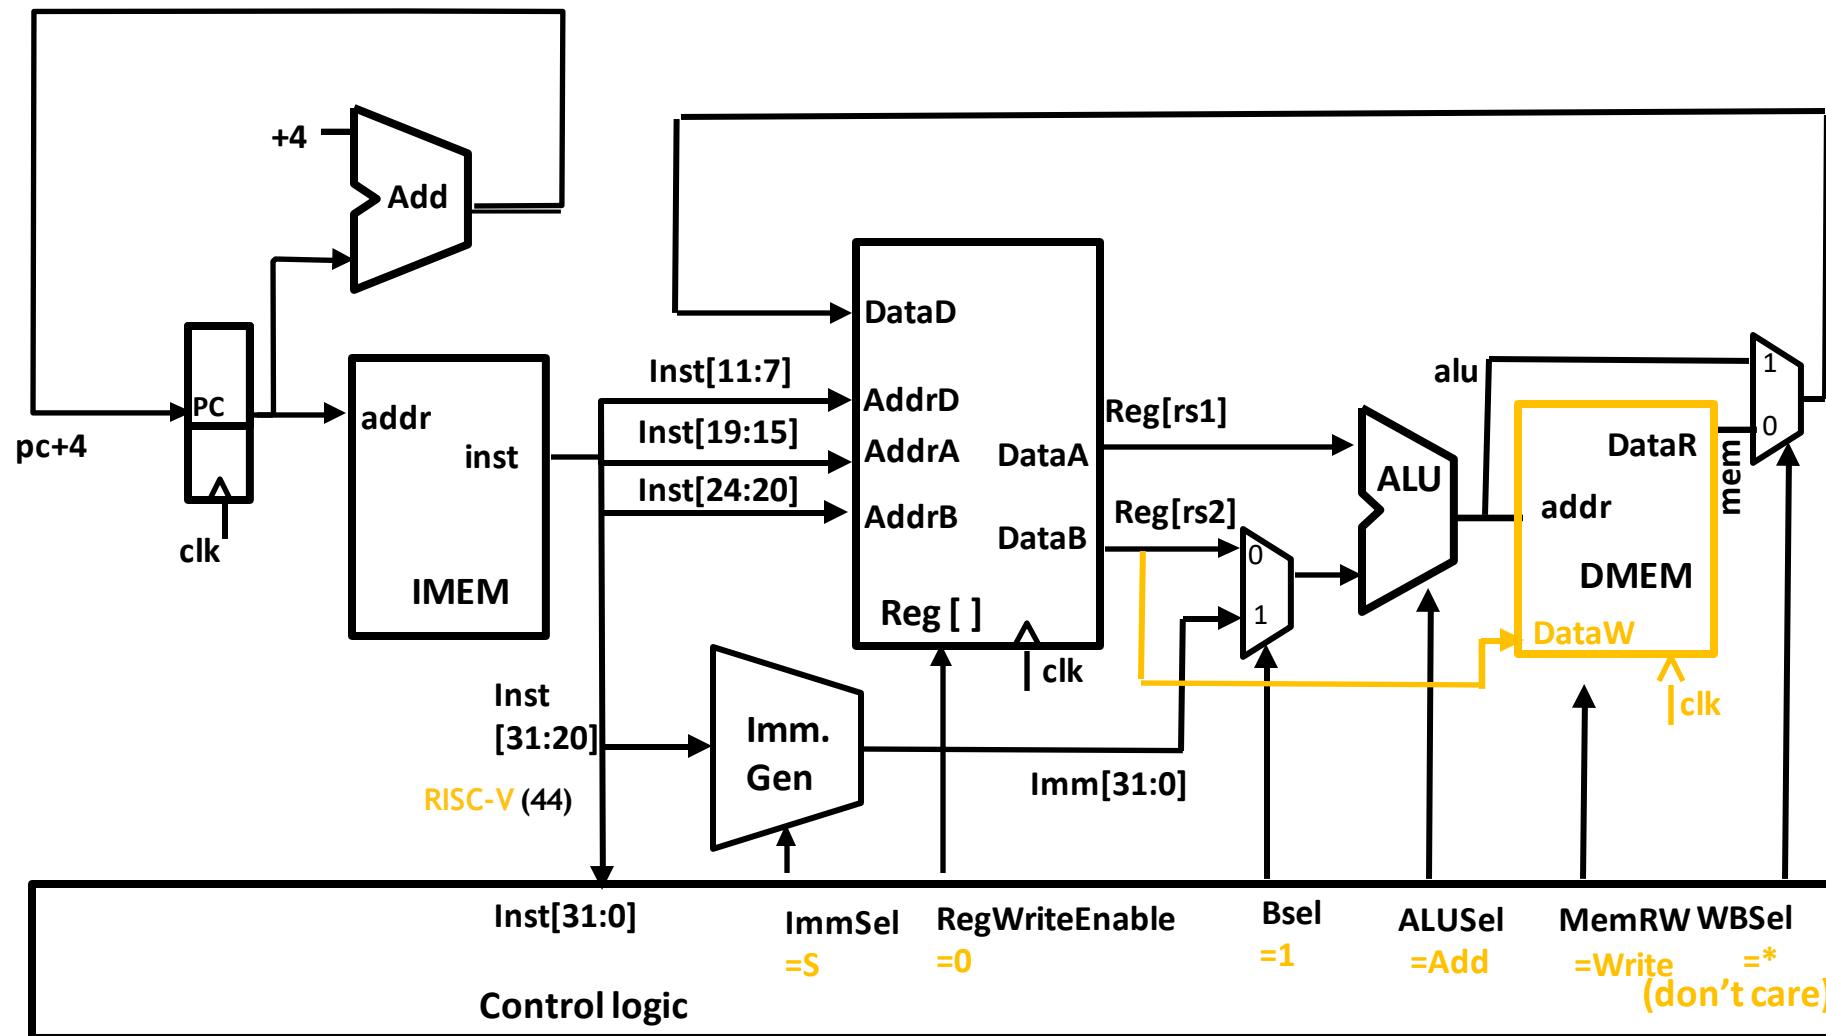
**SW x14, 8(x2)**



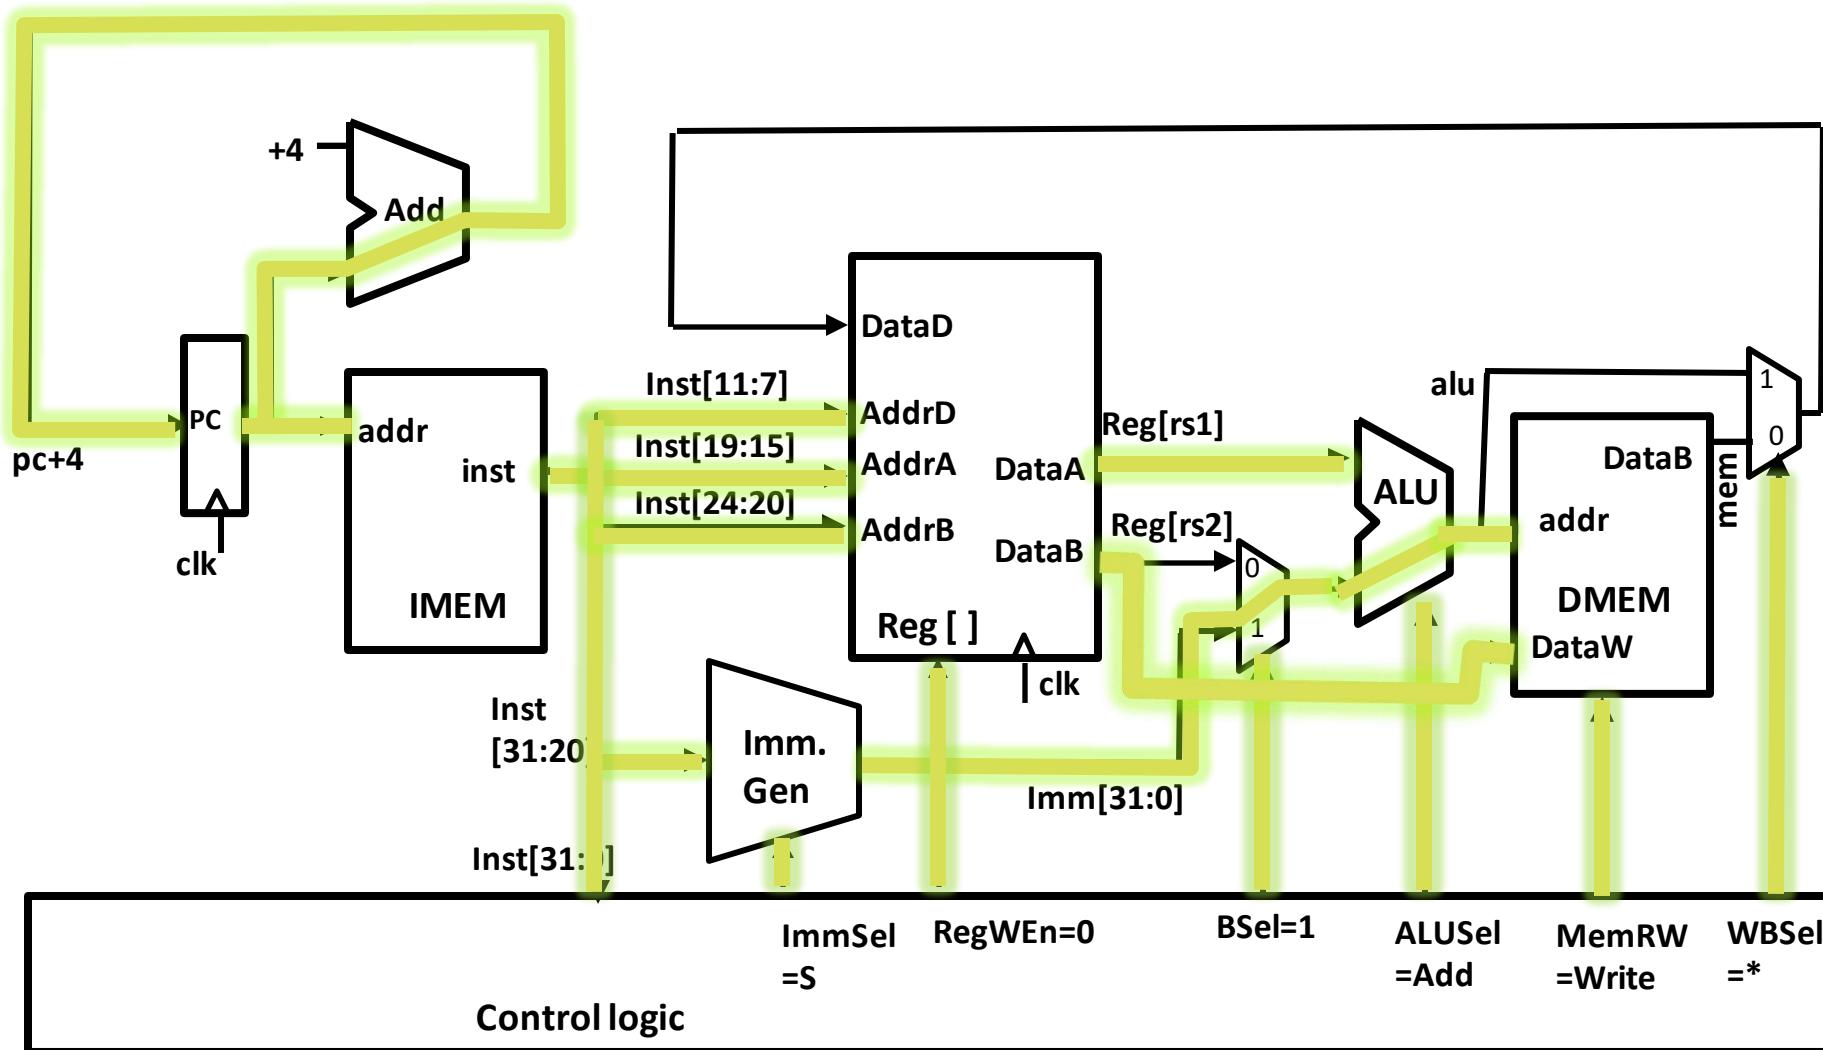
# Datapath with $lw$



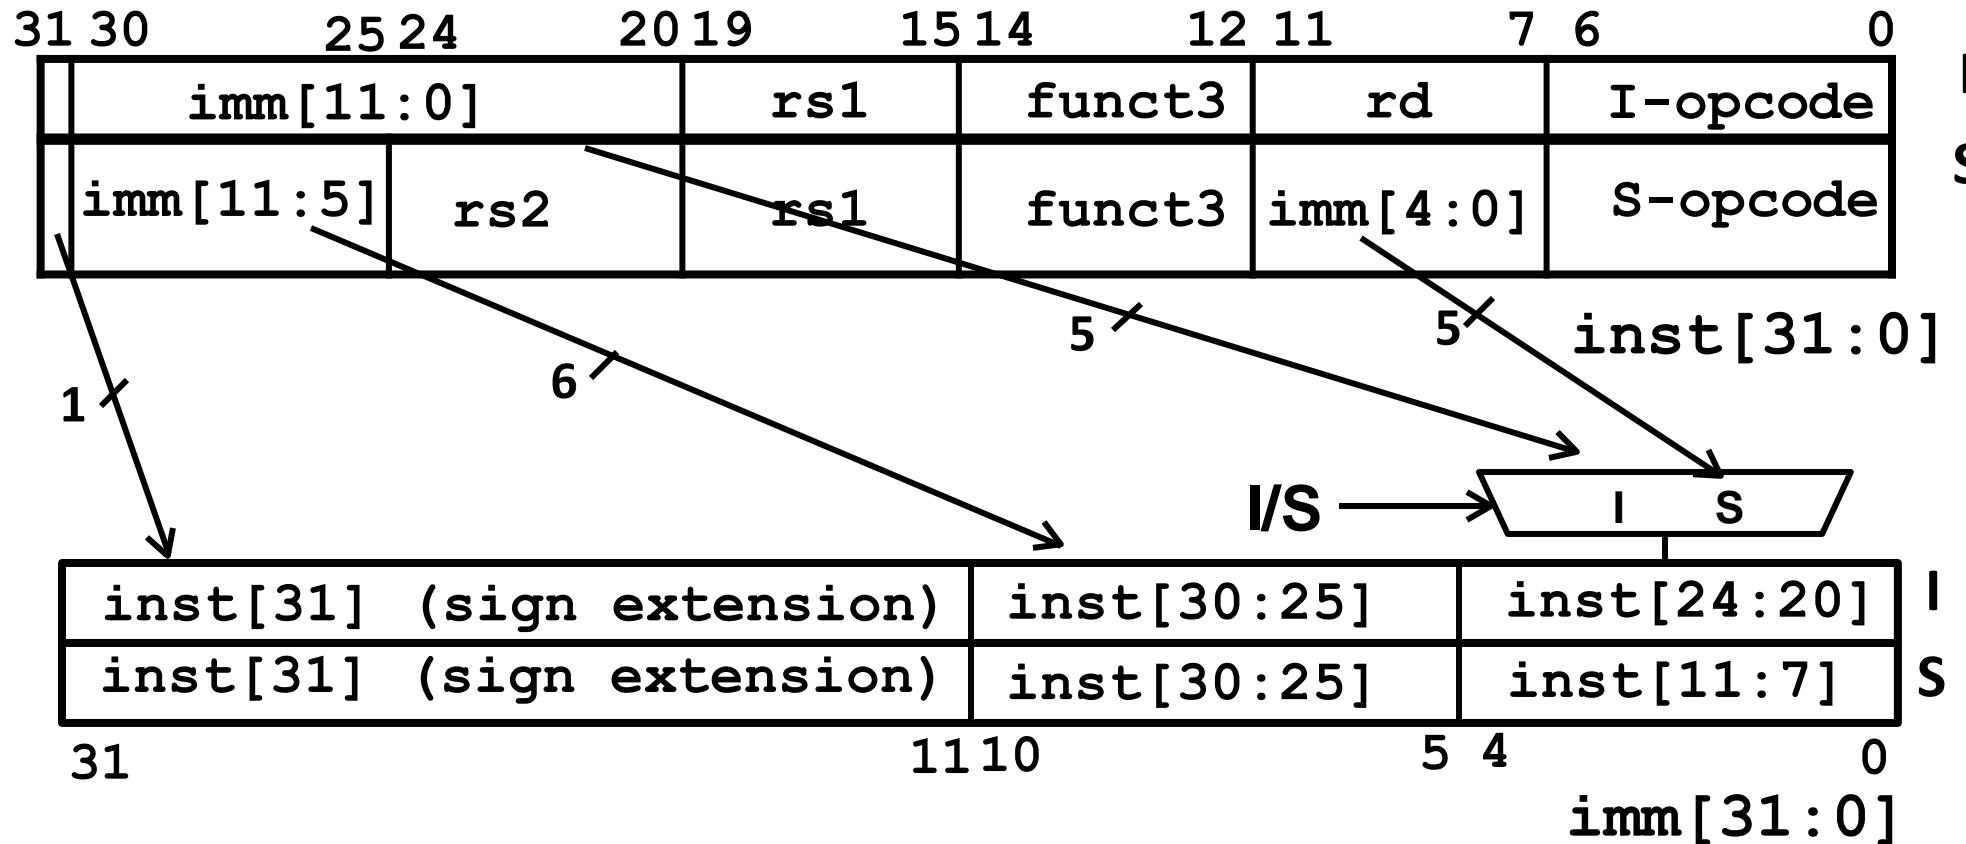
# Adding sw to Datapath



# Adding **sw** to Datapath



# I+S Immediate Generation



- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

# All RV32 Store Instructions

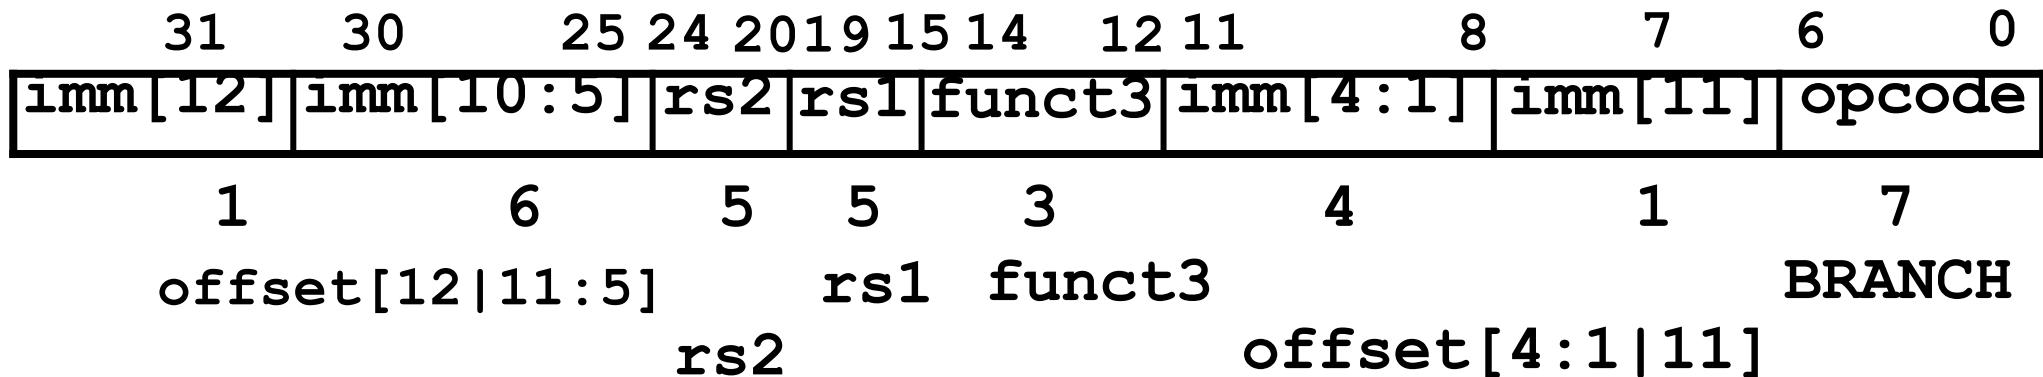
- Store byte writes the low byte to memory
- Store halfword writes the lower two bytes to memory

|           |     |     |     |          |         |           |
|-----------|-----|-----|-----|----------|---------|-----------|
| Imm[11:5] | rs2 | rs1 | 000 | imm[4:0] | 0100011 | <b>sb</b> |
| Imm[11:5] | rs2 | rs1 | 001 | imm[4:0] | 0100011 | <b>sh</b> |
| Imm[11:5] | rs2 | rs1 | 010 | imm[4:0] | 0100011 | <b>sw</b> |

width

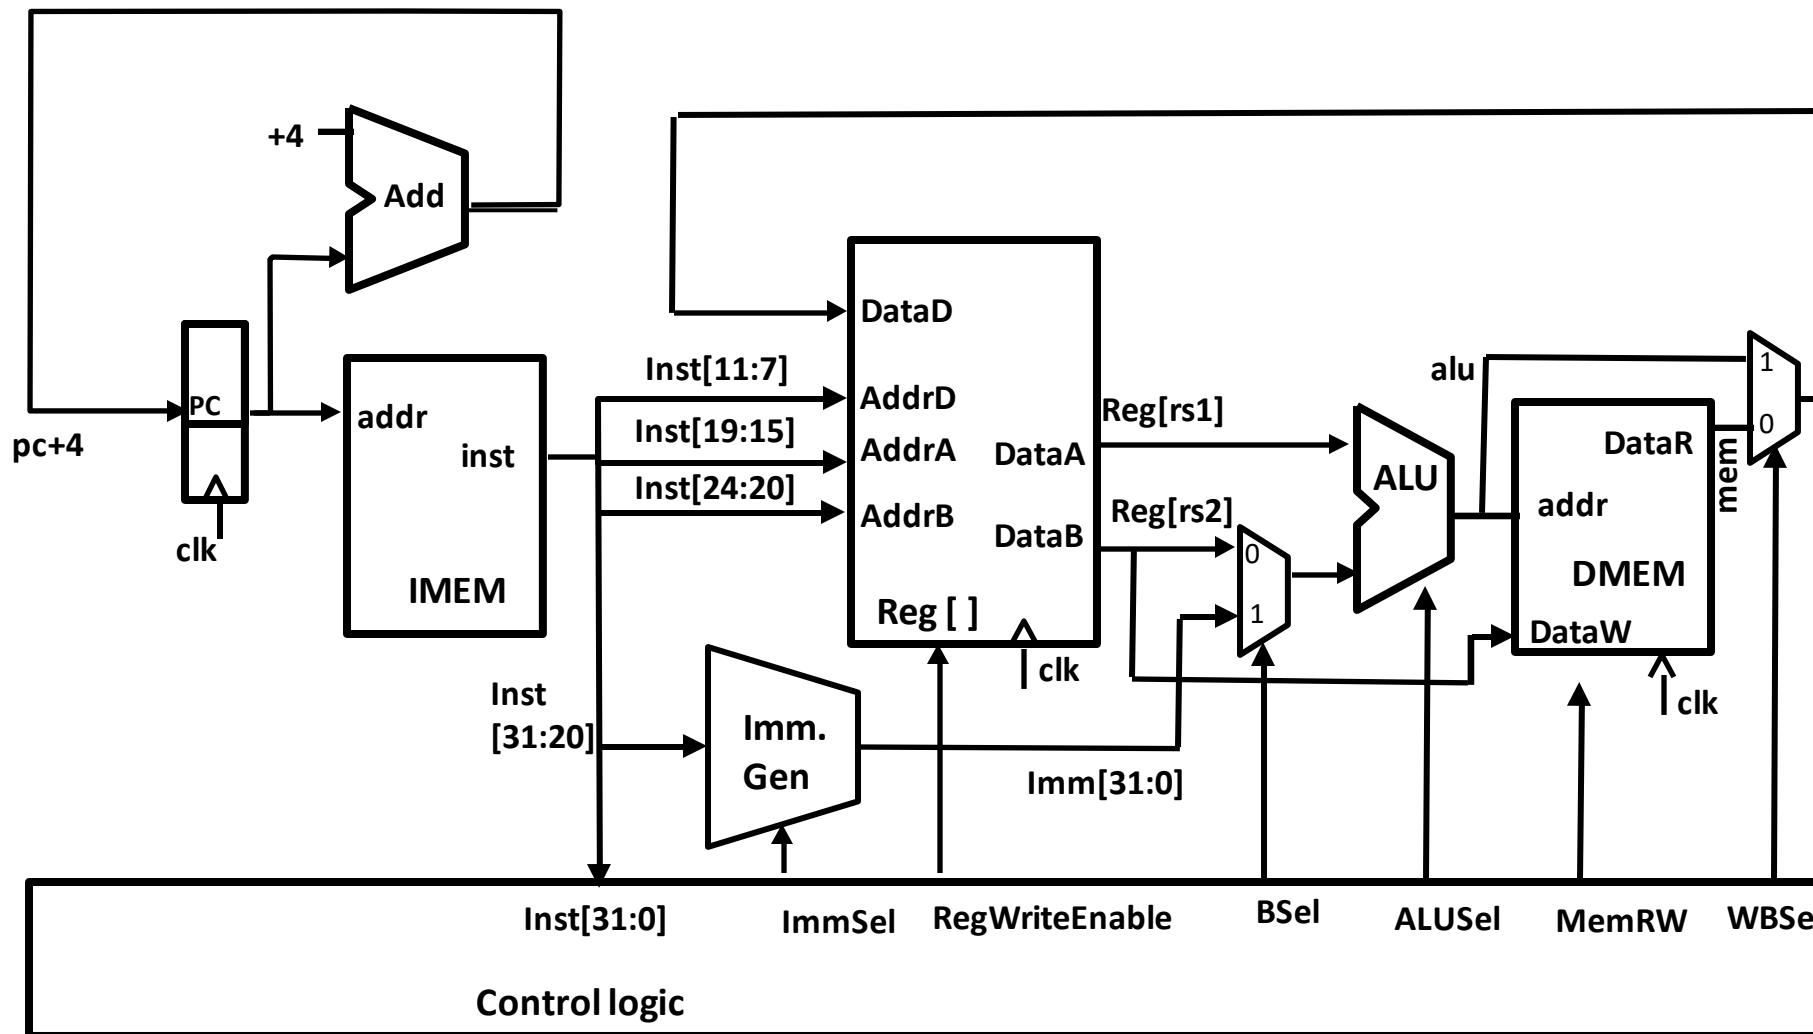
# **Implementing Branches**

# RISC-V B-Format for Branches



- B-format is mostly same as S-Format, with two register sources (**rs1/rs2**) and a 12-bit immediate **imm[12:1]**
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode **even** 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

# Datapath So Far



# To Add Branches

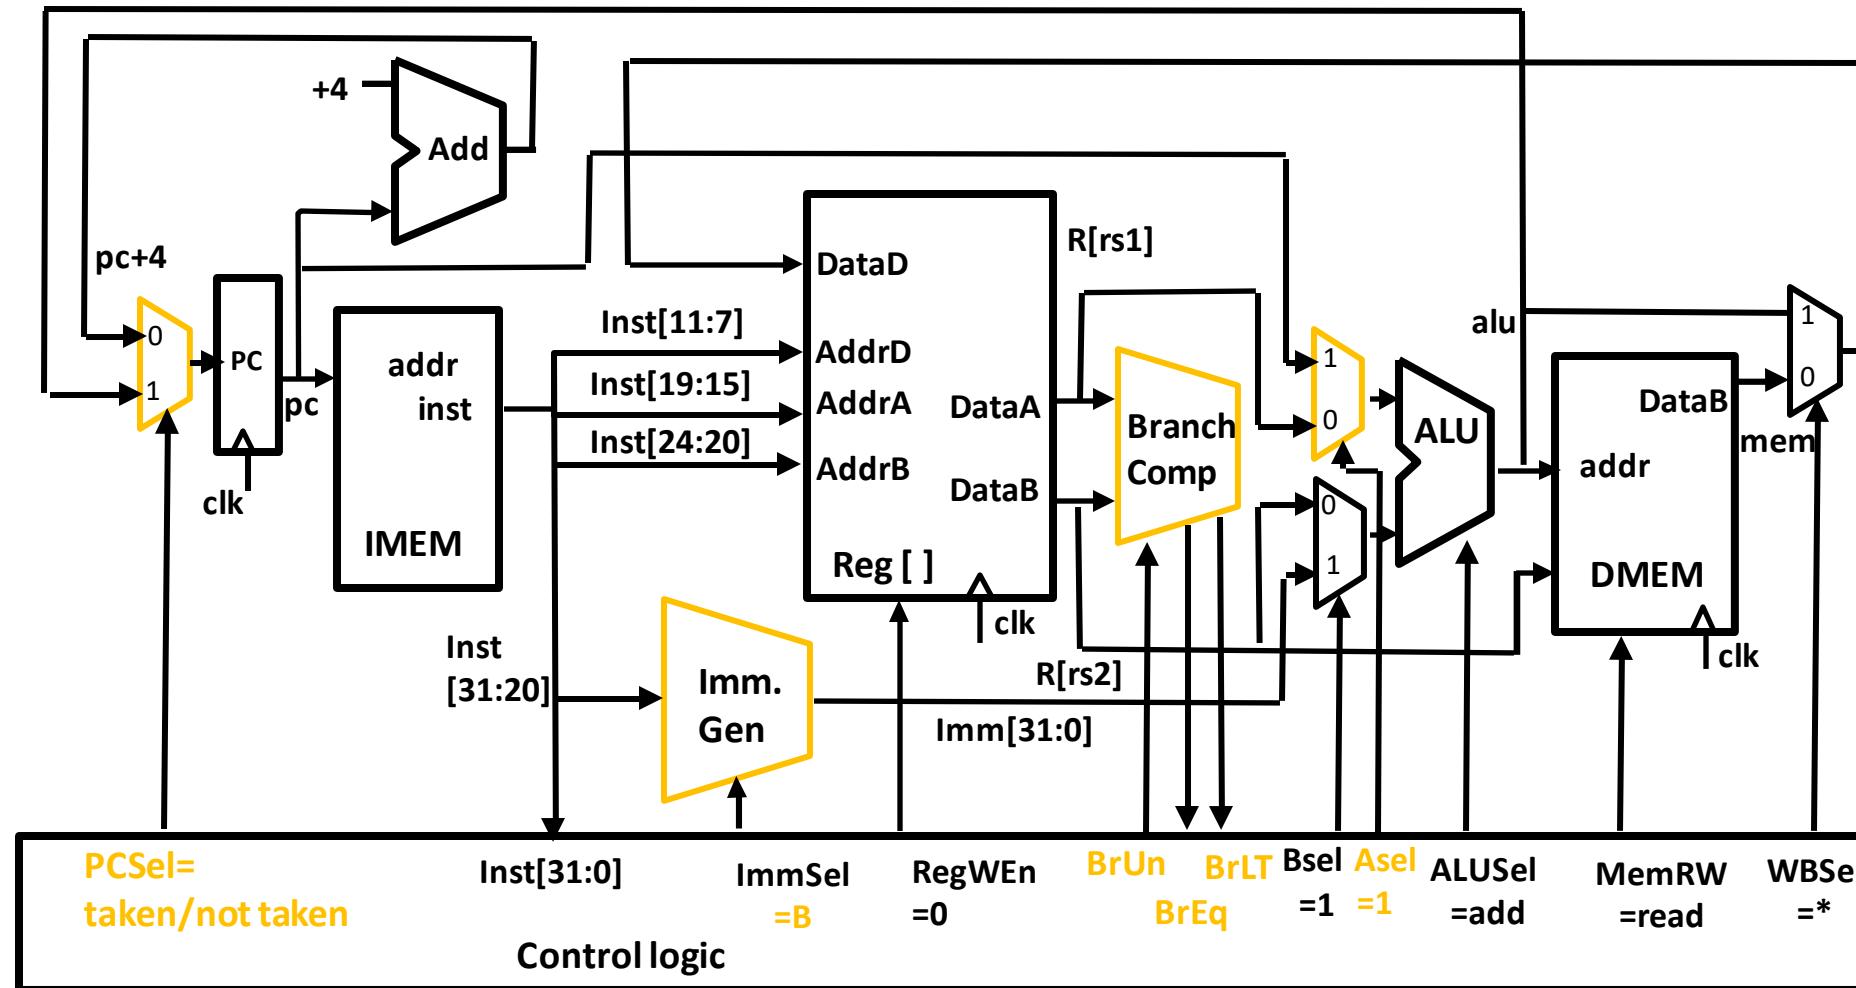
- Different change to the state:

$$\text{PC} = \begin{cases} \text{PC} + 4, & \text{branch not taken} \\ \text{PC} + \text{immediate}, & \text{branch taken} \end{cases}$$

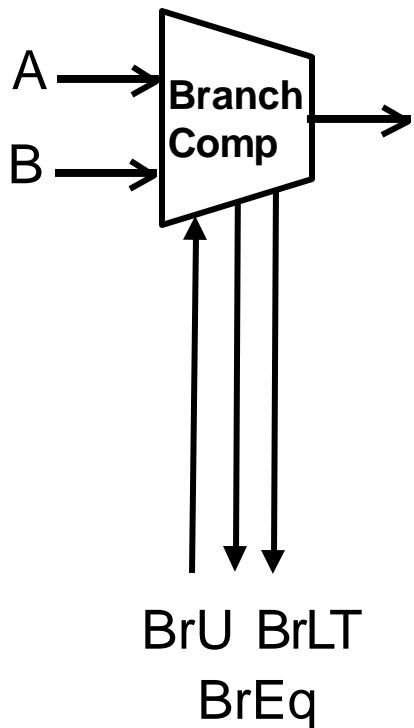
- Six branch instructions: **beq**, **bne**, **blt**, **bge**,  
**bltu**, **bgeu**
- Need to compute **PC + immediate** and to compare values of **rs1** and **rs2**
  - But have only one ALU – need more hardware

| 31              | 30        | 25 24 | 20 19 | 15 14   |  | 12 11          | 8       | 7 | 6      | 0 |
|-----------------|-----------|-------|-------|---------|--|----------------|---------|---|--------|---|
| imm[12]         | imm[10:5] | rs2   | rs1   | funct3  |  | imm[4:1]       | imm[11] |   | opcode |   |
| 1               | 6         | 5     | 5     | 3       |  | 4              | 1       |   | 7      |   |
| offset[12 10:5] |           | src2  | src1  | BEQ/BNE |  | offset[11 4:1] |         |   | BRANCH |   |
| offset[12 10:5] |           | src2  | src1  | BLT[U]  |  | offset[11 4:1] |         |   | BRANCH |   |
| offset[12 10:5] |           | src2  | src1  | BGE[U]  |  | offset[11 4:1] |         |   | BRANCH |   |

# Adding Branches



# Branch Comparator



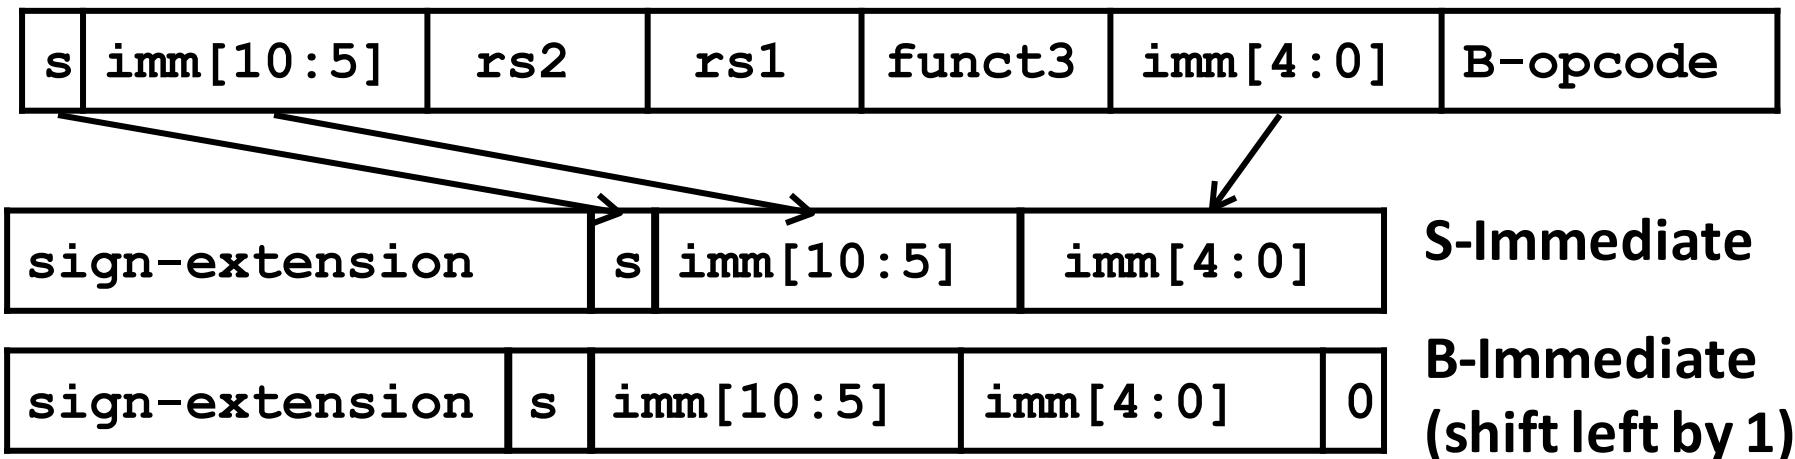
- **BrEq** = 1, if  $A=B$
- **BrLT** = 1, if  $A < B$
- **BrUn** = 1 selects unsigned comparison for **BrLT**, 0=signed

**BGE** branch:  $A \geq B$ , if  $\overline{A < B}$

$$\overline{A < B} = !(A < B)$$

# Branch Immediates (In Other ISAs)

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- Standard approach: Treat immediate as in range -2048..+2047, then shift left by 1 bit to multiply by 2 for branches



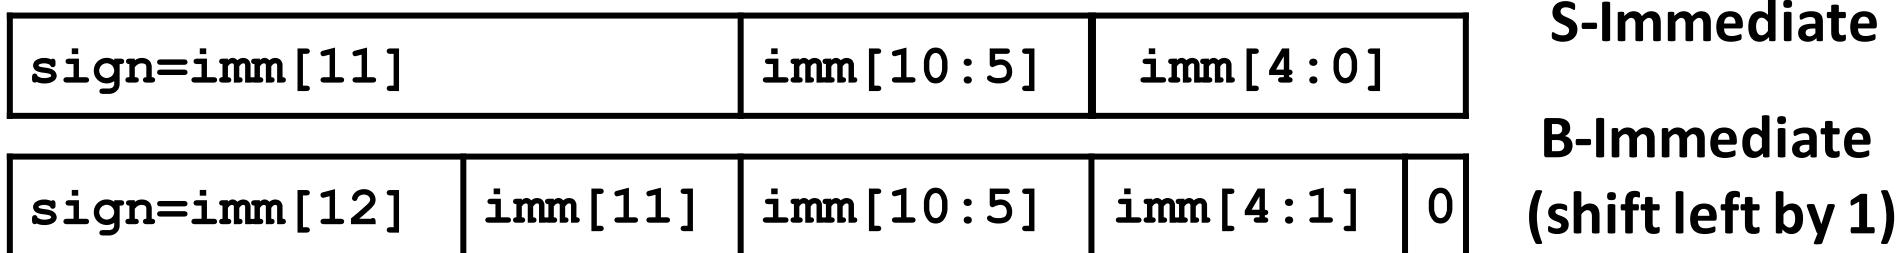
Each instruction immediate bit can appear in one of two places in output immediate value – so need one 2-way mux per bit

# Branch Immediates (In RISC-V ISAs)

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes

**NOTE :** Page. 116 (RISC V Textbook -> get extra information)

- RISC-V approach: keep 11 immediate bits in fixed position in output value, **and rotate LSB of S-format to be bit 12 of B-format**



Only one bit changes position between S and B, so only need a single-bit 2-way mux

The RISC-V architects wanted to support the possibility of instructions that are only 2 bytes long, so the branch instructions represent the number of *halfwords* between the branch and the branch target

# RISC-V Immediate Encoding

## Instruction encodings, $\text{inst}[31:0]$

| 31 | 30           | 25 24 | 20 19 | 15 14  | 12 11       | 8 7 6  | 0 |        |
|----|--------------|-------|-------|--------|-------------|--------|---|--------|
|    | imm[11:0]    |       | rs1   | funct3 | rd          | opcode |   | I-type |
|    | imm[11:5]    | rs2   | rs1   | funct3 | imm[4:0]    | opcode |   | S-type |
|    | imm[12 10:5] | rs2   | rs1   | funct3 | imm[4:1 11] | opcode |   | B-type |

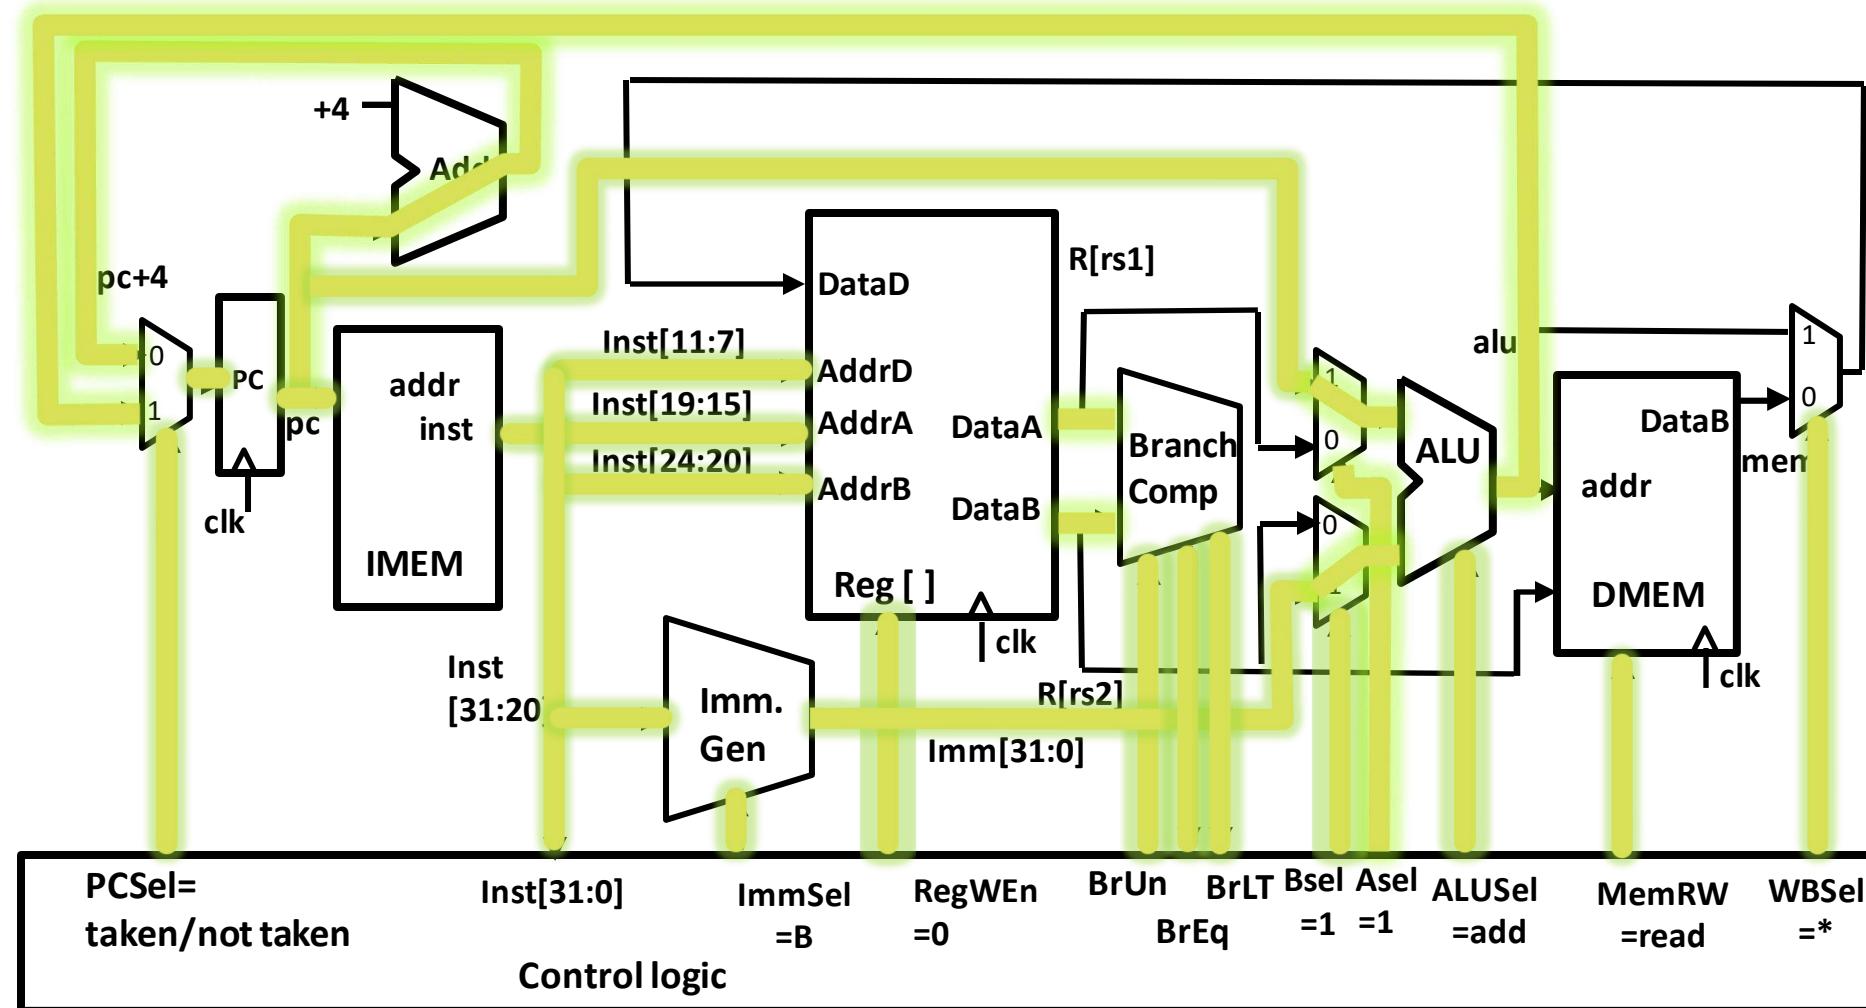
## 32-bit immediates produced, $\text{imm}[31:0]$

| 31 | 25 24      | 12 11   | 10          | 5 4         | 1        | 0 |        |
|----|------------|---------|-------------|-------------|----------|---|--------|
| -  | -inst[31]- |         | inst[30:25] | inst[24:21] | inst[20] |   | I-imm. |
| -  | -inst[31]- |         | inst[30:25] | inst[11:8]  | inst[7]  |   | S-imm. |
| -  | -inst[31]- | inst[7] | inst[30:25] | inst[11:8]  | 0        |   | B-imm. |

Upper bits sign-extended from  $\text{inst}[31]$  always

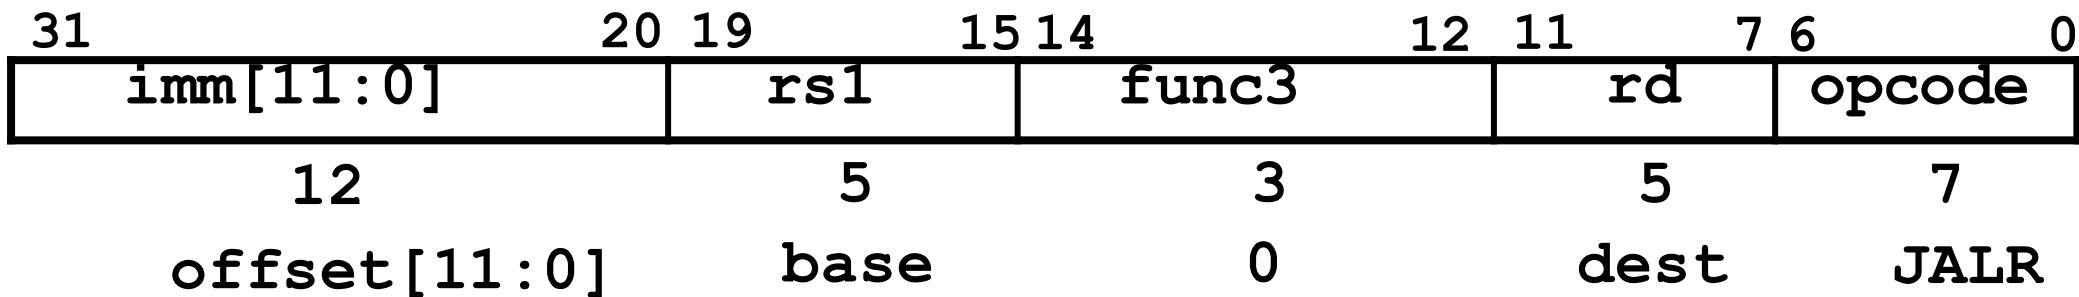
Only bit 7 of instruction changes role in immediate between S and B

# Lighting Up Branch Path



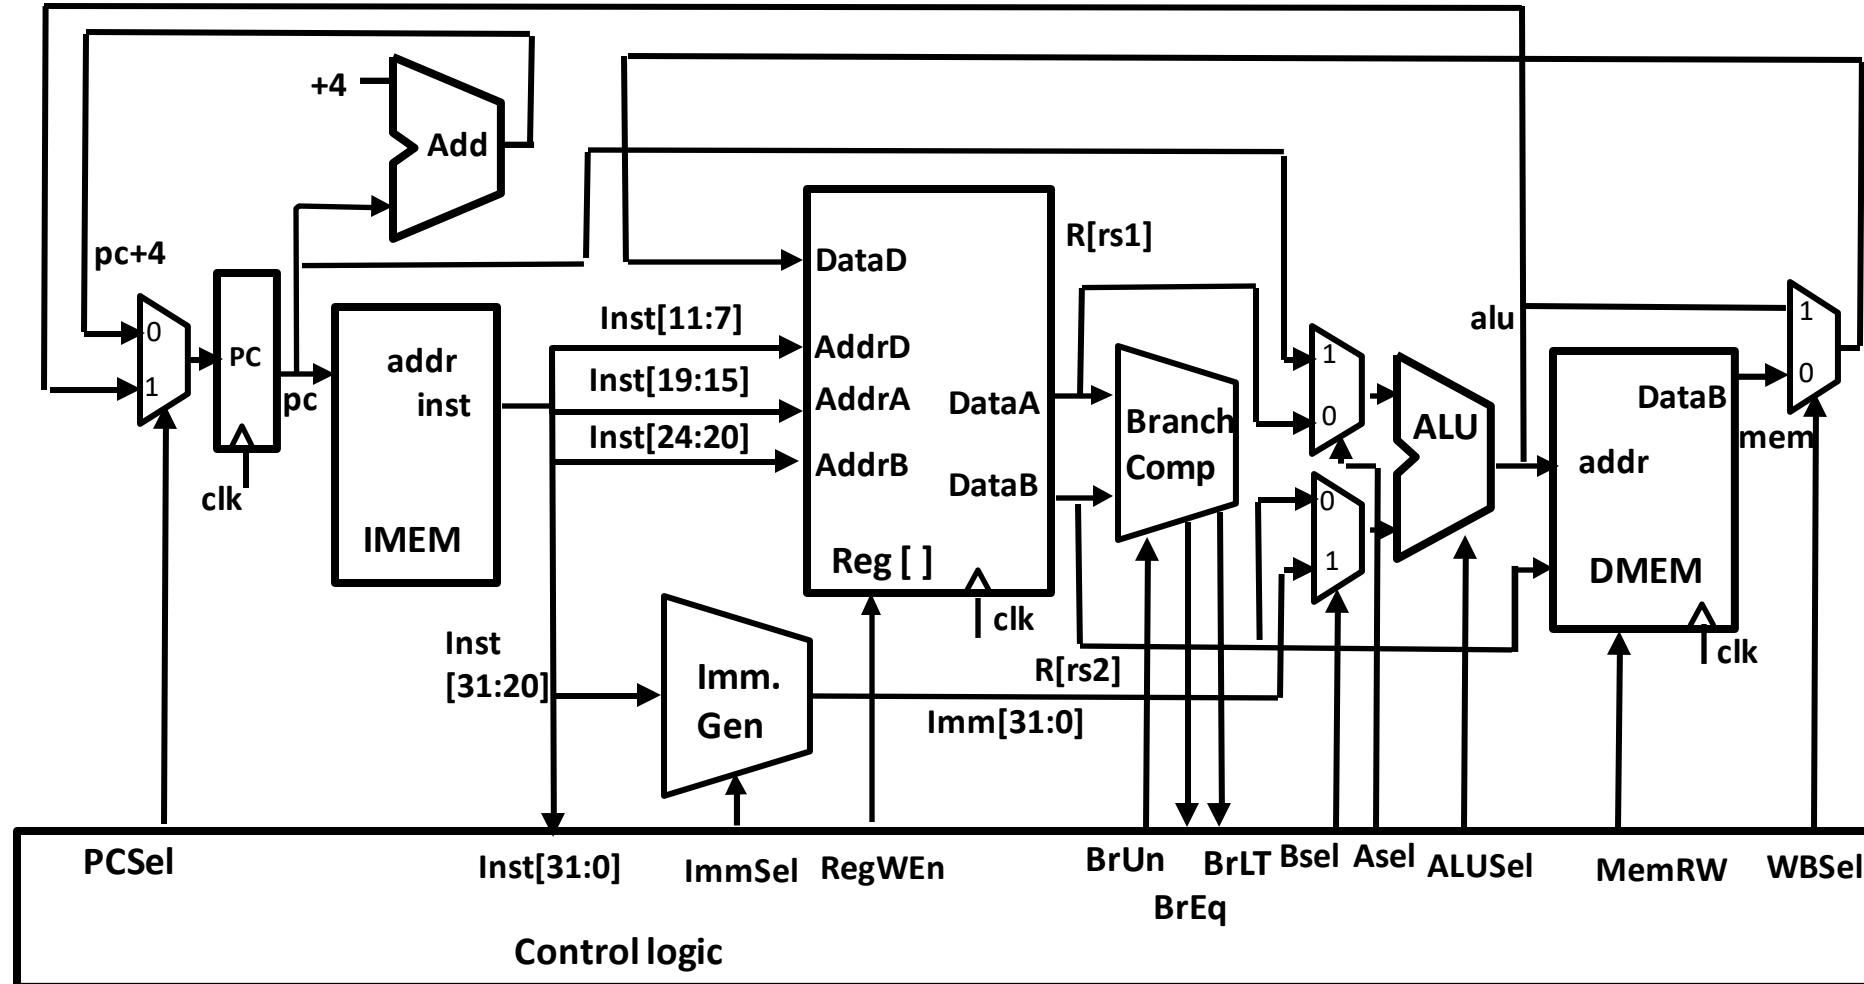
# **Adding JALR to Datapath**

# Let's Add JALR(I-Format)

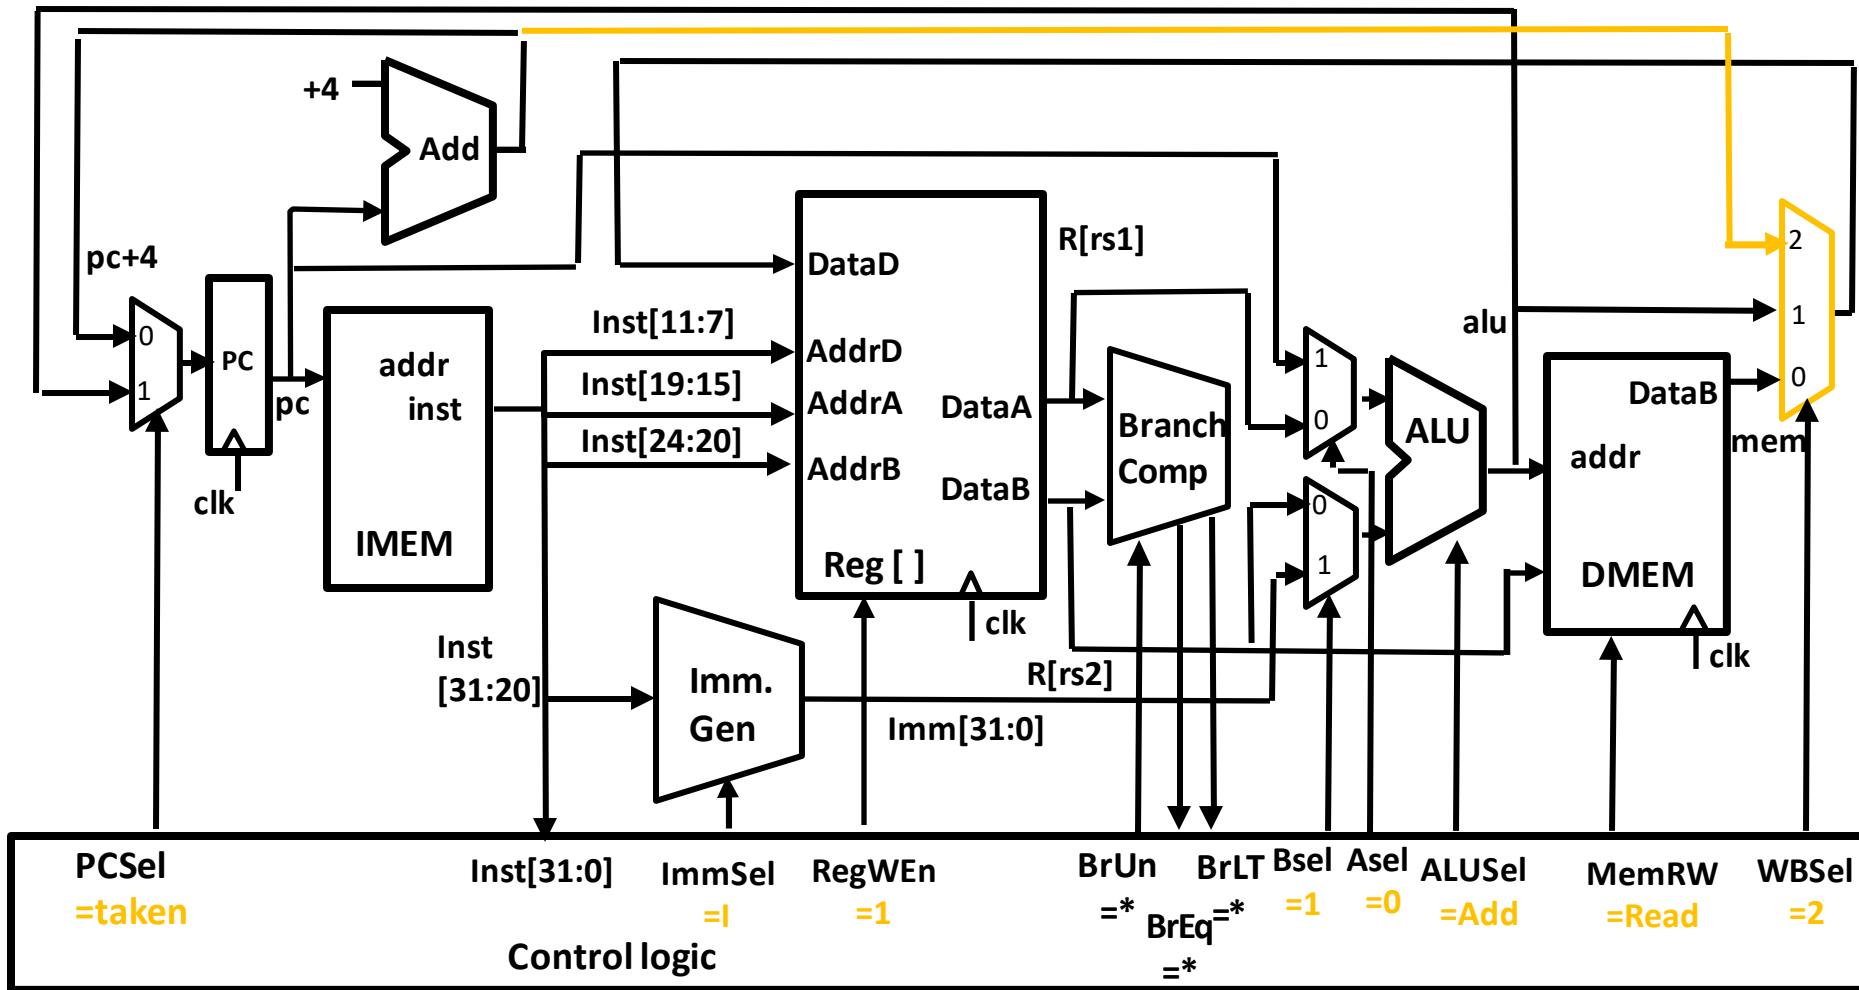


- **JALR rd, rs, immediate**
- Two changes to the state
  - Writes PC+4 to **rd** (return address)
  - Sets PC = **rs1 + immediate**
  - Uses same immediates as arithmetic and loads
    - **no** multiplication by 2 bytes
    - LSB is ignored

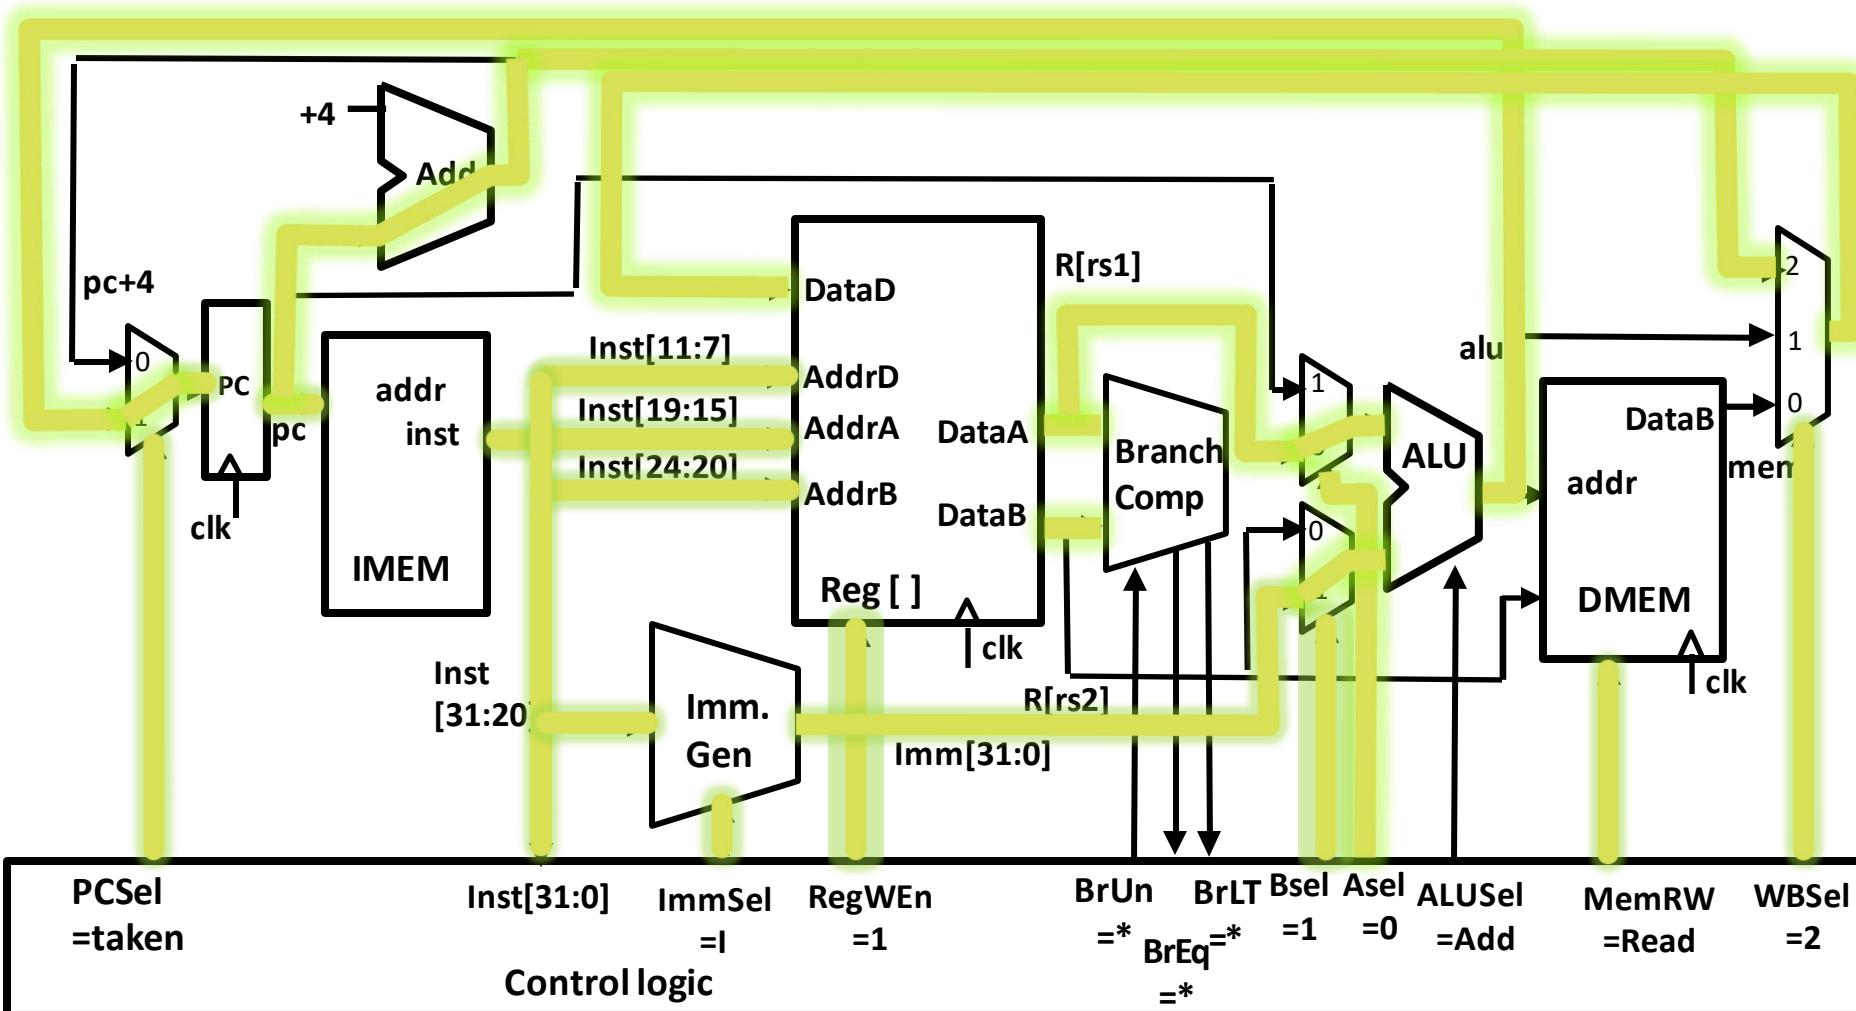
# Datapath So Far, With Branches



# Datapath With JALR

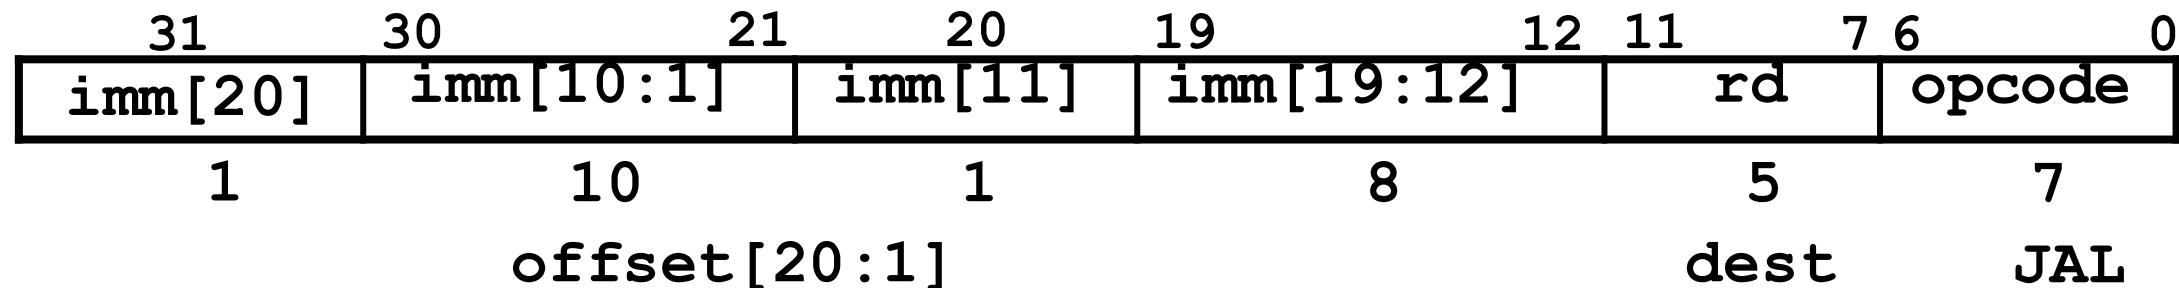


# Datapath With JALR



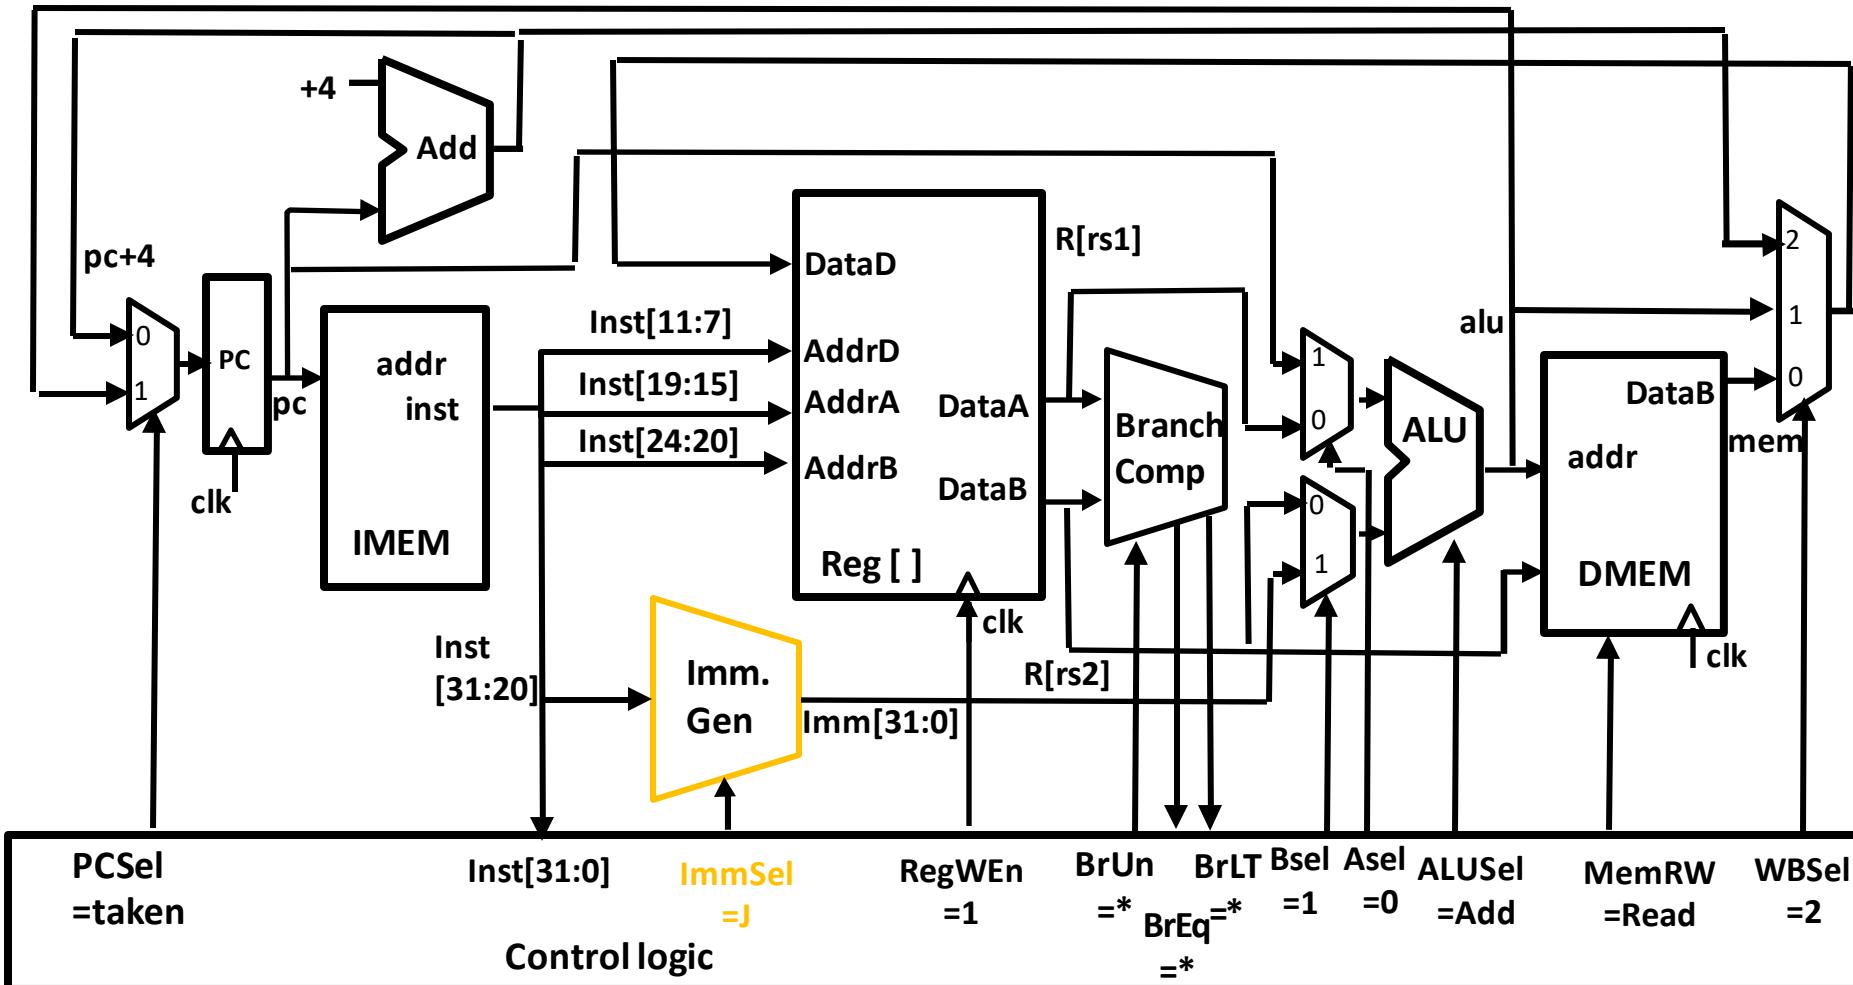
# Adding JAL

# JFormat for Jump Instructions

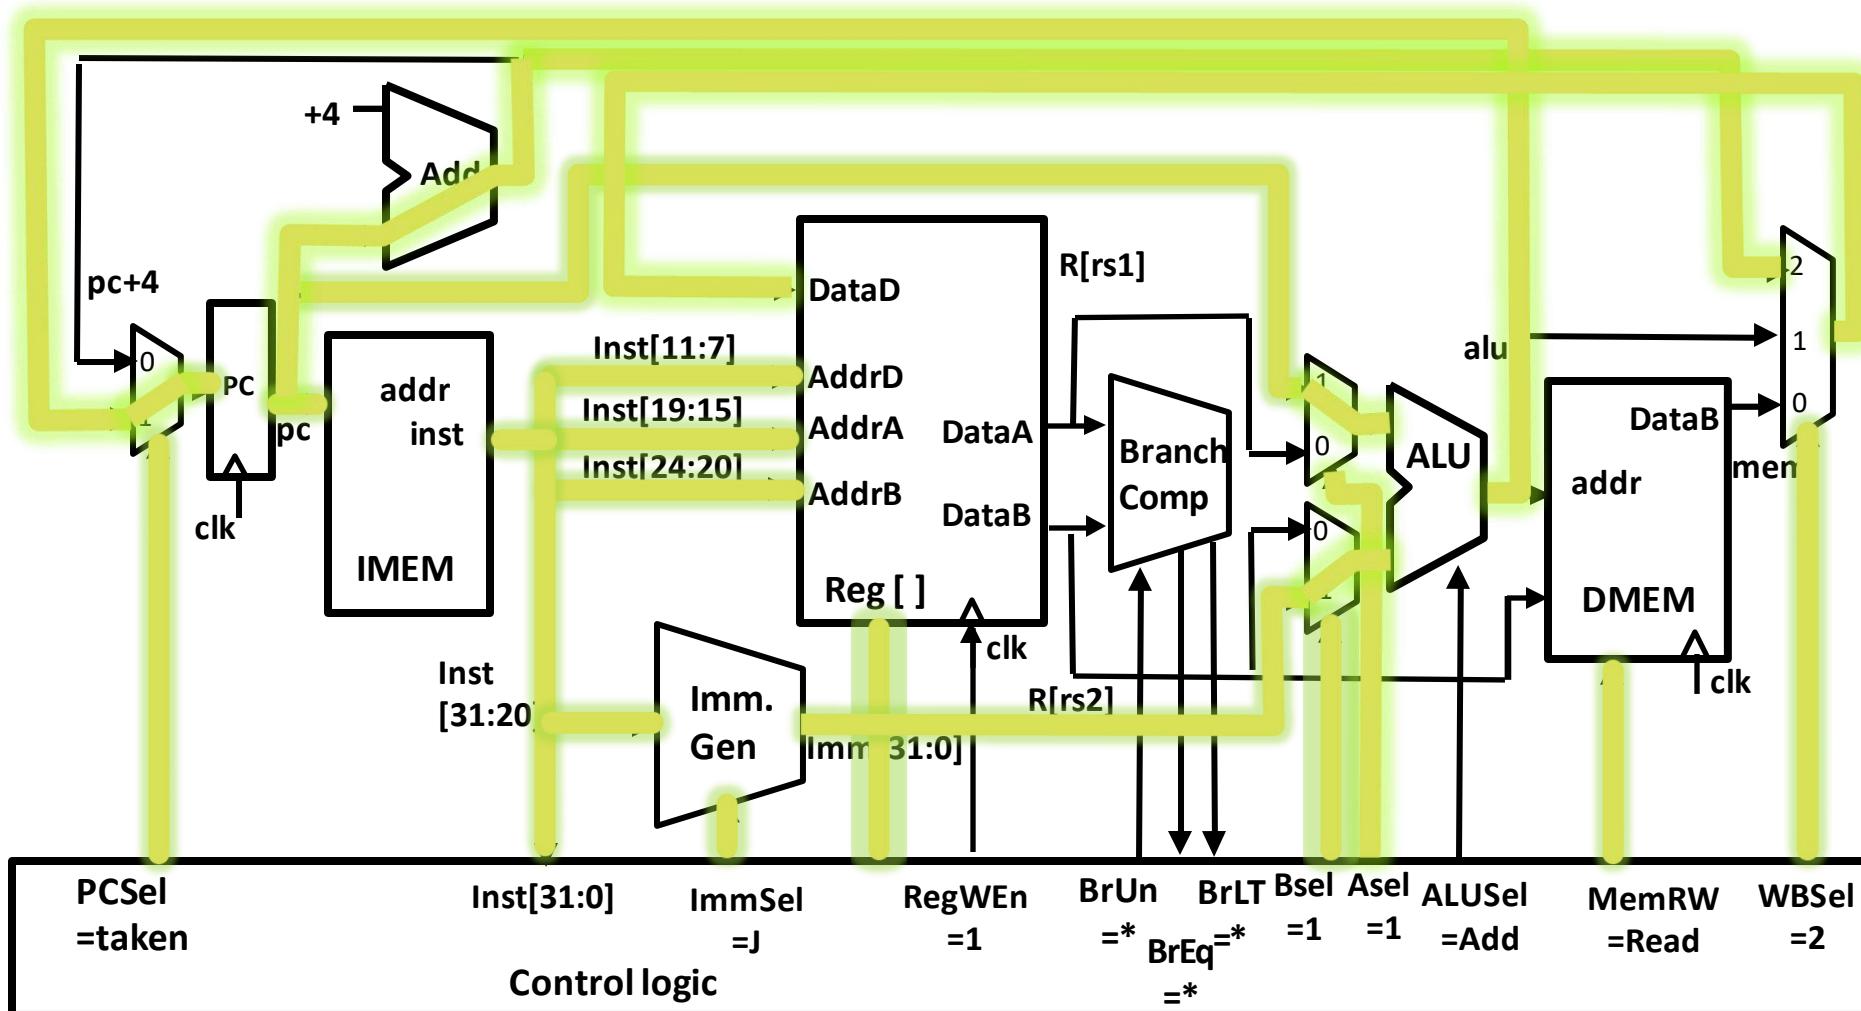


- Two changes to the state
  - `jal` saves PC+4 in register **rd** (the return address)
  - Set PC = PC + offset (PC-relative jump)
- Target somewhere within  $\pm 2^{19}$  locations, 2 bytes apart
  - $\pm 2^{18}$  32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

# Datapath with JAL

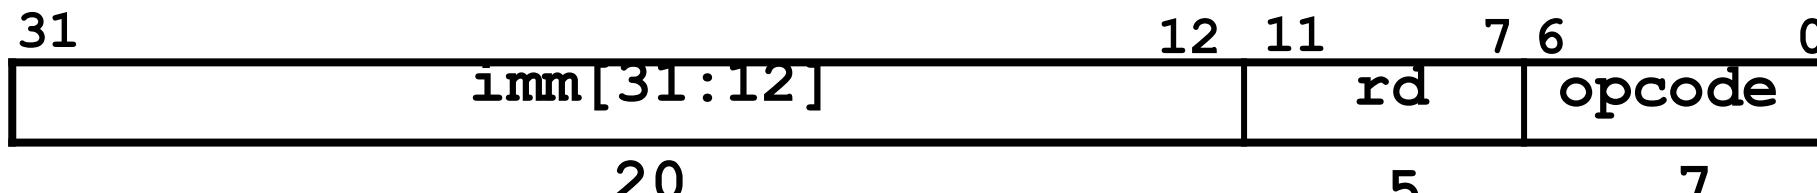


# LightUp JAL Path



# **Adding U-Types**

# U-Format for “Upper Immediate” Instructions



**U-immediate [31:12]**

**dest**

**LUI**

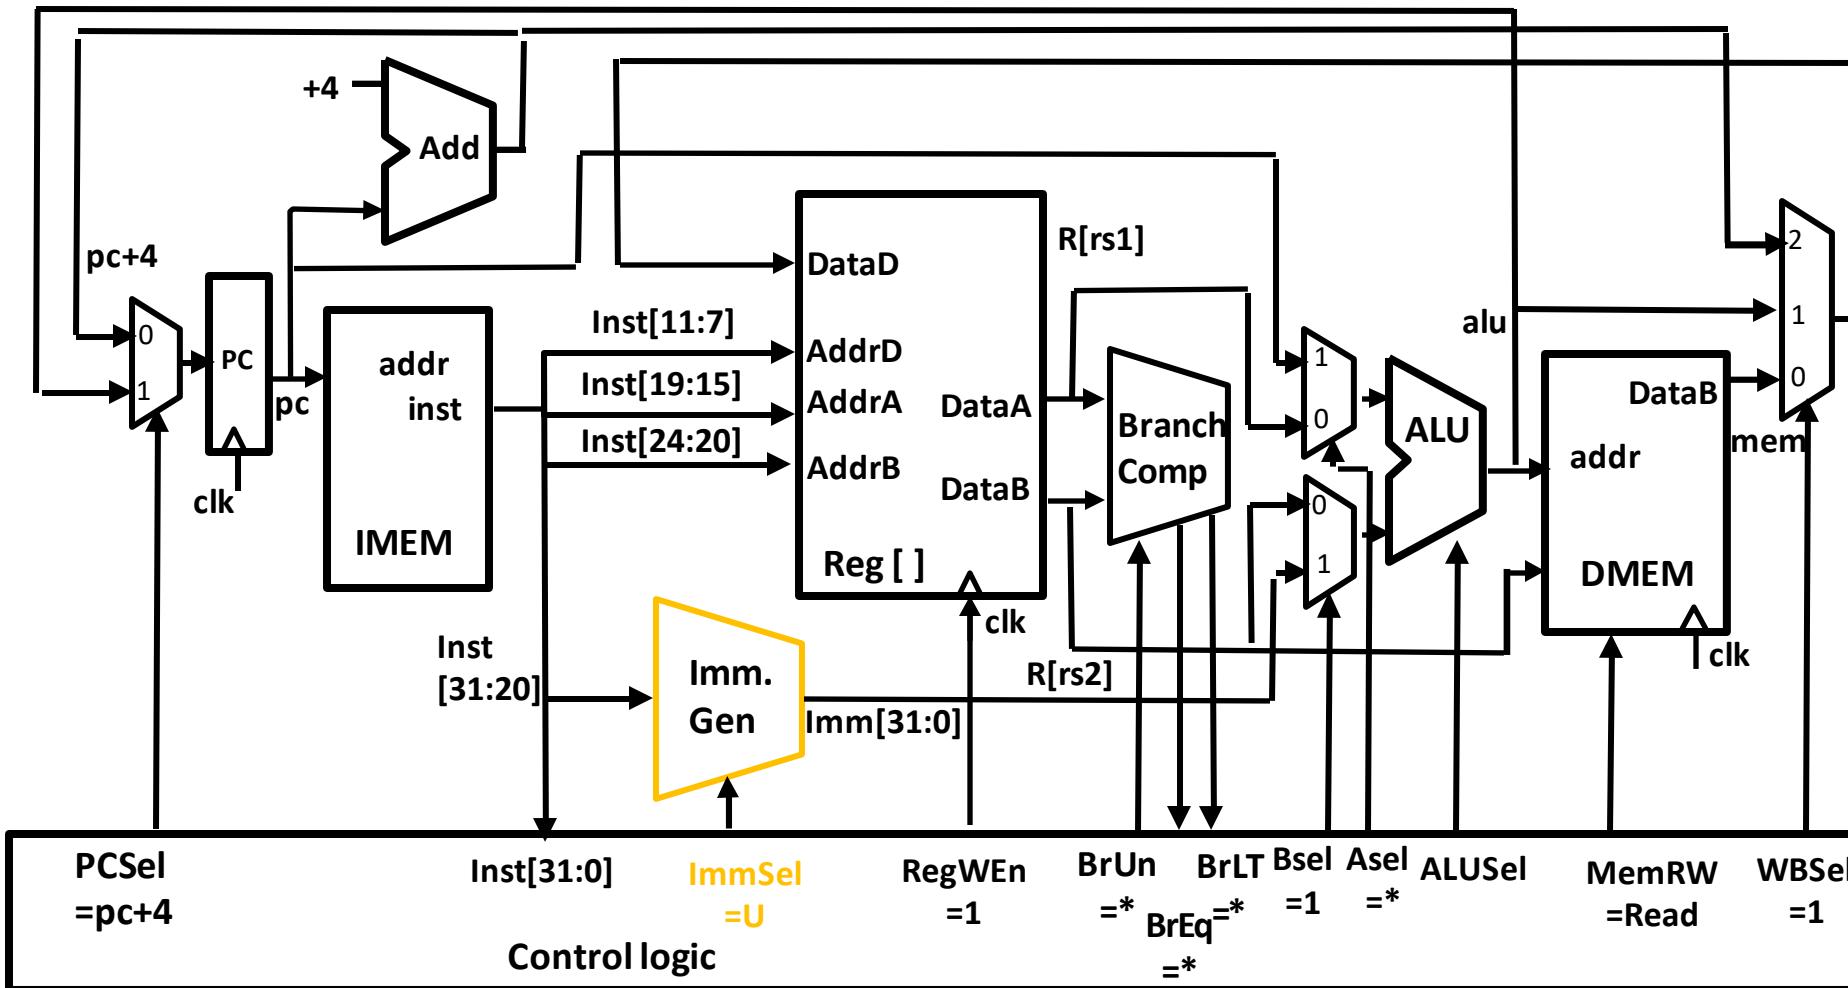
**U-immediate [31:12]**

**dest**

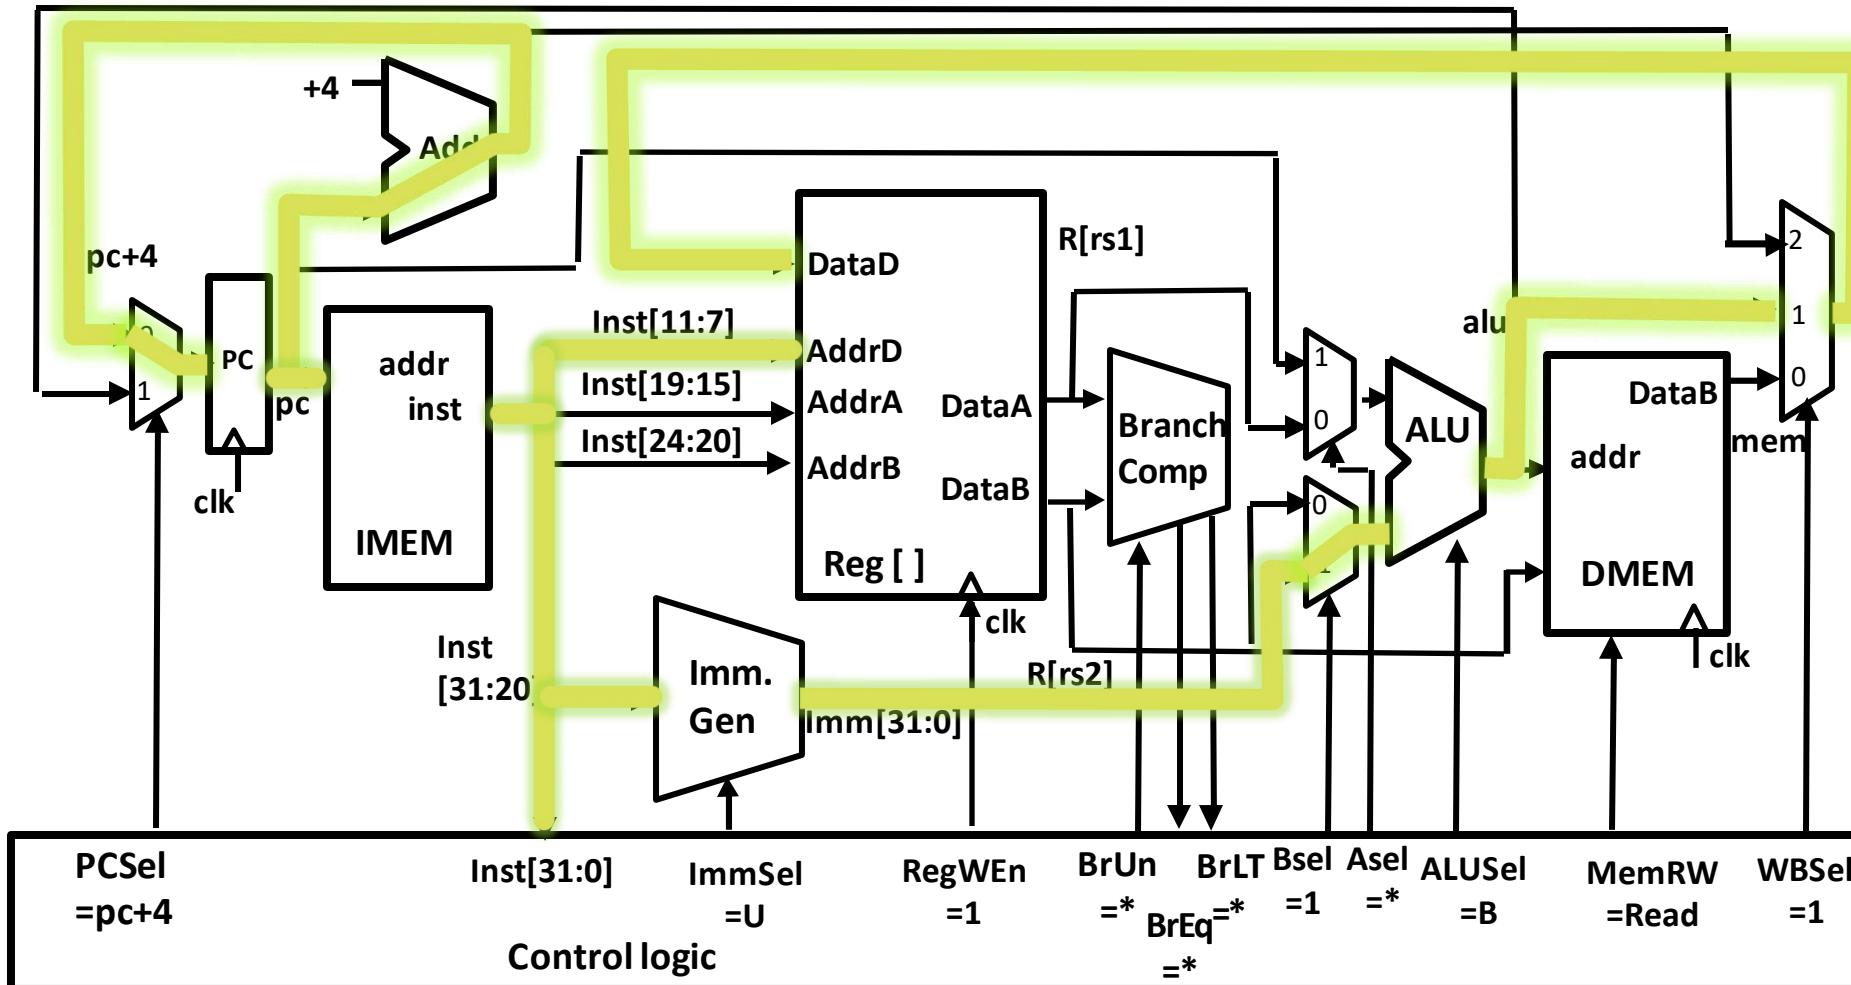
**AUIPC**

- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, **rd**
- Used for two instructions
  - **lui** – Load Upper Immediate
  - **auipc** – Add Upper Immediate to PC

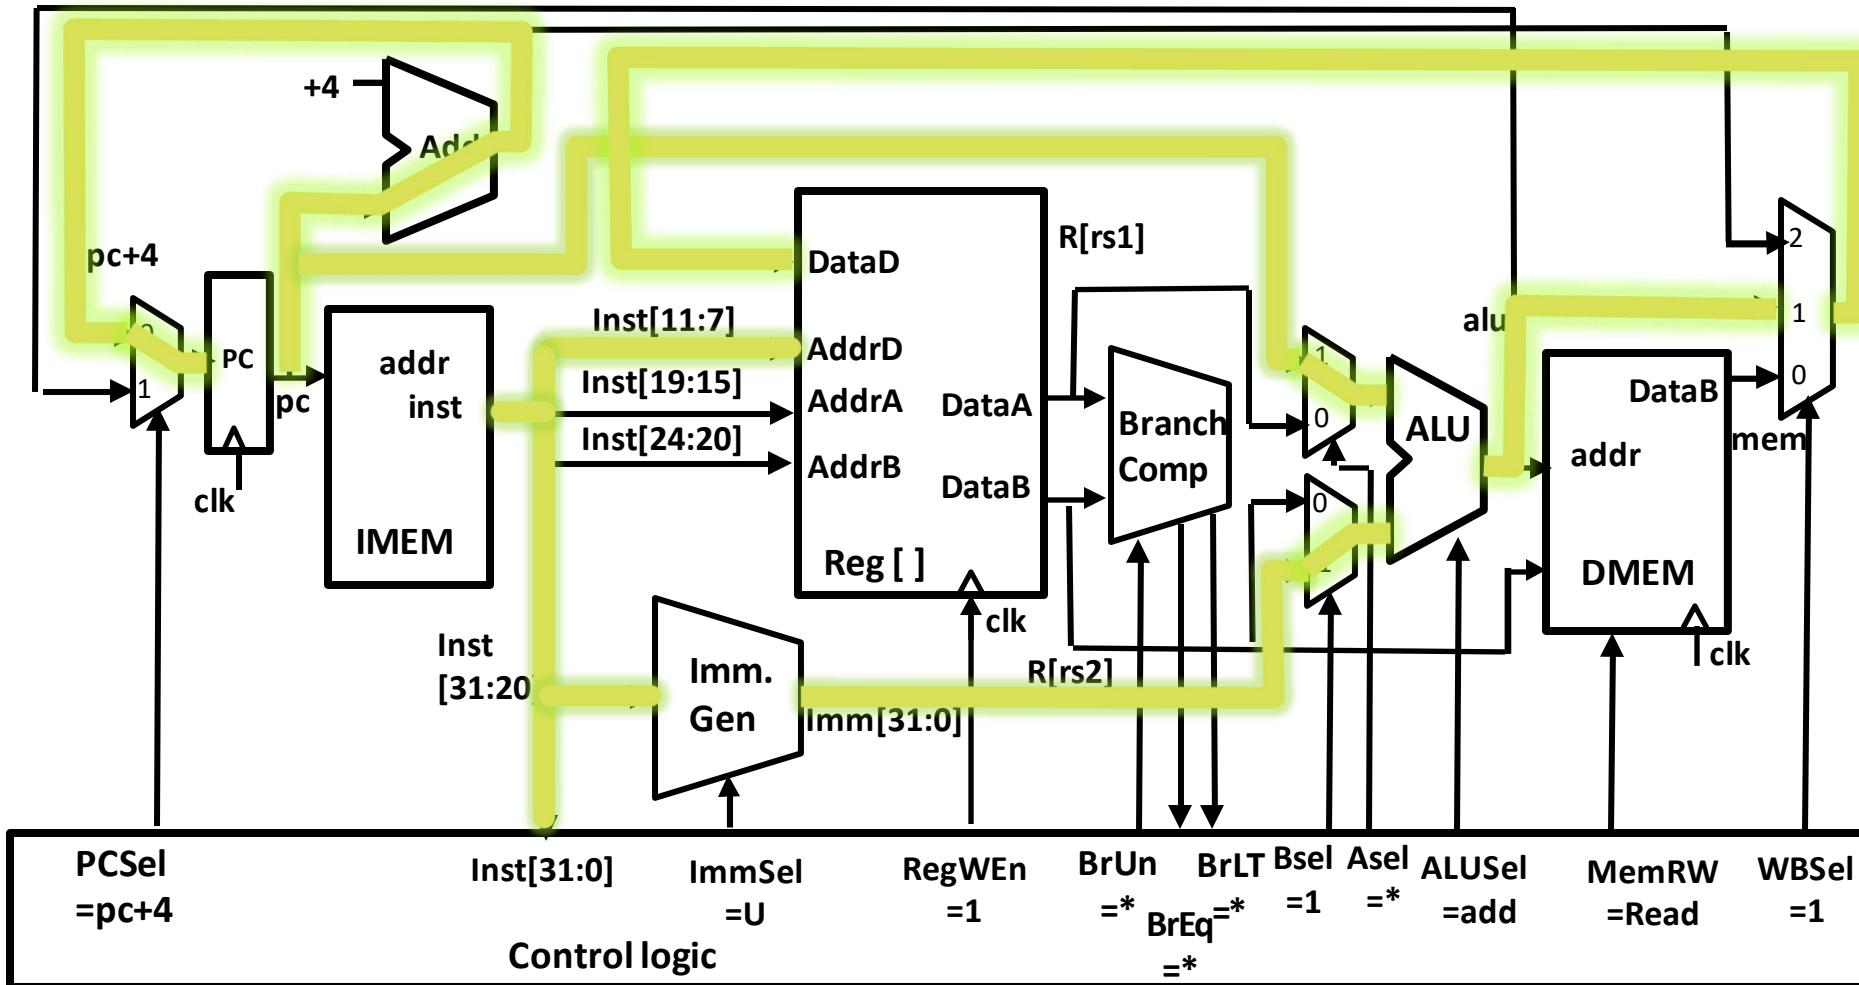
# Datapath With LUI, AUIPC



# Lighting Up LUI

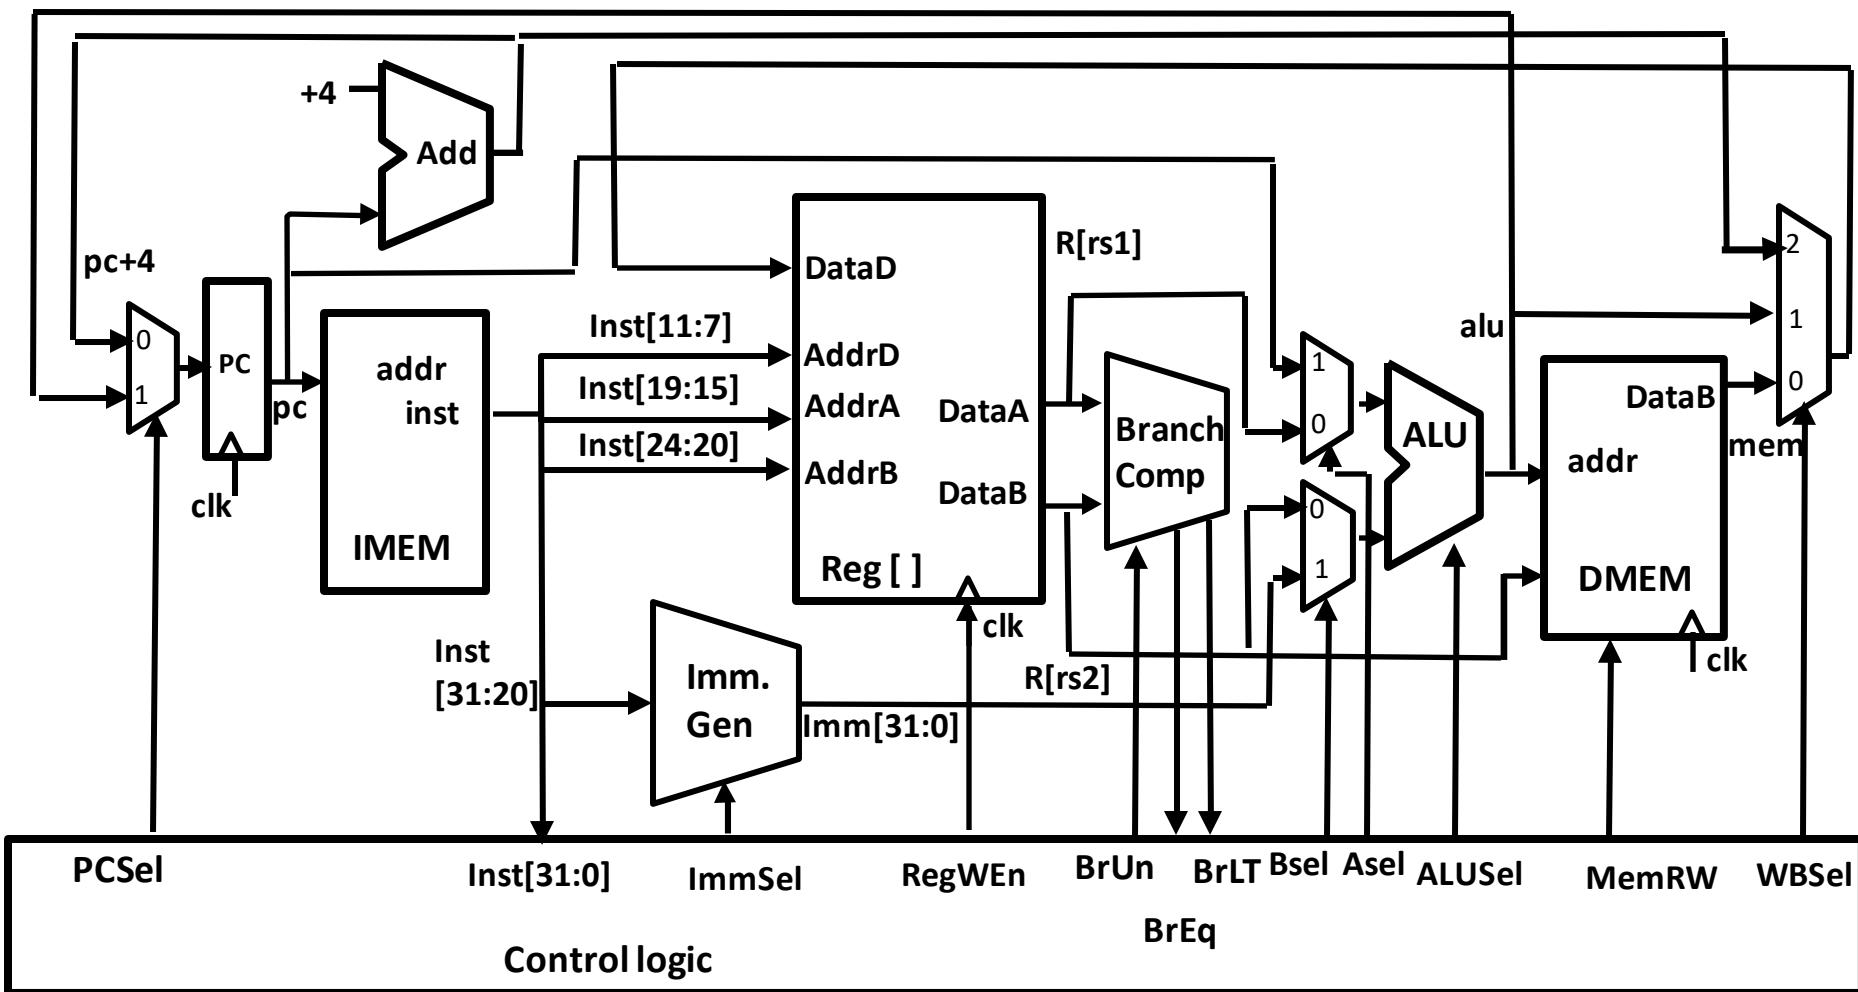


# Lighting Up **AUIPC**



**“And In Conclusion...”**

# Complete RV32I Datapath!





# Complete RV32I ISA!

## Open RISC-V Reference Card

| Base Integer Instructions: RV32I |                         |     |                   |             |                      |
|----------------------------------|-------------------------|-----|-------------------|-------------|----------------------|
| Category                         | Name                    | Fmt | RV32I Base        | Category    | Name                 |
| Shifts                           | Shift Left Logical      | R   | SLL rd,rs1,rs2    | Loads       | Load Byte            |
|                                  | Shift Left Log. Imm.    | I   | SLLI rd,rs1,shamt |             | Load Halfword        |
|                                  | Shift Right Logical     | R   | SRL rd,rs1,rs2    |             | Load Byte Unsigned   |
|                                  | Shift Right Log. Imm.   | I   | SRLI rd,rs1,shamt |             | Load Half Unsigned   |
|                                  | Shift Right Arithmetic  | R   | SRA rd,rs1,rs2    |             | Load Word            |
|                                  | Shift Right Arith. Imm. | I   | SRAI rd,rs1,shamt | Stores      | Store Byte           |
| Arithmetic                       | ADD                     | R   | ADD rd,rs1,rs2    |             | Store Halfword       |
|                                  | ADD Immediate           | I   | ADDI rd,rs1,imm   |             | Store Word           |
|                                  | SUBtract                | R   | SUB rd,rs1,rs2    | Branches    | Branch =             |
|                                  | Load Upper Imm          | U   | LUI rd,imm        |             | Branch ≠             |
|                                  | Add Upper Imm to PC     | U   | AUIPC rd,imm      |             | Branch <             |
| Logical                          | XOR                     | R   | XOR rd,rs1,rs2    |             | Branch ≥             |
|                                  | XOR Immediate           | I   | XORI rd,rs1,imm   |             | Branch < Unsigned    |
|                                  | OR                      | R   | OR rd,rs1,rs2     |             | Branch ≥ Unsigned    |
|                                  | OR Immediate            | I   | ORI rd,rs1,imm    | Jump & Link | J&L                  |
|                                  | AND                     | R   | AND rd,rs1,rs2    |             | Jump & Link Register |
|                                  | AND Immediate           | I   | ANDI rd,rs1,imm   |             |                      |
| Compare                          | Set <                   | R   | SLT rd,rs1,rs2    | Synch       | Synch thread         |
|                                  | Set < Immediate         | I   | SLTI rd,rs1,imm   |             |                      |
|                                  | Set < Unsigned          | R   | SLTU rd,rs1,rs2   | Environment | CALL                 |
|                                  | Set < Imm Unsigned      | I   | SLTIU rd,rs1,imm  |             | BREAK                |

Not in  
61C

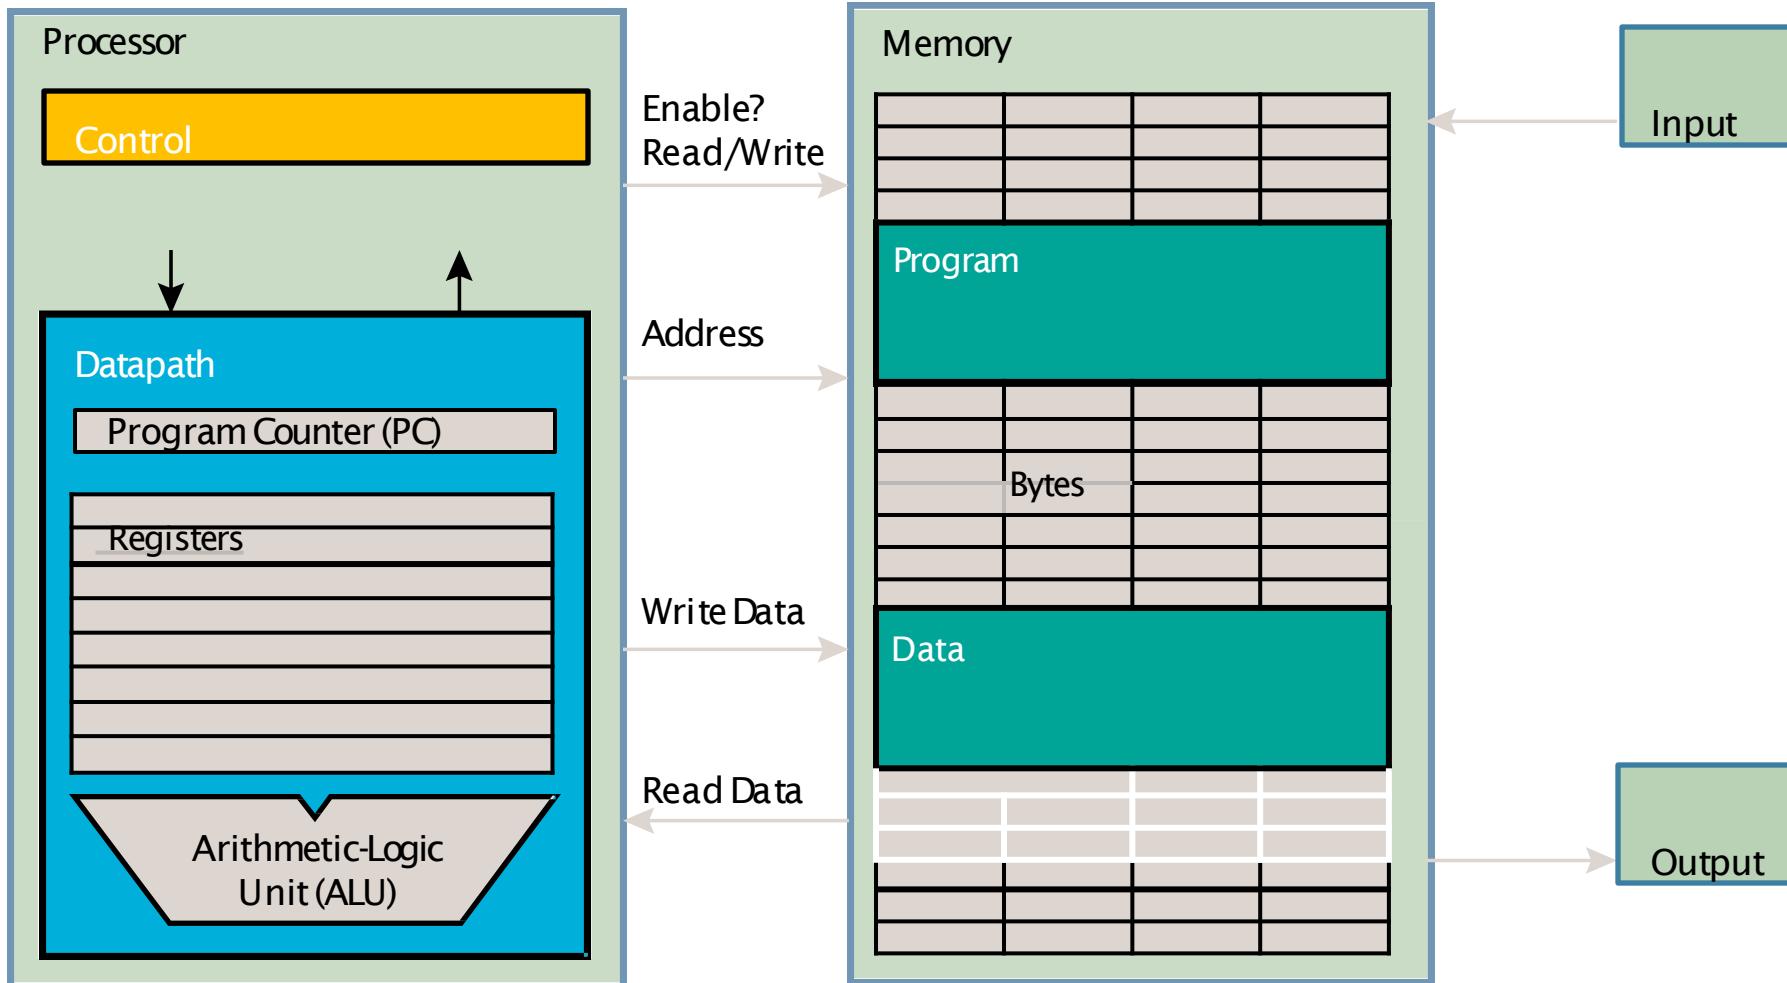
# Review

---

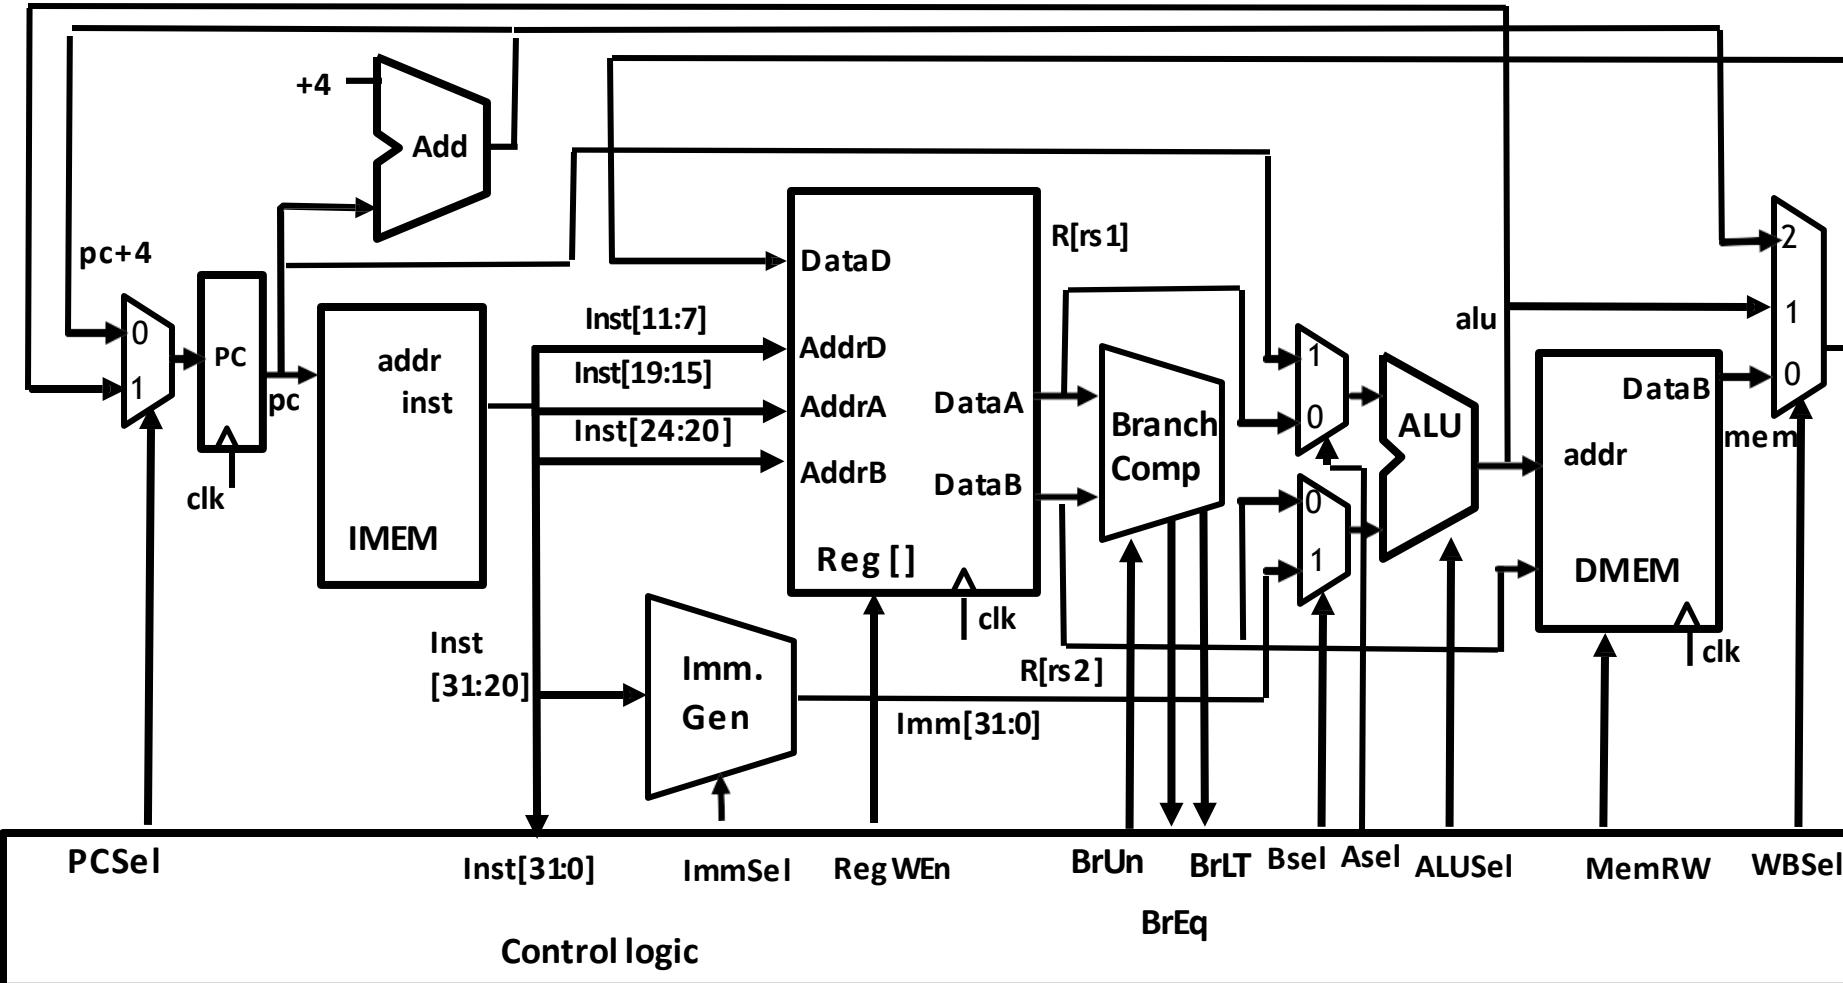
- We have designed a complete datapath
  - Capable of executing all RISC-V instructions in one cycle each
  - Not all units (hardware) used by all instructions
- 5 Phases of execution
  - IF, ID, EX, MEM, WB
  - Not all instructions are active in all phases
- Controller specifies how to execute instructions
  - We still need to design it

# Datapath Control

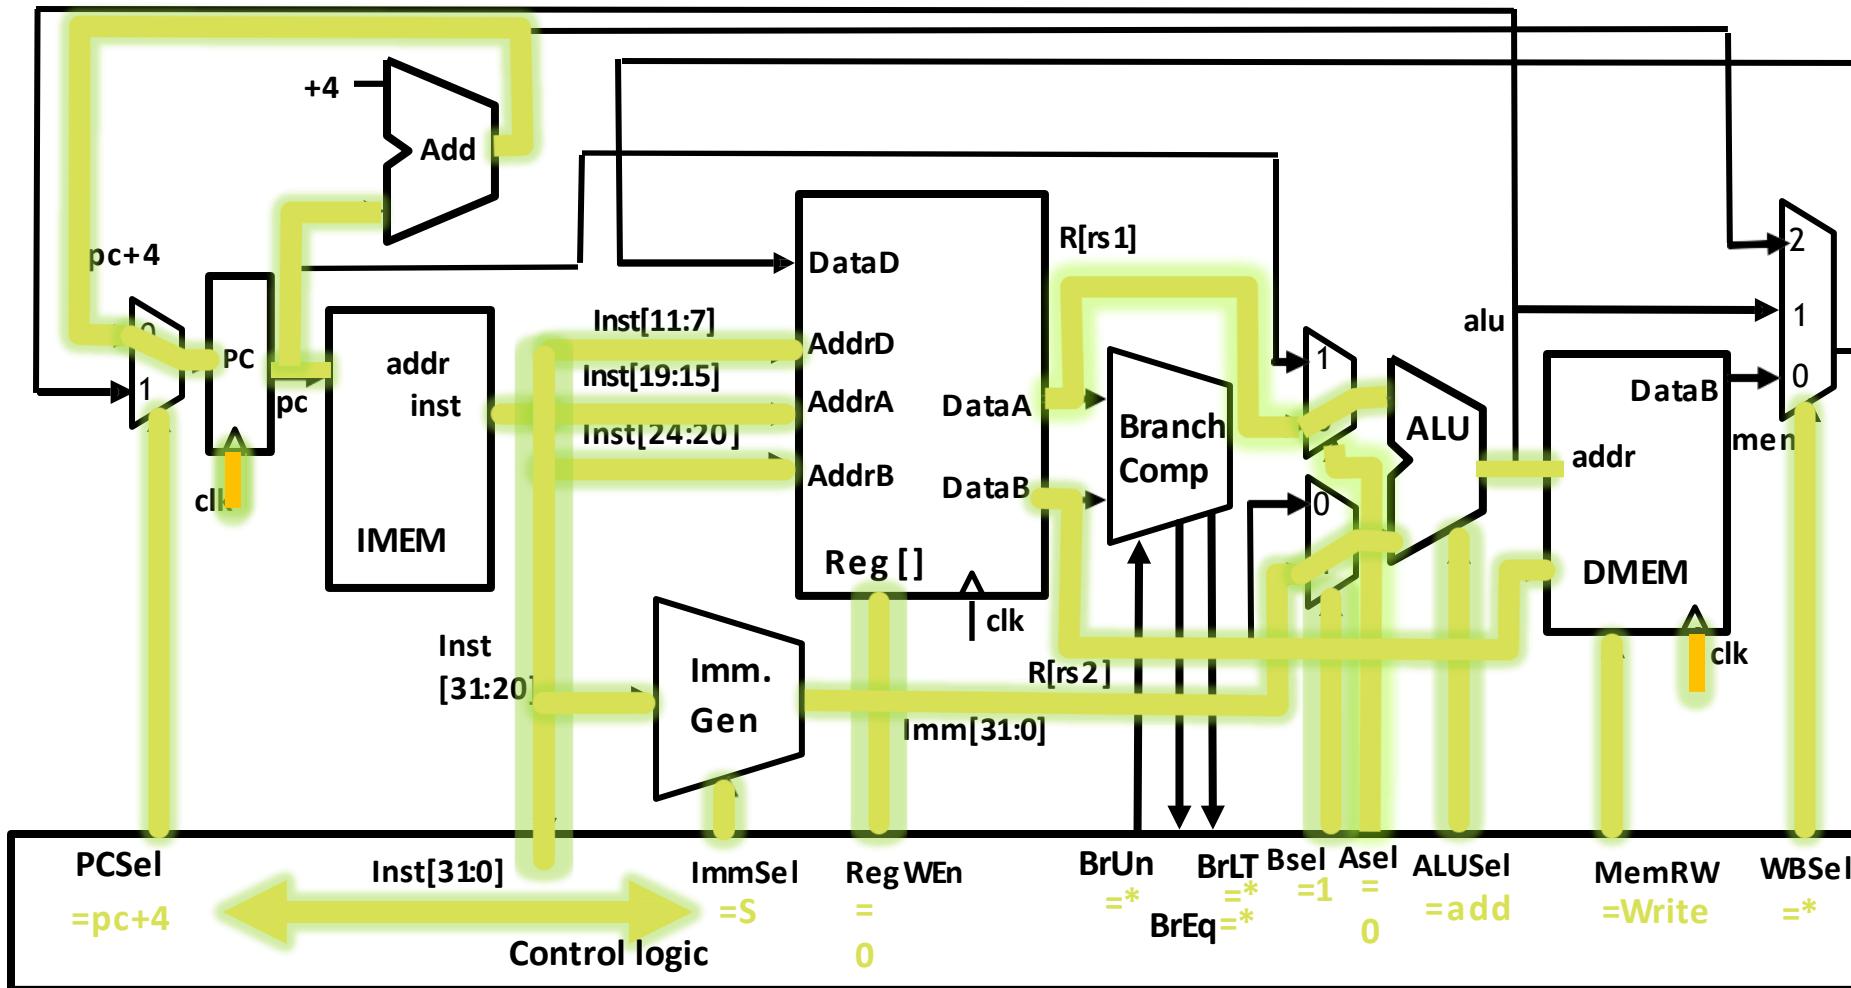
# Our Single-Core Processor



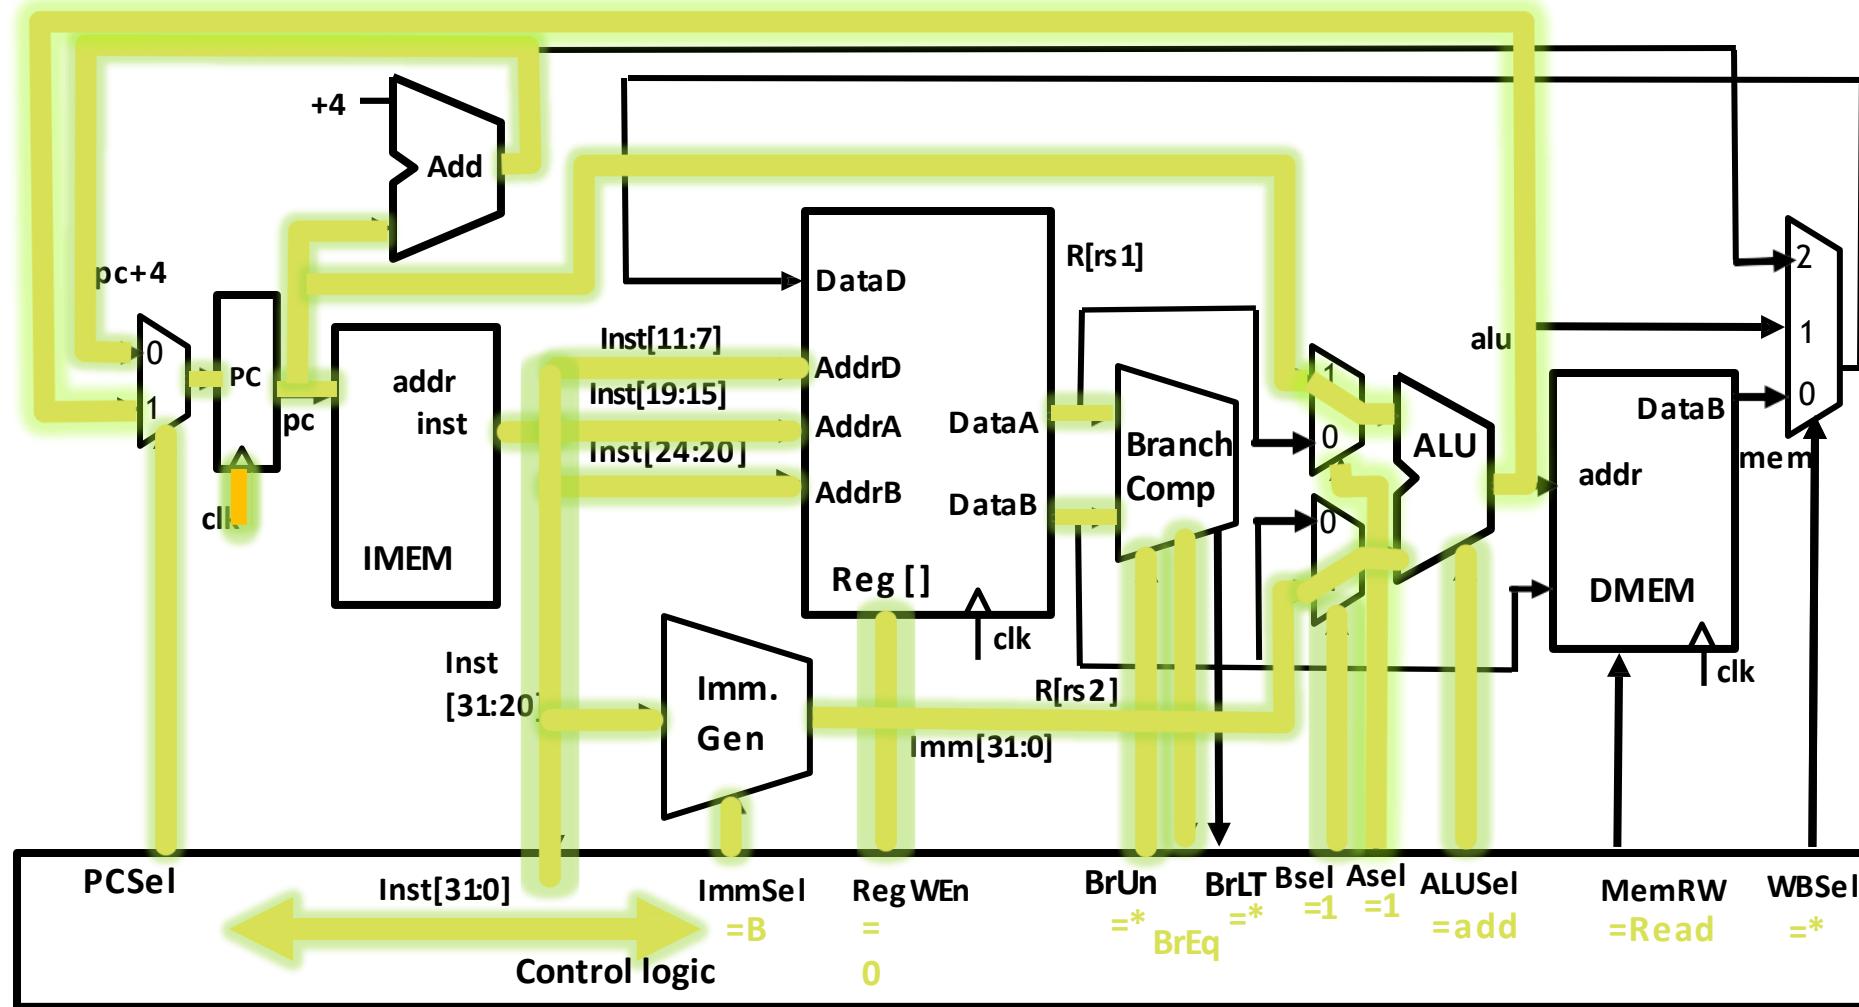
# Single-Cycle RV32I Datapath and Control



# Example: sw

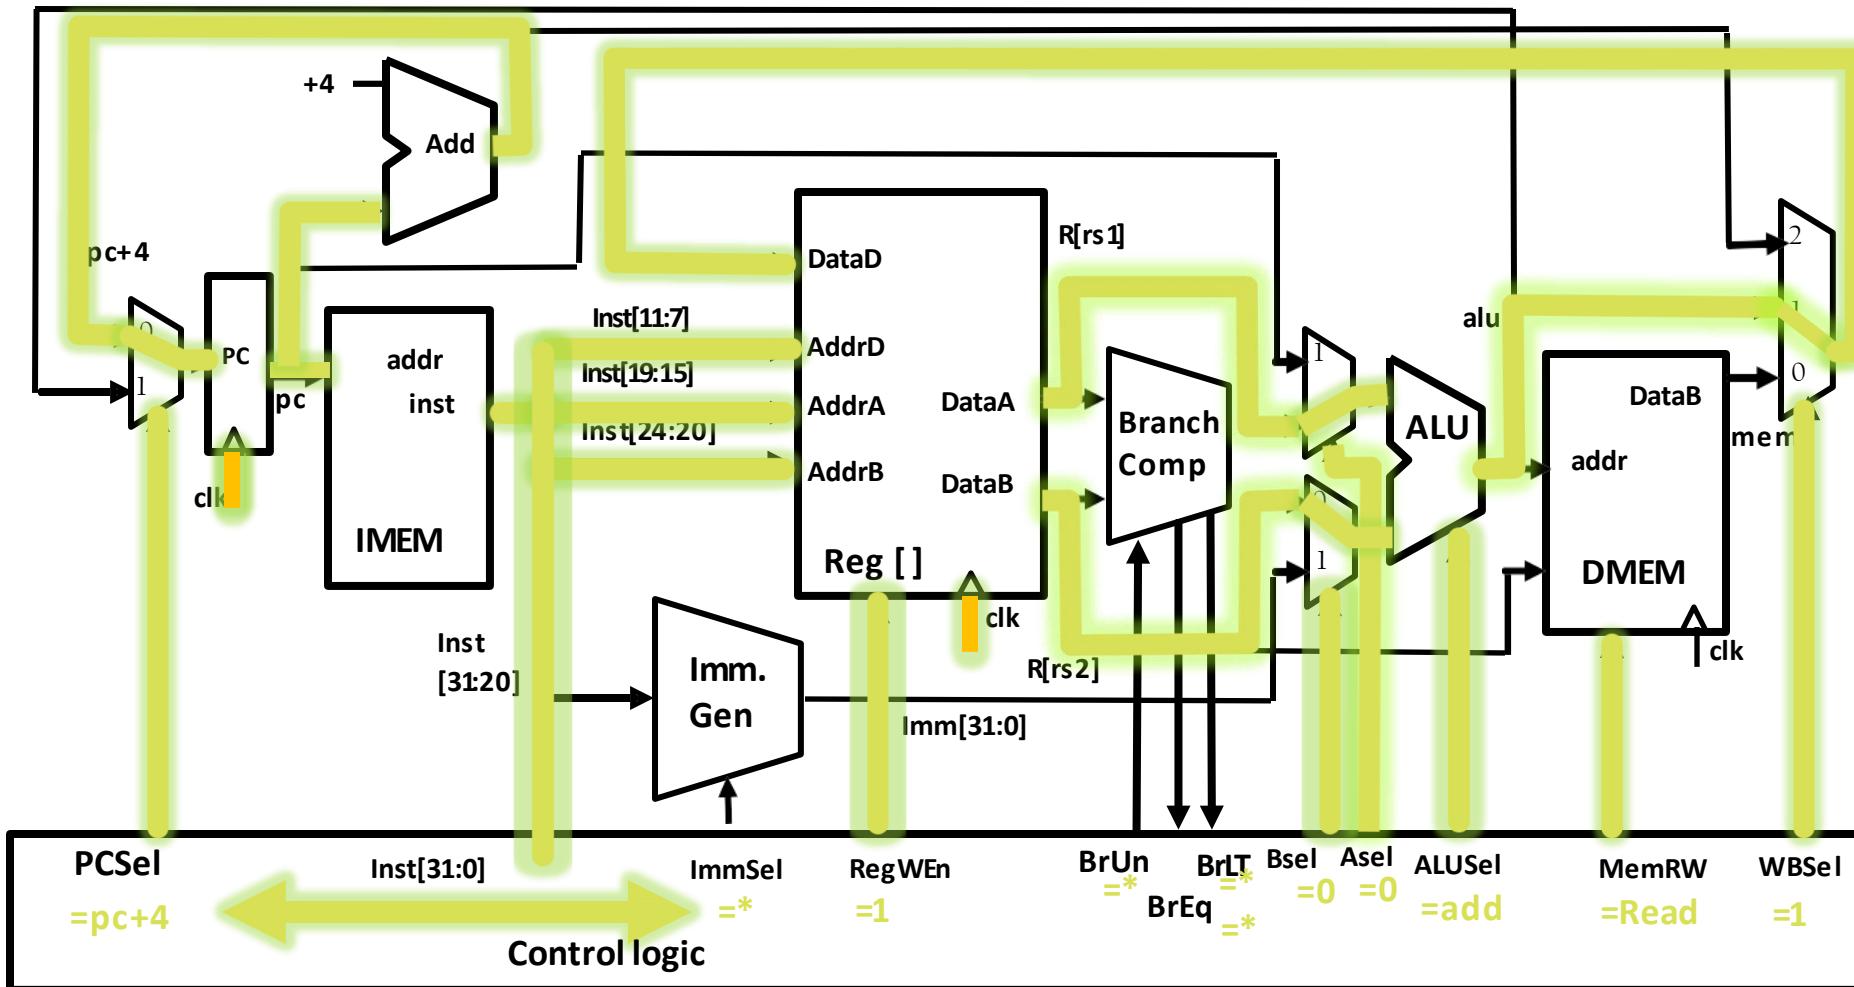


# Example: beq

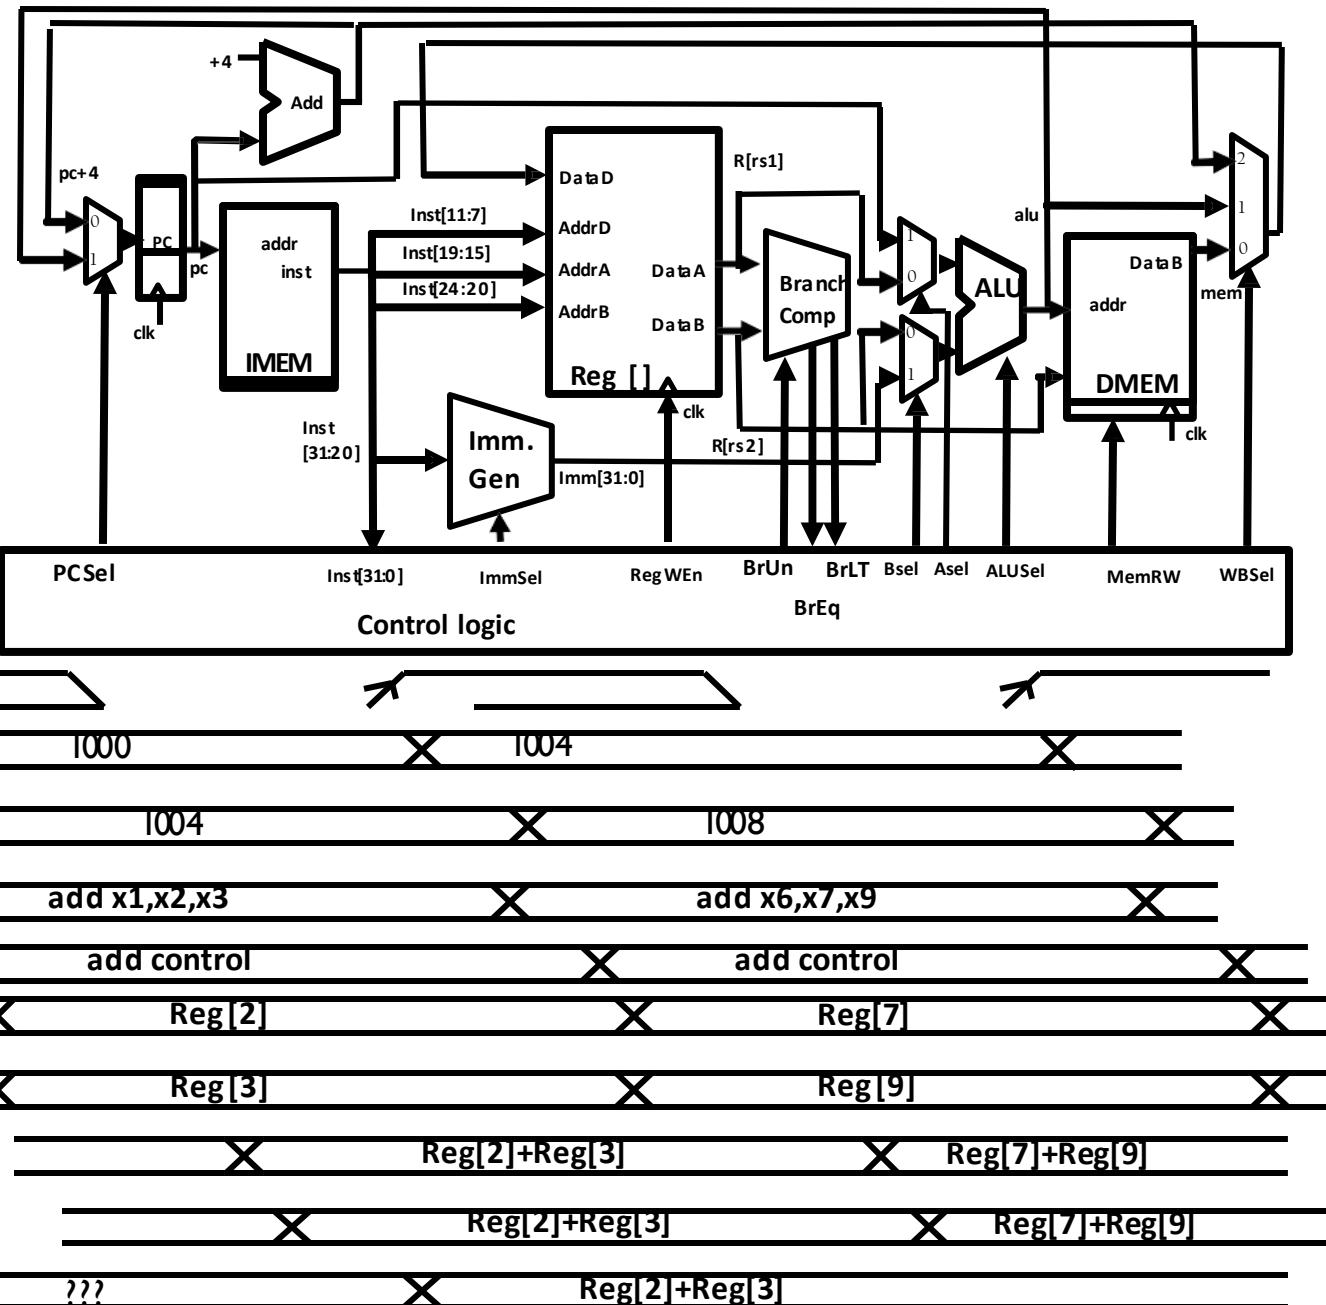


# Instruction Timing

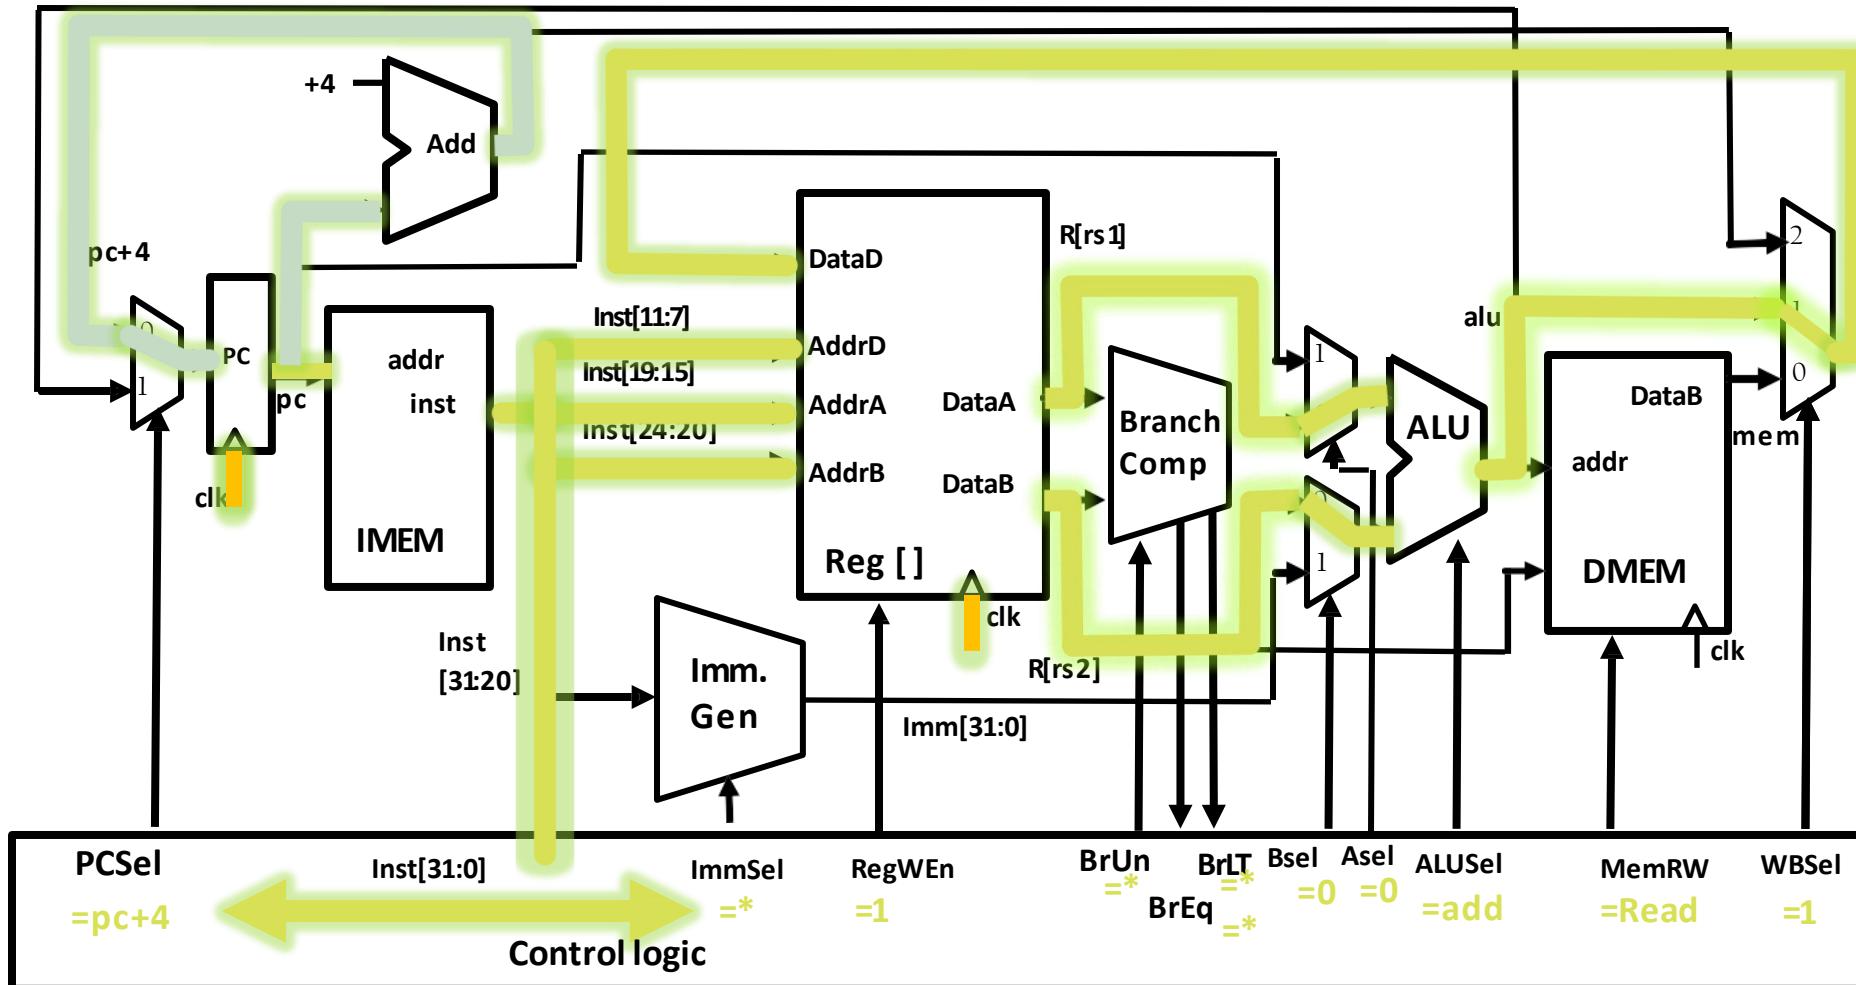
# Example: add



# Add Execution

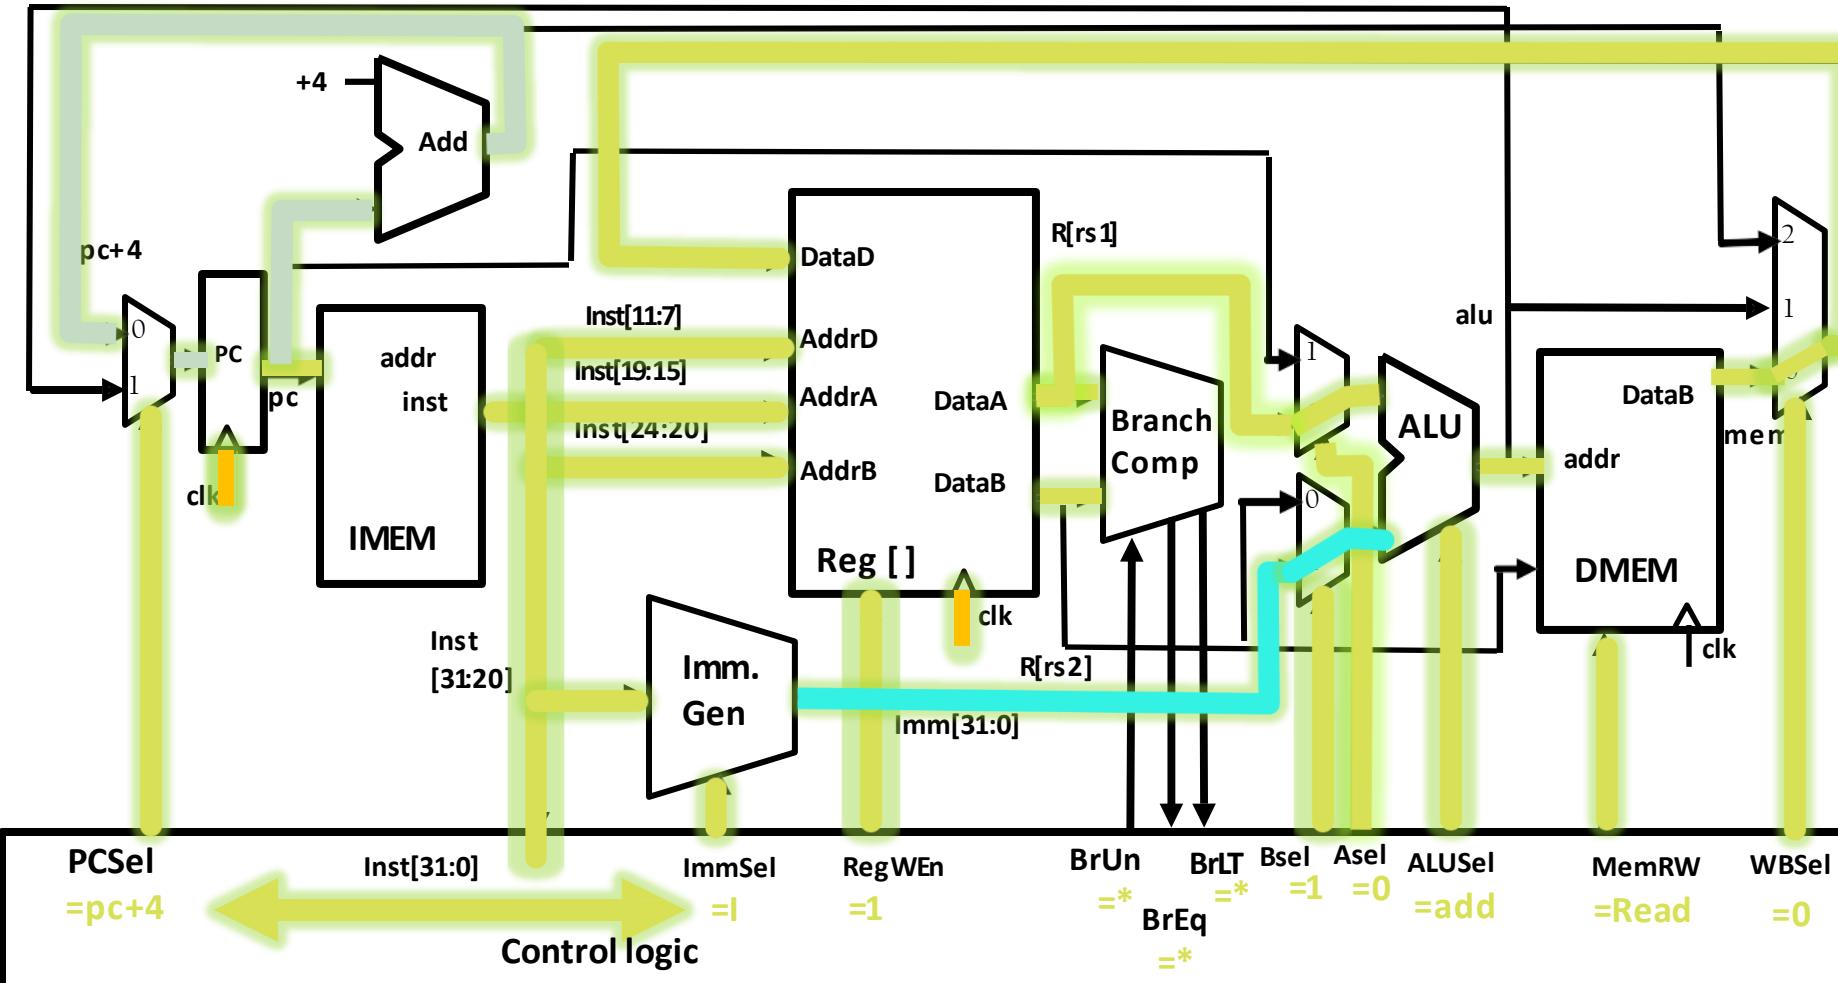


# Example: add timing



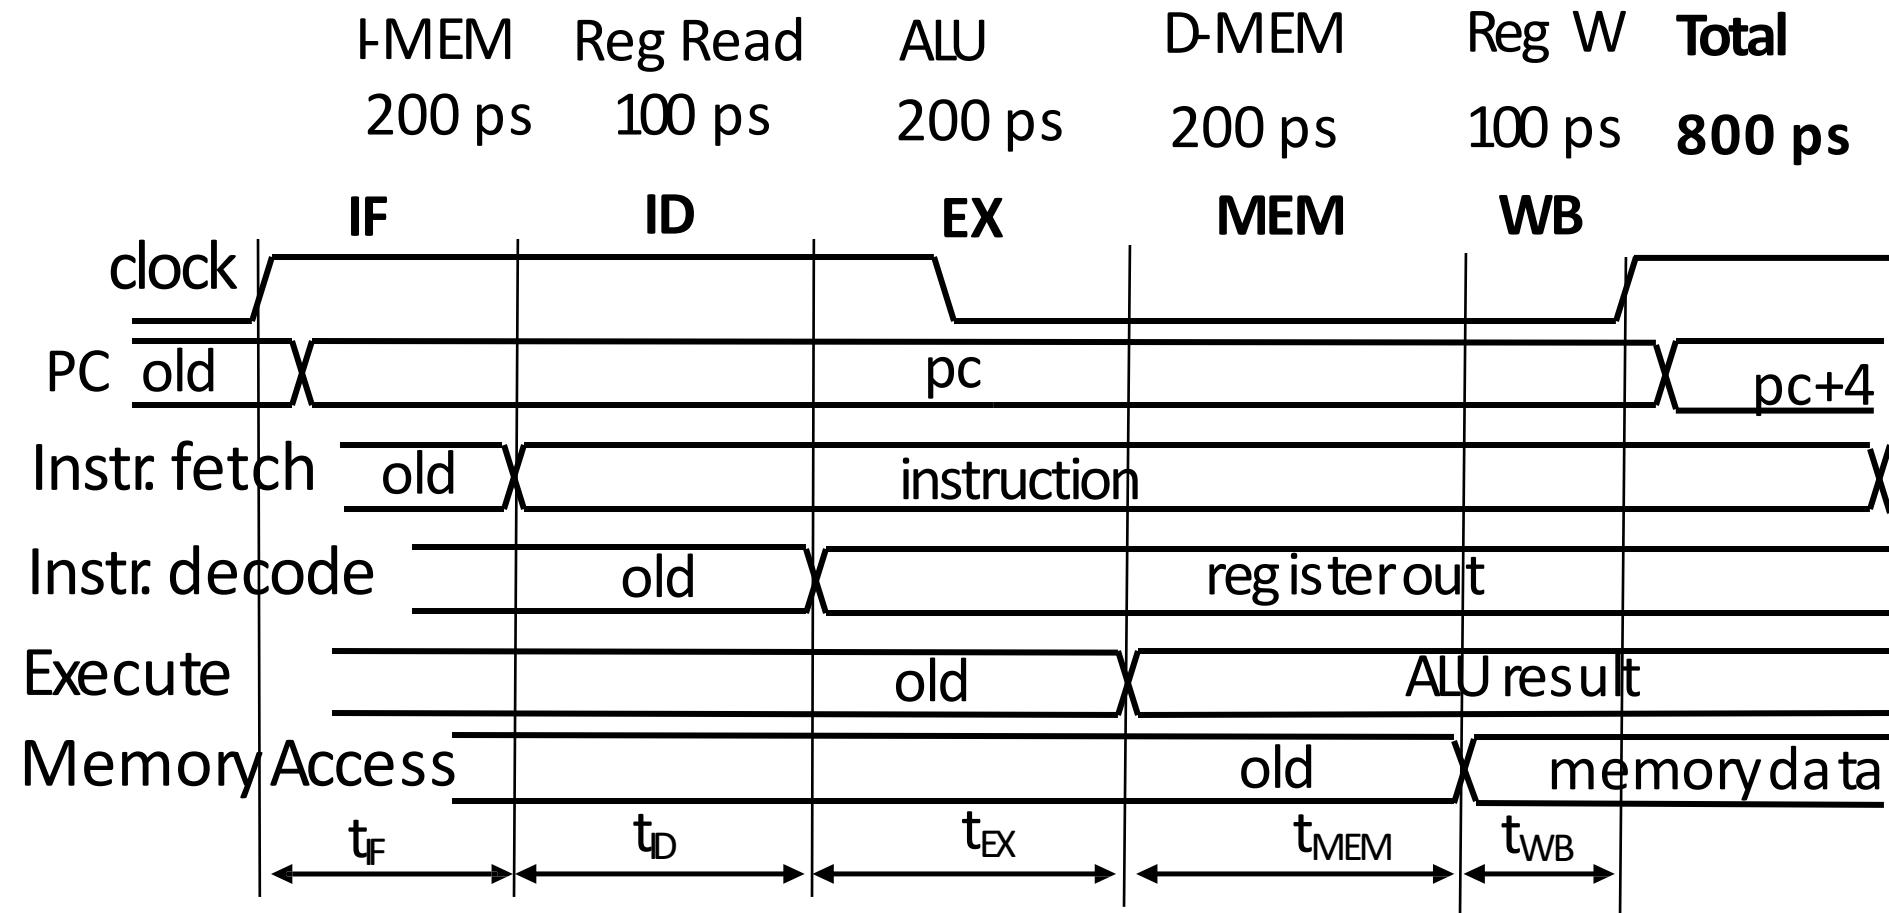
$$\begin{aligned}
 \text{Critical path} &= t_{\text{clk-q}} + \max \{ t_{\text{Add}} + t_{\text{mux}}, t_{\text{MEM}} + t_{\text{Reg}} + t_{\text{mux}} + t_{\text{ALU}} + t_{\text{mux}} \} + t_{\text{setup}} \\
 &= t_{\text{clk-q}} + t_{\text{MEM}} + t_{\text{Reg}} + t_{\text{mux}} + t_{\text{ALU}} + t_{\text{mux}} + t_{\text{setup}}
 \end{aligned}$$

# Example: lw



Critical path =  $t_{clk-q} + \max \{ t_{Add} + t_{mux}, t_{MEM} + t_{Imm} + t_{mux} + t_{ALU} + t_{DMEM} + t_{mux}, t_{IMEM} + t_{Reg} + t_{mux} + t_{ALU} + t_{DMEM} + t_{mux} \} + t_{setup}$

# Instruction Timing



# Instruction Timing

| Instr | IF = 200ps | ID = 100ps | ALU = 200ps | MEM=200ps | WB = 100ps | Total |
|-------|------------|------------|-------------|-----------|------------|-------|
| add   | X          | X          | X           |           | X          | 600ps |
| beq   | X          | X          | X           |           |            | 500ps |
| jal   | X          | X          | X           |           | X          | 600ps |
| lw    | X          | X          | X           | X         | X          | 800ps |
| sw    | X          | X          | X           | X         |            | 700ps |

- Maximum clock frequency
  - $f_{max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
  - E.g.  $f_{max,ALU} = 1/200\text{ps} = 5 \text{ GHz!}$

# Control Logic Design

# Control Logic Truth Table

| Inst[31:0]      | BrEq | BrLT | PCSel | ImmSel | BrUn | ASel | BSel | ALUSel | MemRW | RegWEn | WBSel |     |
|-----------------|------|------|-------|--------|------|------|------|--------|-------|--------|-------|-----|
| <b>add</b>      | *    | *    | +4    |        | *    | *    | Reg  | Reg    | Add   | Read   | 1     | ALU |
| <b>sub</b>      | *    | *    | +4    |        | *    | *    | Reg  | Reg    | Sub   | Read   | 1     | ALU |
| <b>(R-R Op)</b> | *    | *    | +4    |        | *    | *    | Reg  | Reg    | (Op)  | Read   | 1     | ALU |
| <b>addi</b>     | *    | *    | +4    | I      | *    | Reg  | Imm  | Add    | Read  | 1      | ALU   |     |
| <b>lw</b>       | *    | *    | +4    | I      | *    | Reg  | Imm  | Add    | Read  | 1      | Mem   |     |
| <b>sw</b>       | *    | *    | +4    | S      | *    | Reg  | Imm  | Add    | Write | 0      | *     |     |
| <b>beq</b>      | 0    | *    | +4    | B      | *    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>beq</b>      | 1    | *    | ALU   | B      | *    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>bne</b>      | 0    | *    | ALU   | B      | *    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>bne</b>      | 1    | *    | +4    | B      | *    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>blt</b>      | *    | 1    | ALU   | B      | 0    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>bltu</b>     | *    | 1    | ALU   | B      | 1    | PC   | Imm  | Add    | Read  | 0      | *     |     |
| <b>jalr</b>     | *    | *    | ALU   | I      | *    | Reg  | Imm  | Add    | Read  | 1      | PC+4  |     |
| <b>jal</b>      | *    | *    | ALU   | J      | *    | PC   | Imm  | Add    | Read  | 1      | PC+4  |     |
| <b>auipc</b>    | *    | *    | +4    | U      | *    | PC   | Imm  | Add    | Read  | 1      | ALU   |     |

# Control Realization Options

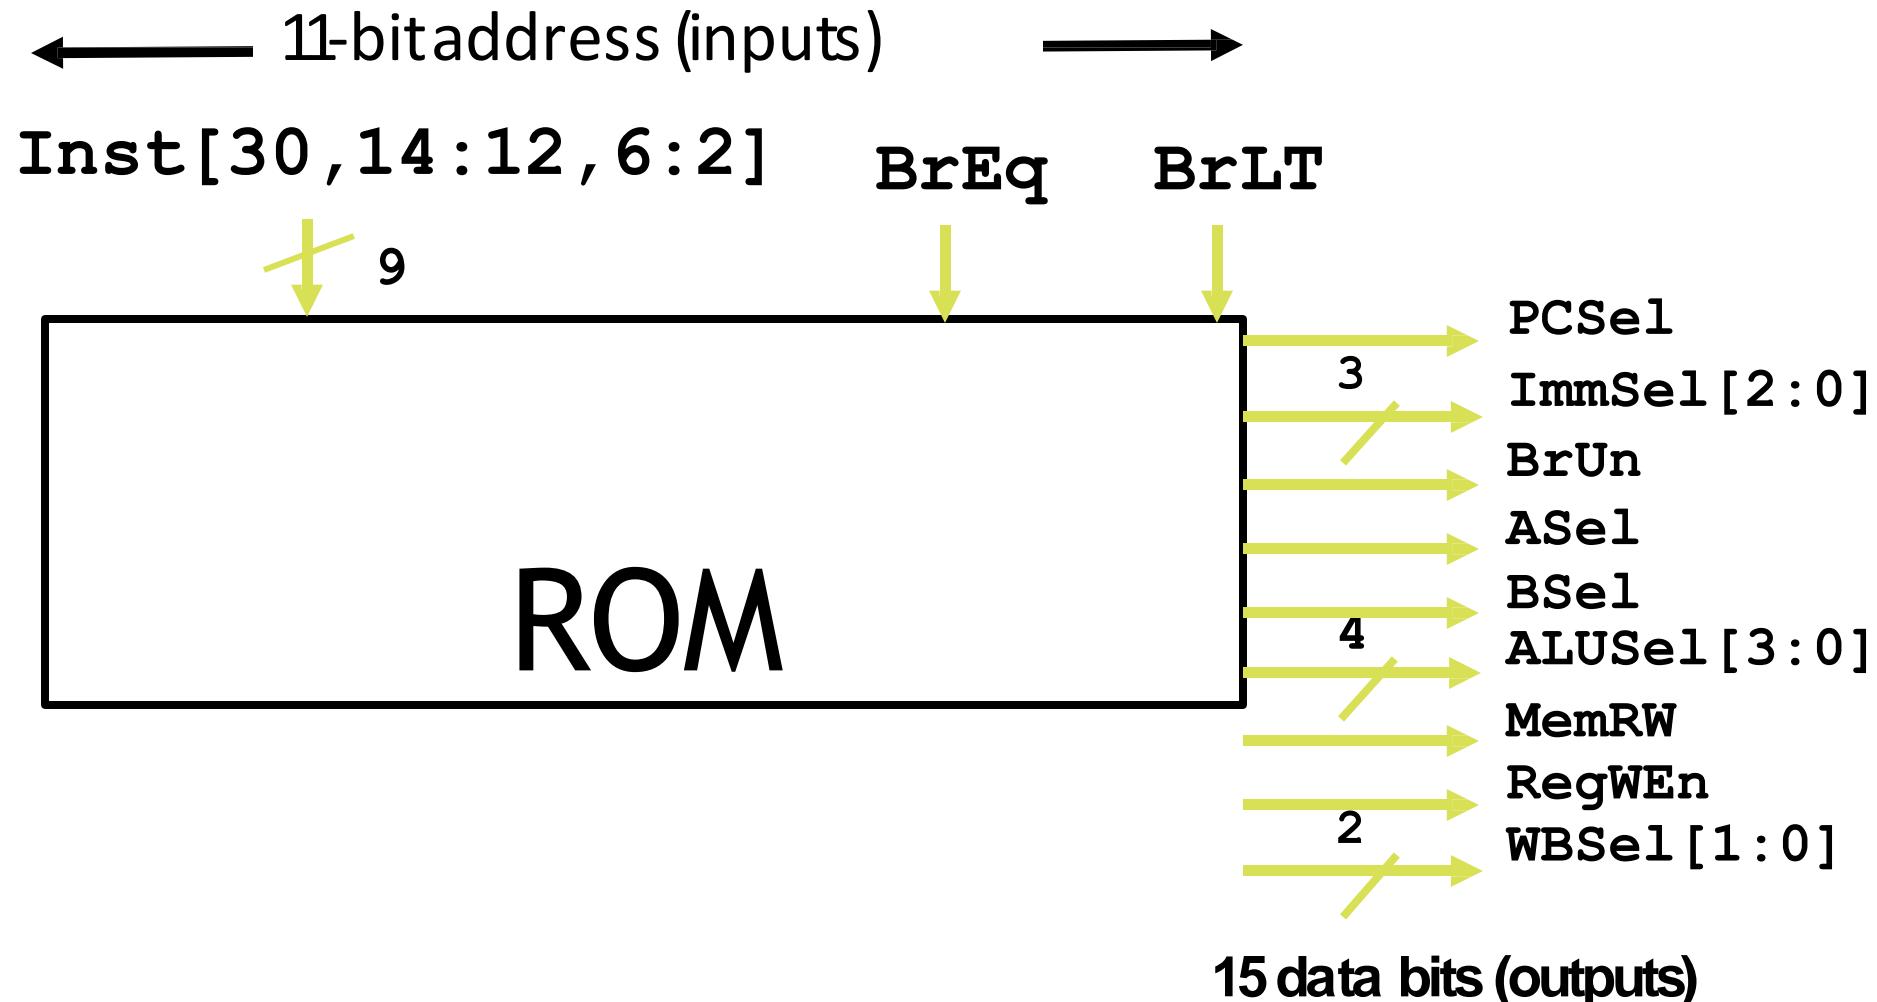
- ROM
  - “Read-Only Memory”
  - Regular structure
  - Can be easily reprogrammed
    - fix errors
    - add instructions
  - Popular when designing control logic manually
- Combinatorial Logic
  - Today, chip designers use logic synthesis tools to convert truth tables to networks of gates

# RV32I, A Nine-Bit ISA!

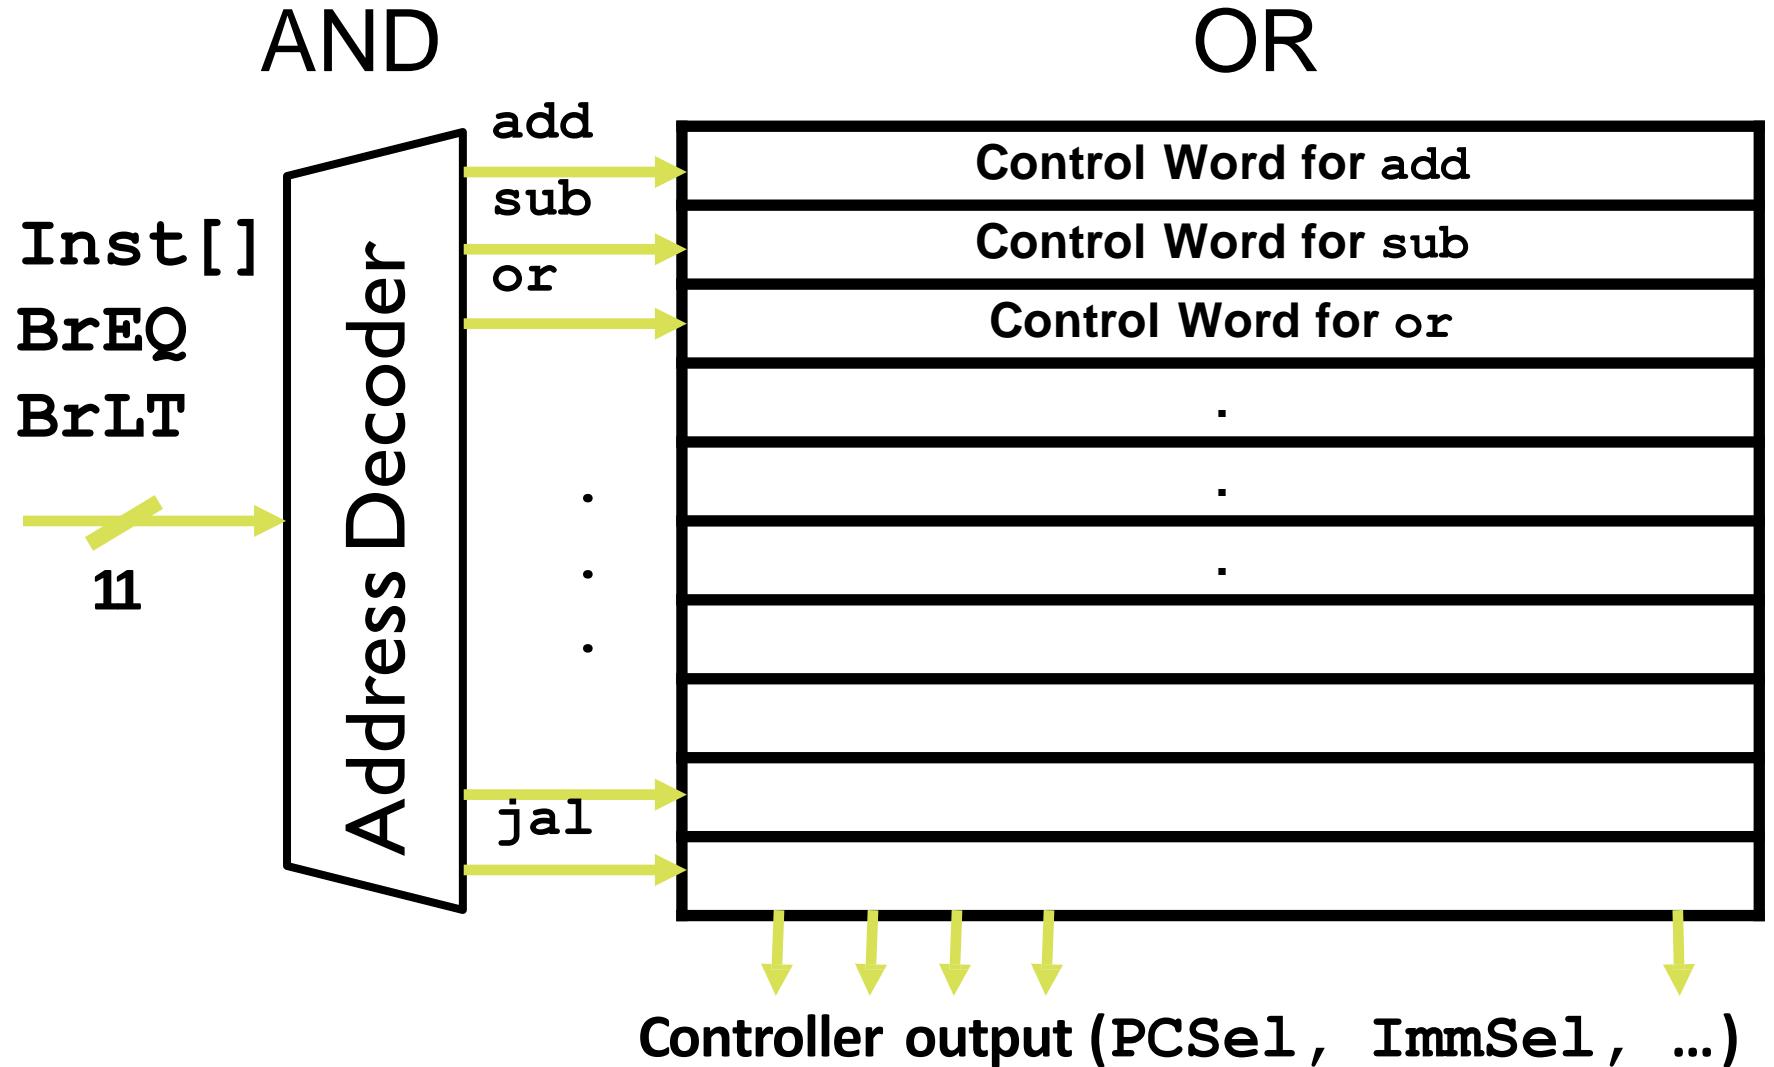
|                       |       |      |            |        |        |
|-----------------------|-------|------|------------|--------|--------|
| imm[31:12]            |       |      | rd         | 011011 | LUI    |
| imm[31:12]            |       |      | rd         | 001011 | AUIPC  |
| imm[20:10:1]11119:12] |       |      | rd         | 110111 | JAL    |
| imm[11:0]             | rs1   | 000  | rd         | 110011 | JALR   |
| imm[12:10:5]          | rs2   | 000  | imm[4:1]11 | 110001 | BEQ    |
| imm[12:10:5]          | rs2   | 001  | Imm[4:1]11 | 110001 | BNE    |
| imm[12:10:5]          | rs2   | 100  | imm[4:1]11 | 110001 | BLT    |
| imm[12:10:5]          | rs2   | 101  | imm[4:1]11 | 110001 | BGE    |
| imm[12:10:5]          | rs2   | 110  | Imm[4:1]11 | 110001 | BLTU   |
| imm[12:10:5]          | rs2   | 111  | imm[4:1]11 | 110001 | BGEU   |
| imm[11:0]             | rs1   | 000  | rd         | 000001 | LB     |
| imm[11:0]             | rs1   | 001  | rd         | 000001 | LH     |
| imm[11:0]             | rs1   | 010  | rd         | 000001 | LW     |
| imm[11:0]             | rs1   | 100  | rd         | 000001 | LBU    |
| imm[11:0]             | rs1   | 101  | rd         | 000001 | LHU    |
| imm[11:5]             | rs2   | 000  | imm[4:0]   | 010001 | SB     |
| imm[11:5]             | rs2   | 001  | imm[4:0]   | 010001 | SH     |
| imm[11:5]             | rs2   | 010  | imm[4:0]   | 010001 | SW     |
| imm[11:0]             | rs1   | 000  | rd         | 001001 | ADDI   |
| imm[11:0]             | rs1   | 010  | rd         | 001001 | SLTI   |
| imm[11:0]             | rs1   | 011  | rd         | 001001 | SLTIU  |
| imm[11:0]             | rs1   | 100  | rd         | 001001 | XORI   |
| imm[11:0]             | rs1   | 110  | rd         | 001001 | ORI    |
| imm[11:0]             | rs1   | 111  | rd         | 001001 | ANDI   |
| 000000                | shamt | rs1  | 001        | rd     | SLLI   |
| 000000                | shamt | rs1  | 101        | rd     | SRLI   |
| 000000                | shamt | rs1  | 101        | rd     | SRAT   |
| 000000                | rs2   | rs1  | 000        | rd     | ADD    |
| 000000                | rs2   | rs1  | 000        | rd     | SUB    |
| 000000                | rs2   | rs1  | 001        | rd     | SLL    |
| 000000                | rs2   | rs1  | 010        | rd     | SLT    |
| 000000                | rs2   | rs1  | 011        | rd     | SLTU   |
| 000000                | rs2   | rs1  | 100        | rd     | XOR    |
| 000000                | rs2   | rs1  | 101        | rd     | SRL    |
| 000000                | rs2   | rs1  | 110        | rd     | SRA    |
| 000000                | rs2   | rs1  | 111        | rd     | OR     |
| fm                    | pred  | succ | rs1        | 000    | AND    |
| 0000000000000000      |       |      | 00000      | 000    | FENCE  |
| 0000000000000001      |       |      | 00000      | 000    | ECALL  |
|                       |       |      | 00000      | 000    | EBREAK |

- Instruction type encoded using only 9 bits:  
**inst[30],**  
**inst[14:12],**  
**inst[6:2]**
- inst[6:2]**
- inst[14:12]**
- inst[30]**

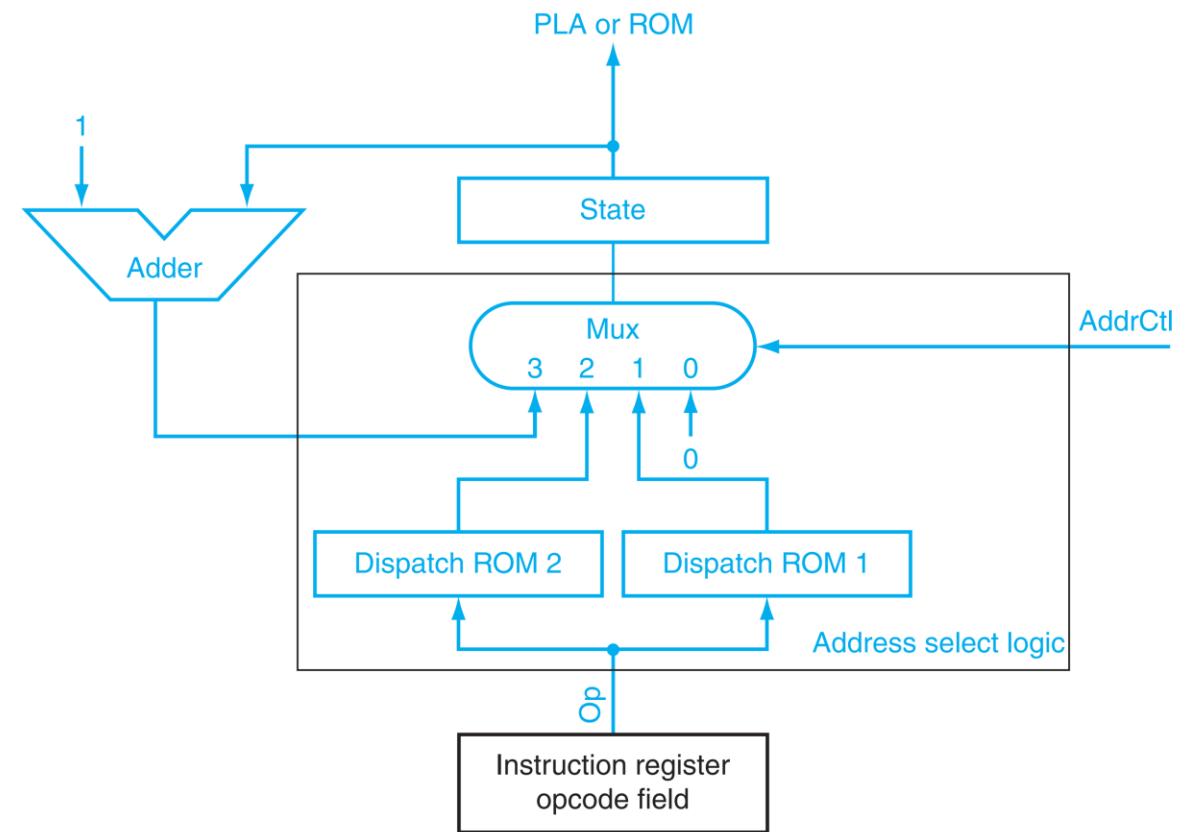
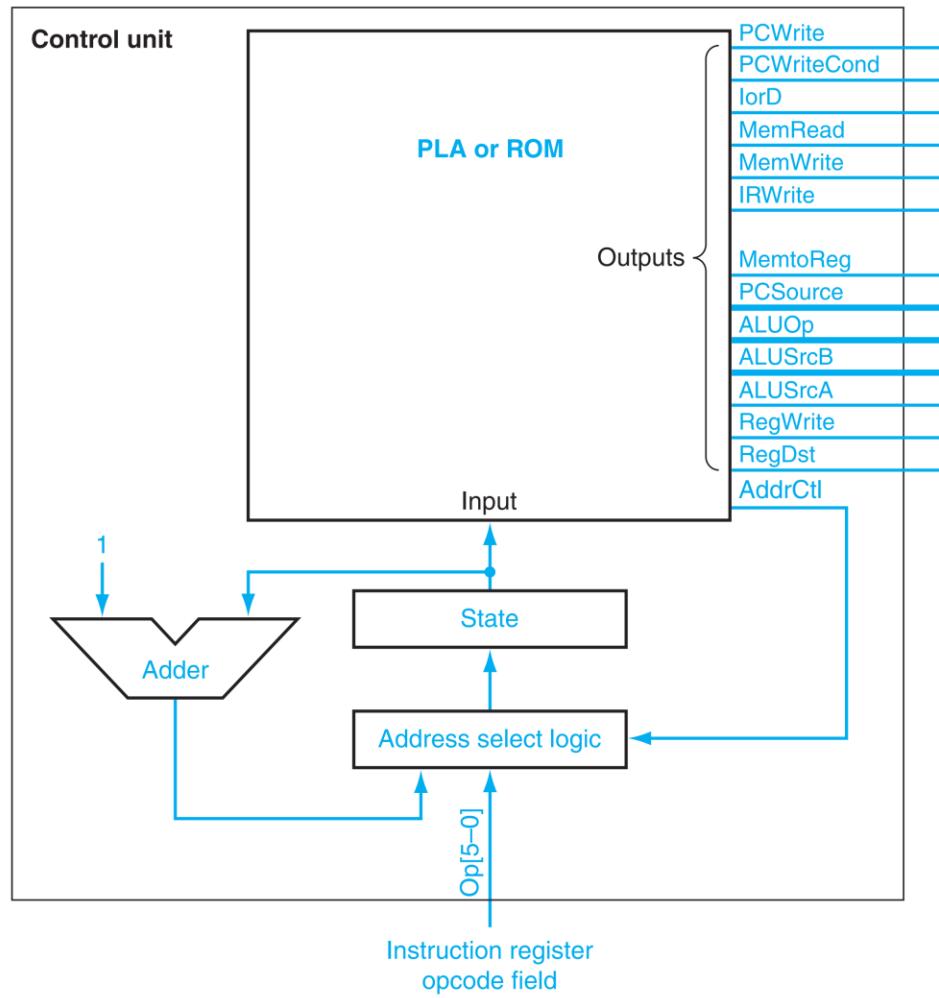
# ROM-based Control

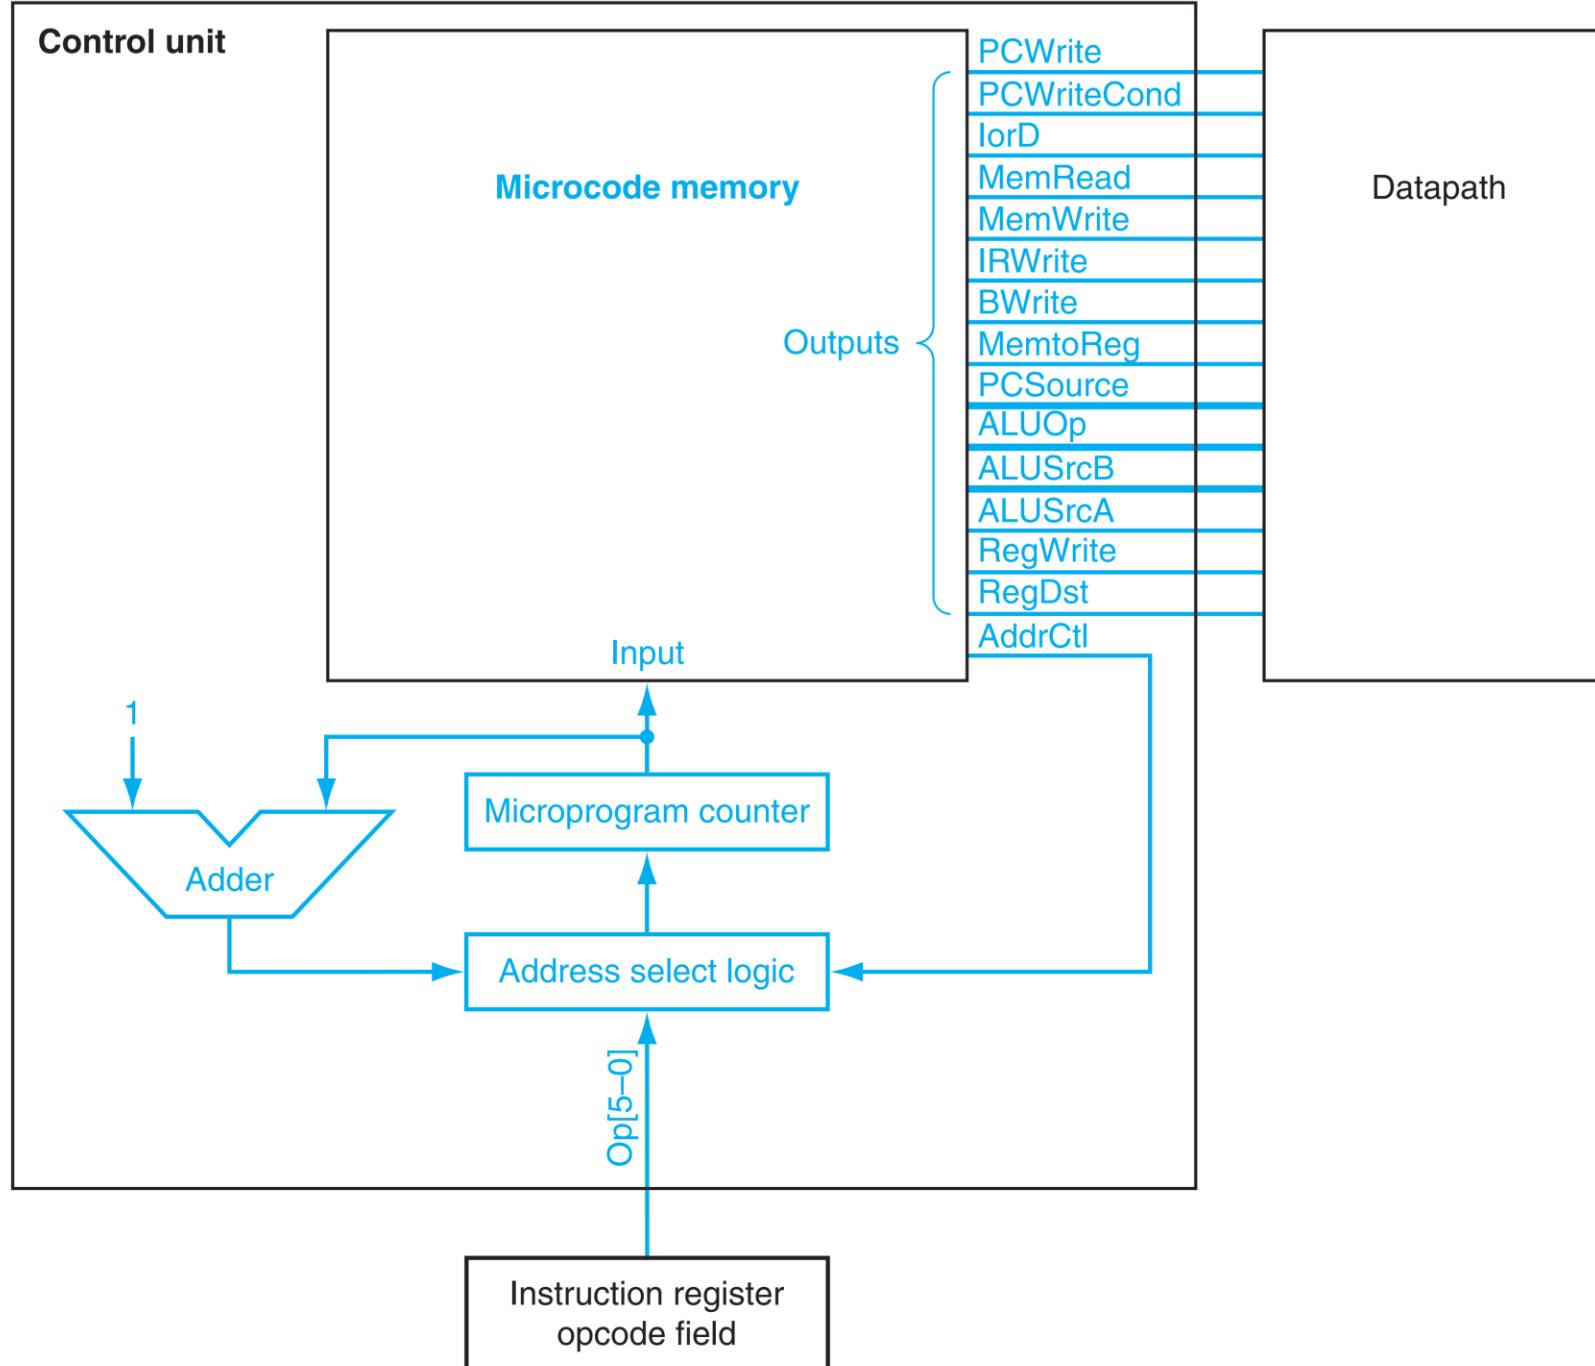


# ROM Controller Implementation



The control unit using an explicit counter to compute the next state





# Combinational Logic Control

- Simplest example: BrUn

|              |     |     | inst[14:12] | inst[6:2]   |         |      |
|--------------|-----|-----|-------------|-------------|---------|------|
| imm[12 10:5] | rs2 | rs1 | 000         | imm[4:1 11] | 1100011 | BE Q |
| imm[12 10:5] | rs2 | rs1 | 001         | imm[4:1 11] | 1100011 | BNE  |
| imm[12 10:5] | rs2 | rs1 | 100         | imm[4:1 11] | 1100011 | BLT  |
| imm[12 10:5] | rs2 | rs1 | 101         | imm[4:1 11] | 1100011 | BGE  |
| imm[12 10:5] | rs2 | rs1 | 110         | imm[4:1 11] | 1100011 | BLTU |
| imm[12 10:5] | rs2 | rs1 | 111         | imm[4:1 11] | 1100011 | BGEU |

- How to decode whether BrUn is 1?

**BrUn = Inst [13] • Branch**

# Control Logic to Decode add

$\text{add} = i[30] \cdot i[14] \cdot i[13] \cdot i[12] \cdot \text{R-type}$   
inst[30] inst[14:12] inst[6:2]

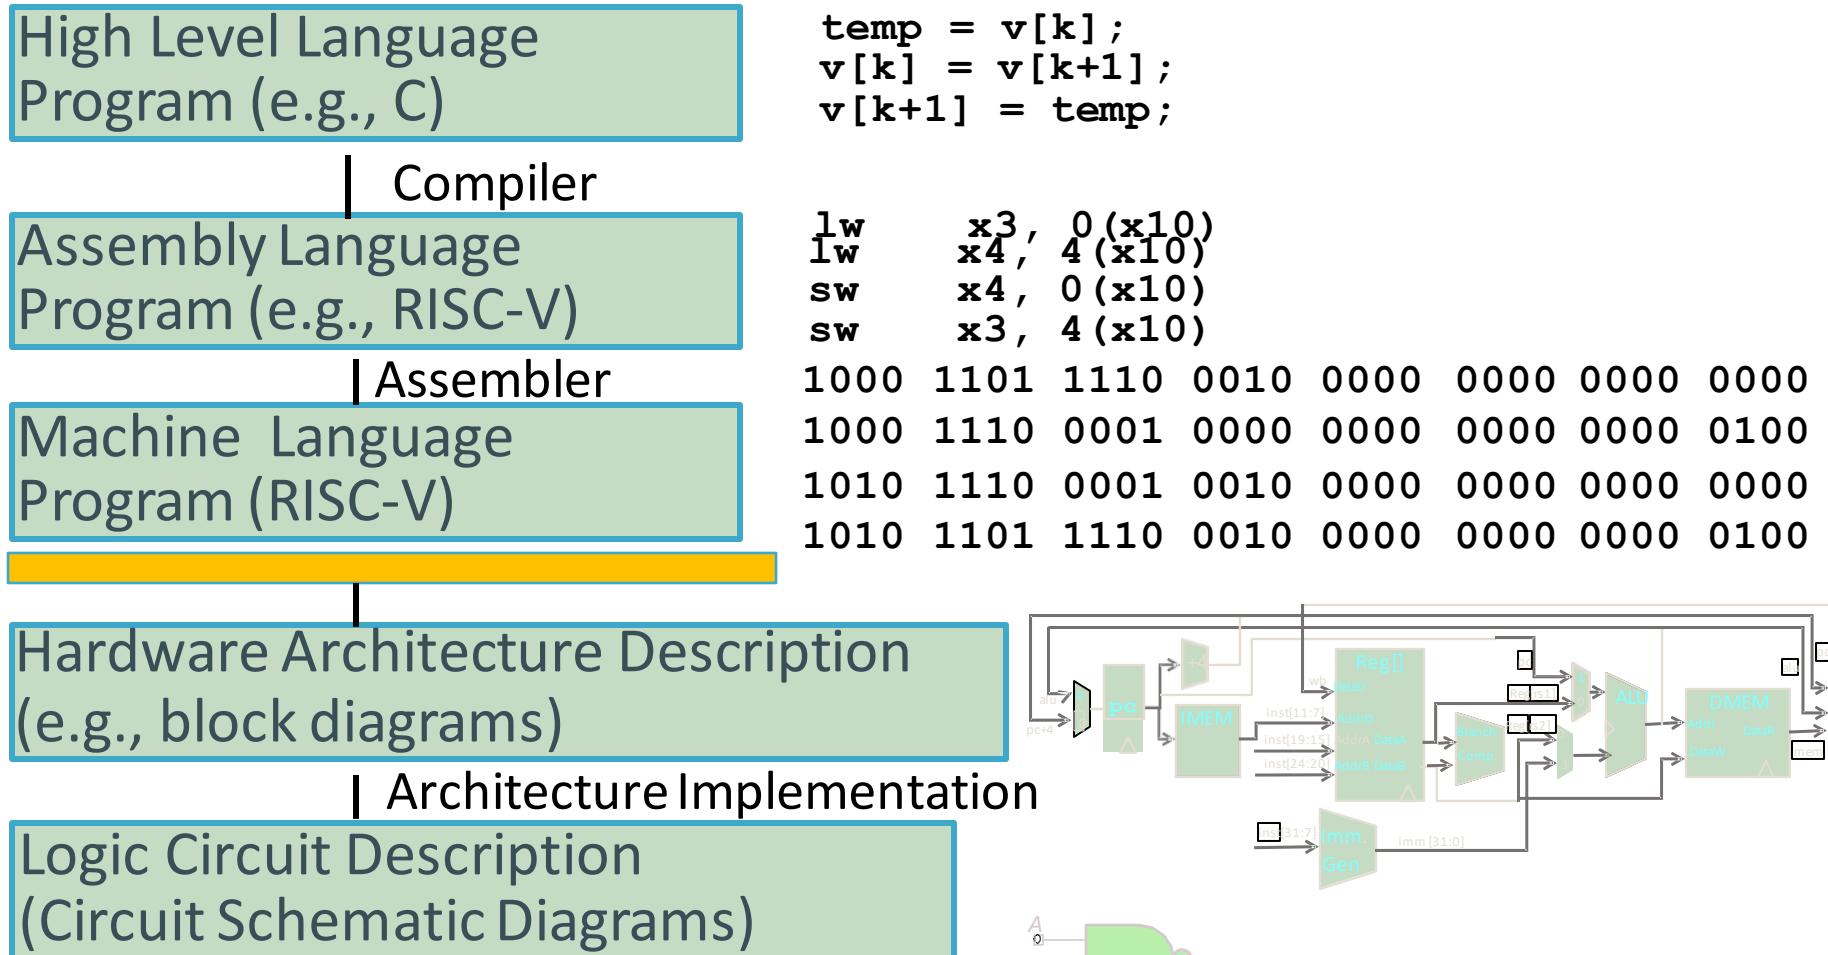
|        |       |     |     |    |         |      |
|--------|-------|-----|-----|----|---------|------|
| 000000 | shamt | rs1 | 001 | rd | 0010011 | SLLI |
| 000000 | shamt | rs1 | 101 | rd | 0010011 | SRLI |
| 010000 | shamt | rs1 | 101 | rd | 0010011 | SRAI |
| 000000 | rs2   | rs1 | 000 | rd | 0110011 | ADD  |
| 010000 | rs2   | rs1 | 000 | rd | 0110011 | SUB  |
| 000000 | rs2   | rs1 | 001 | rd | 0110011 | SLL  |
| 000000 | rs2   | rs1 | 010 | rd | 0110011 | SLT  |
| 000000 | rs2   | rs1 | 011 | rd | 0110011 | SLTU |
| 000000 | rs2   | rs1 | 100 | rd | 0110011 | XOR  |
| 000000 | rs2   | rs1 | 101 | rd | 0110011 | SRL  |
| 010000 | rs2   | rs1 | 101 | rd | 0110011 | SRA  |
| 000000 | rs2   | rs1 | 110 | rd | 0110011 | OR   |
| 000000 | rs2   | rs1 | 111 | rd | 0110011 | AND  |

R-type =  $\overline{i[6]} \cdot i[5] \cdot i[4] \cdot \overline{i[3]} \cdot \overline{i[2]} \cdot \text{RV32I}$

RV32I =  $i[1] \cdot i[0]$

**“And In  
Conclusion...”**

# Call home, we've made HW/SW contact!



# “And In conclusion...”

- We have built a processor!
  - Capable of executing all RISC-V instructions in one cycle each
  - Not all units (hardware) used by all instructions
  - Critical path changes
- 5 Phases of execution
  - IF, ID, EX, MEM, WB
  - Not all instructions are active in all phases
- Controller specifies how to execute instructions
  - Implemented as ROM or logic